



UNIVERSITY OF TECHNOLOGY  
IN THE EUROPEAN CAPITAL OF CULTURE  
CHEMNITZ

# Neurocomputing

Linear classification

Julien Vitay

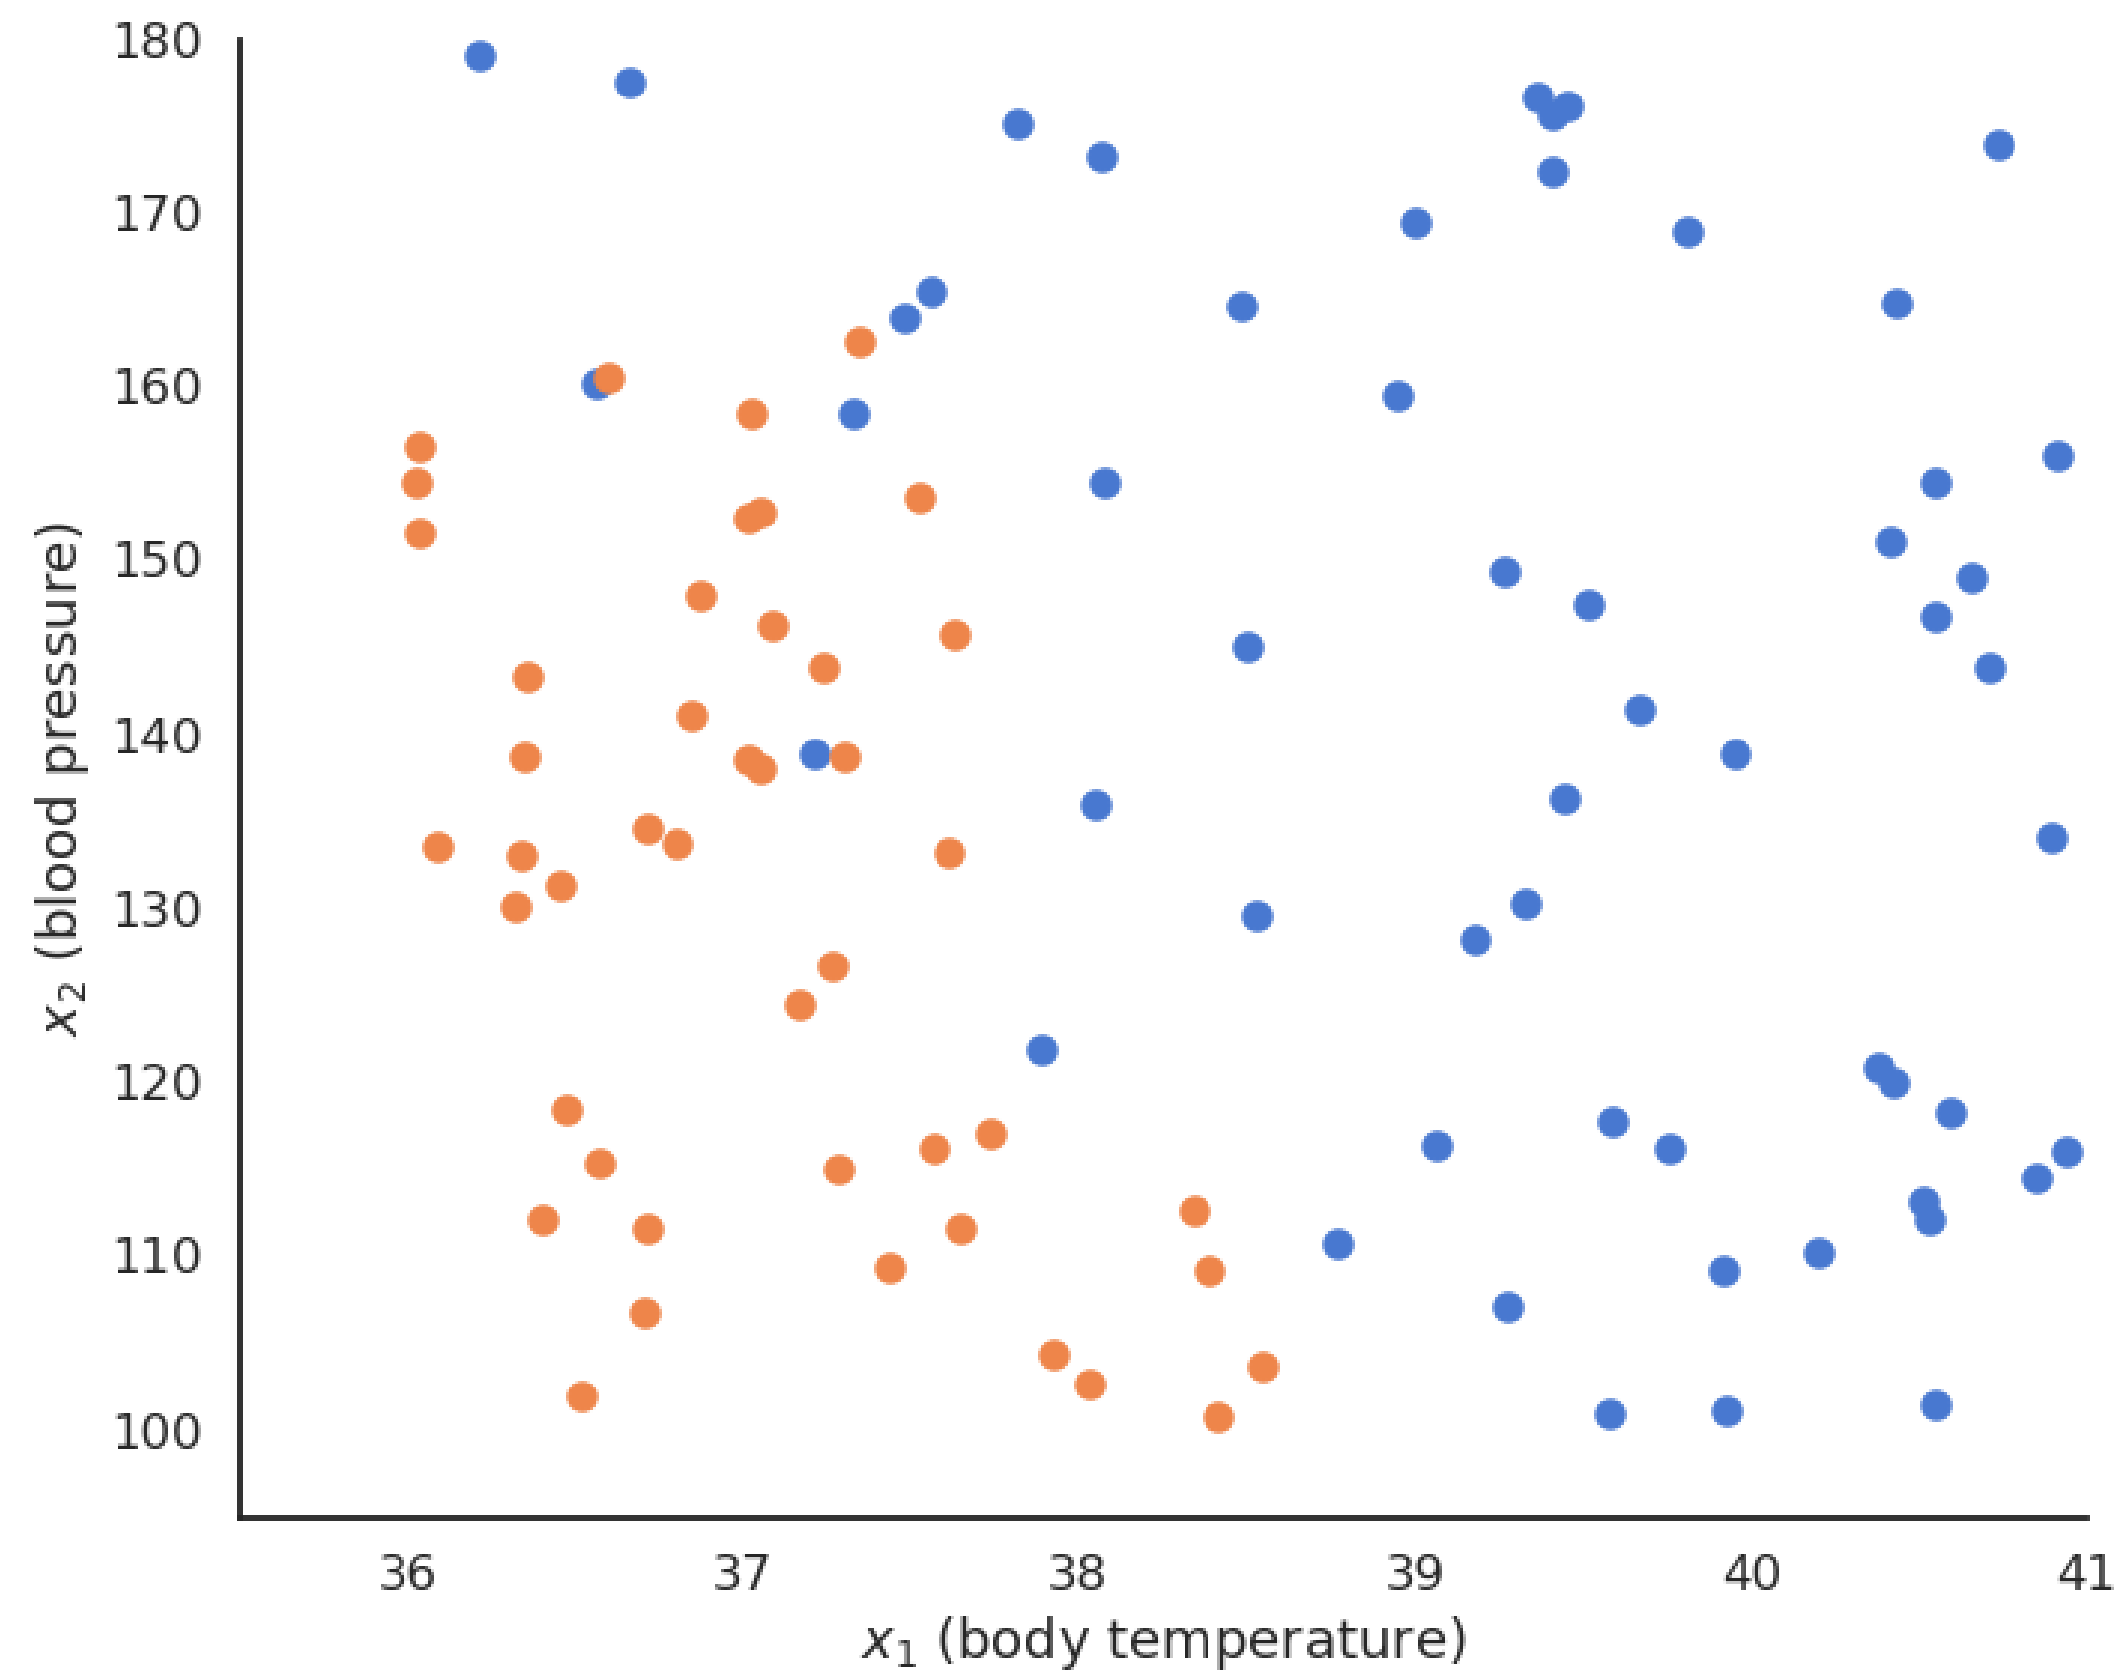
Professur für Künstliche Intelligenz - Fakultät für Informatik

<https://tu-chemnitz.de/informatik/KI/edu/neurocomputing>

# 1 - Hard linear classification

# Binary classification

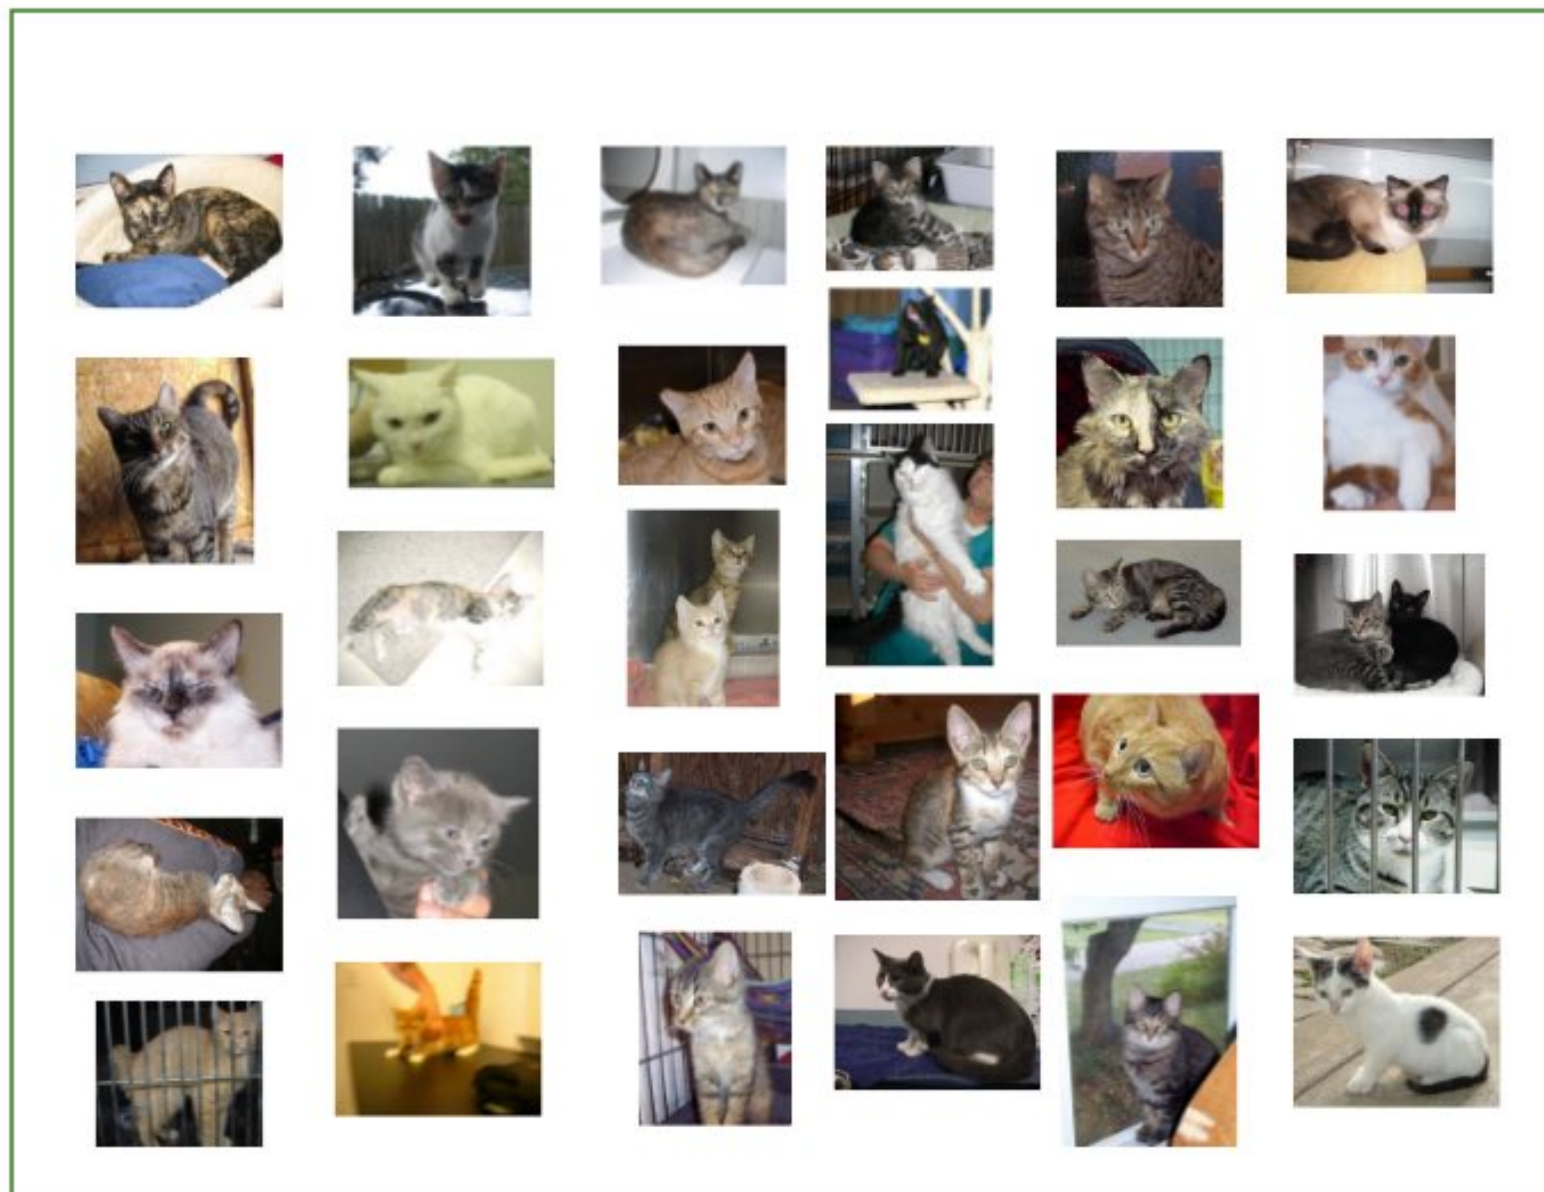
- The training data  $\mathcal{D}$  is composed of  $N$  examples  $(\mathbf{x}_i, t_i)_{i=1..N}$ , with a  $d$ -dimensional input vector  $\mathbf{x}_i \in \mathbb{R}^d$  and a binary output  $t_i \in \{-1, +1\}$
- The data points where  $t = +1$  are called the **positive class**, the other the **negative class**.



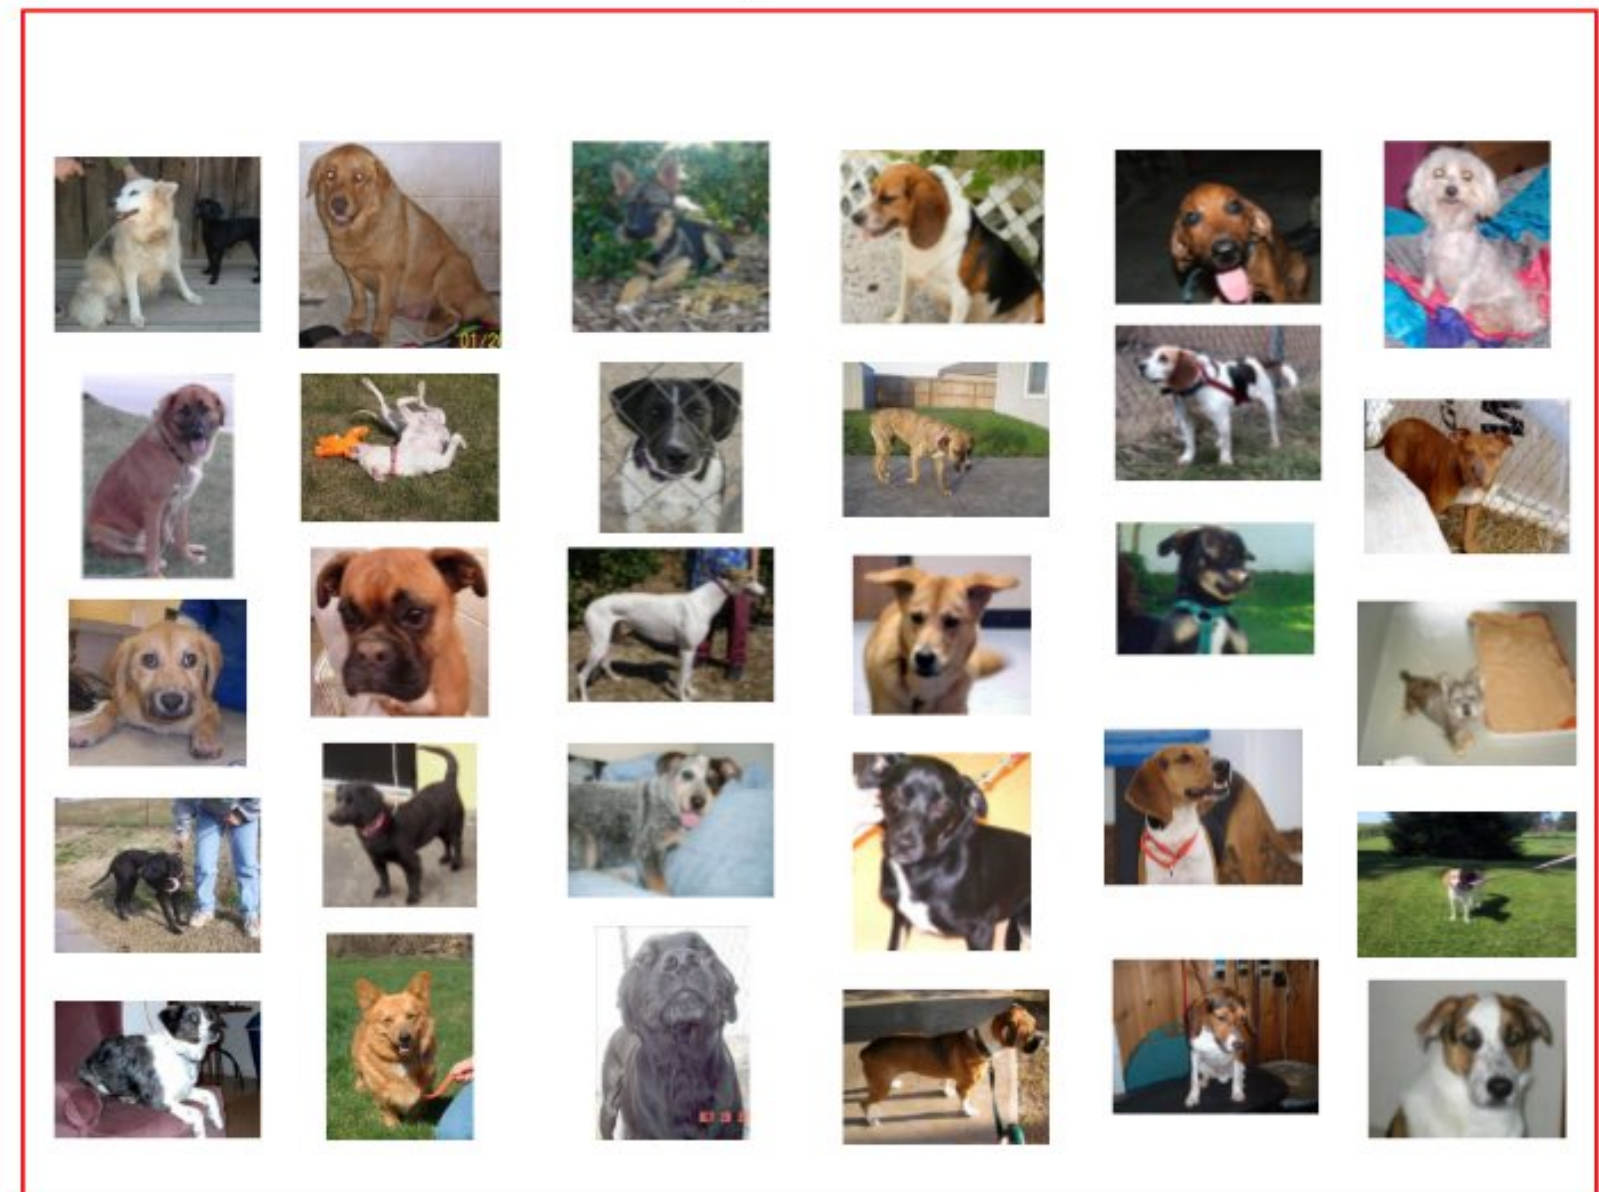
# Binary classification

- For example, the inputs  $\mathbf{x}_i$  can be images (one dimension per pixel) and the positive class corresponds to cats ( $t_i = +1$ ), the negative class to dogs ( $t_i = -1$ ).

Cats



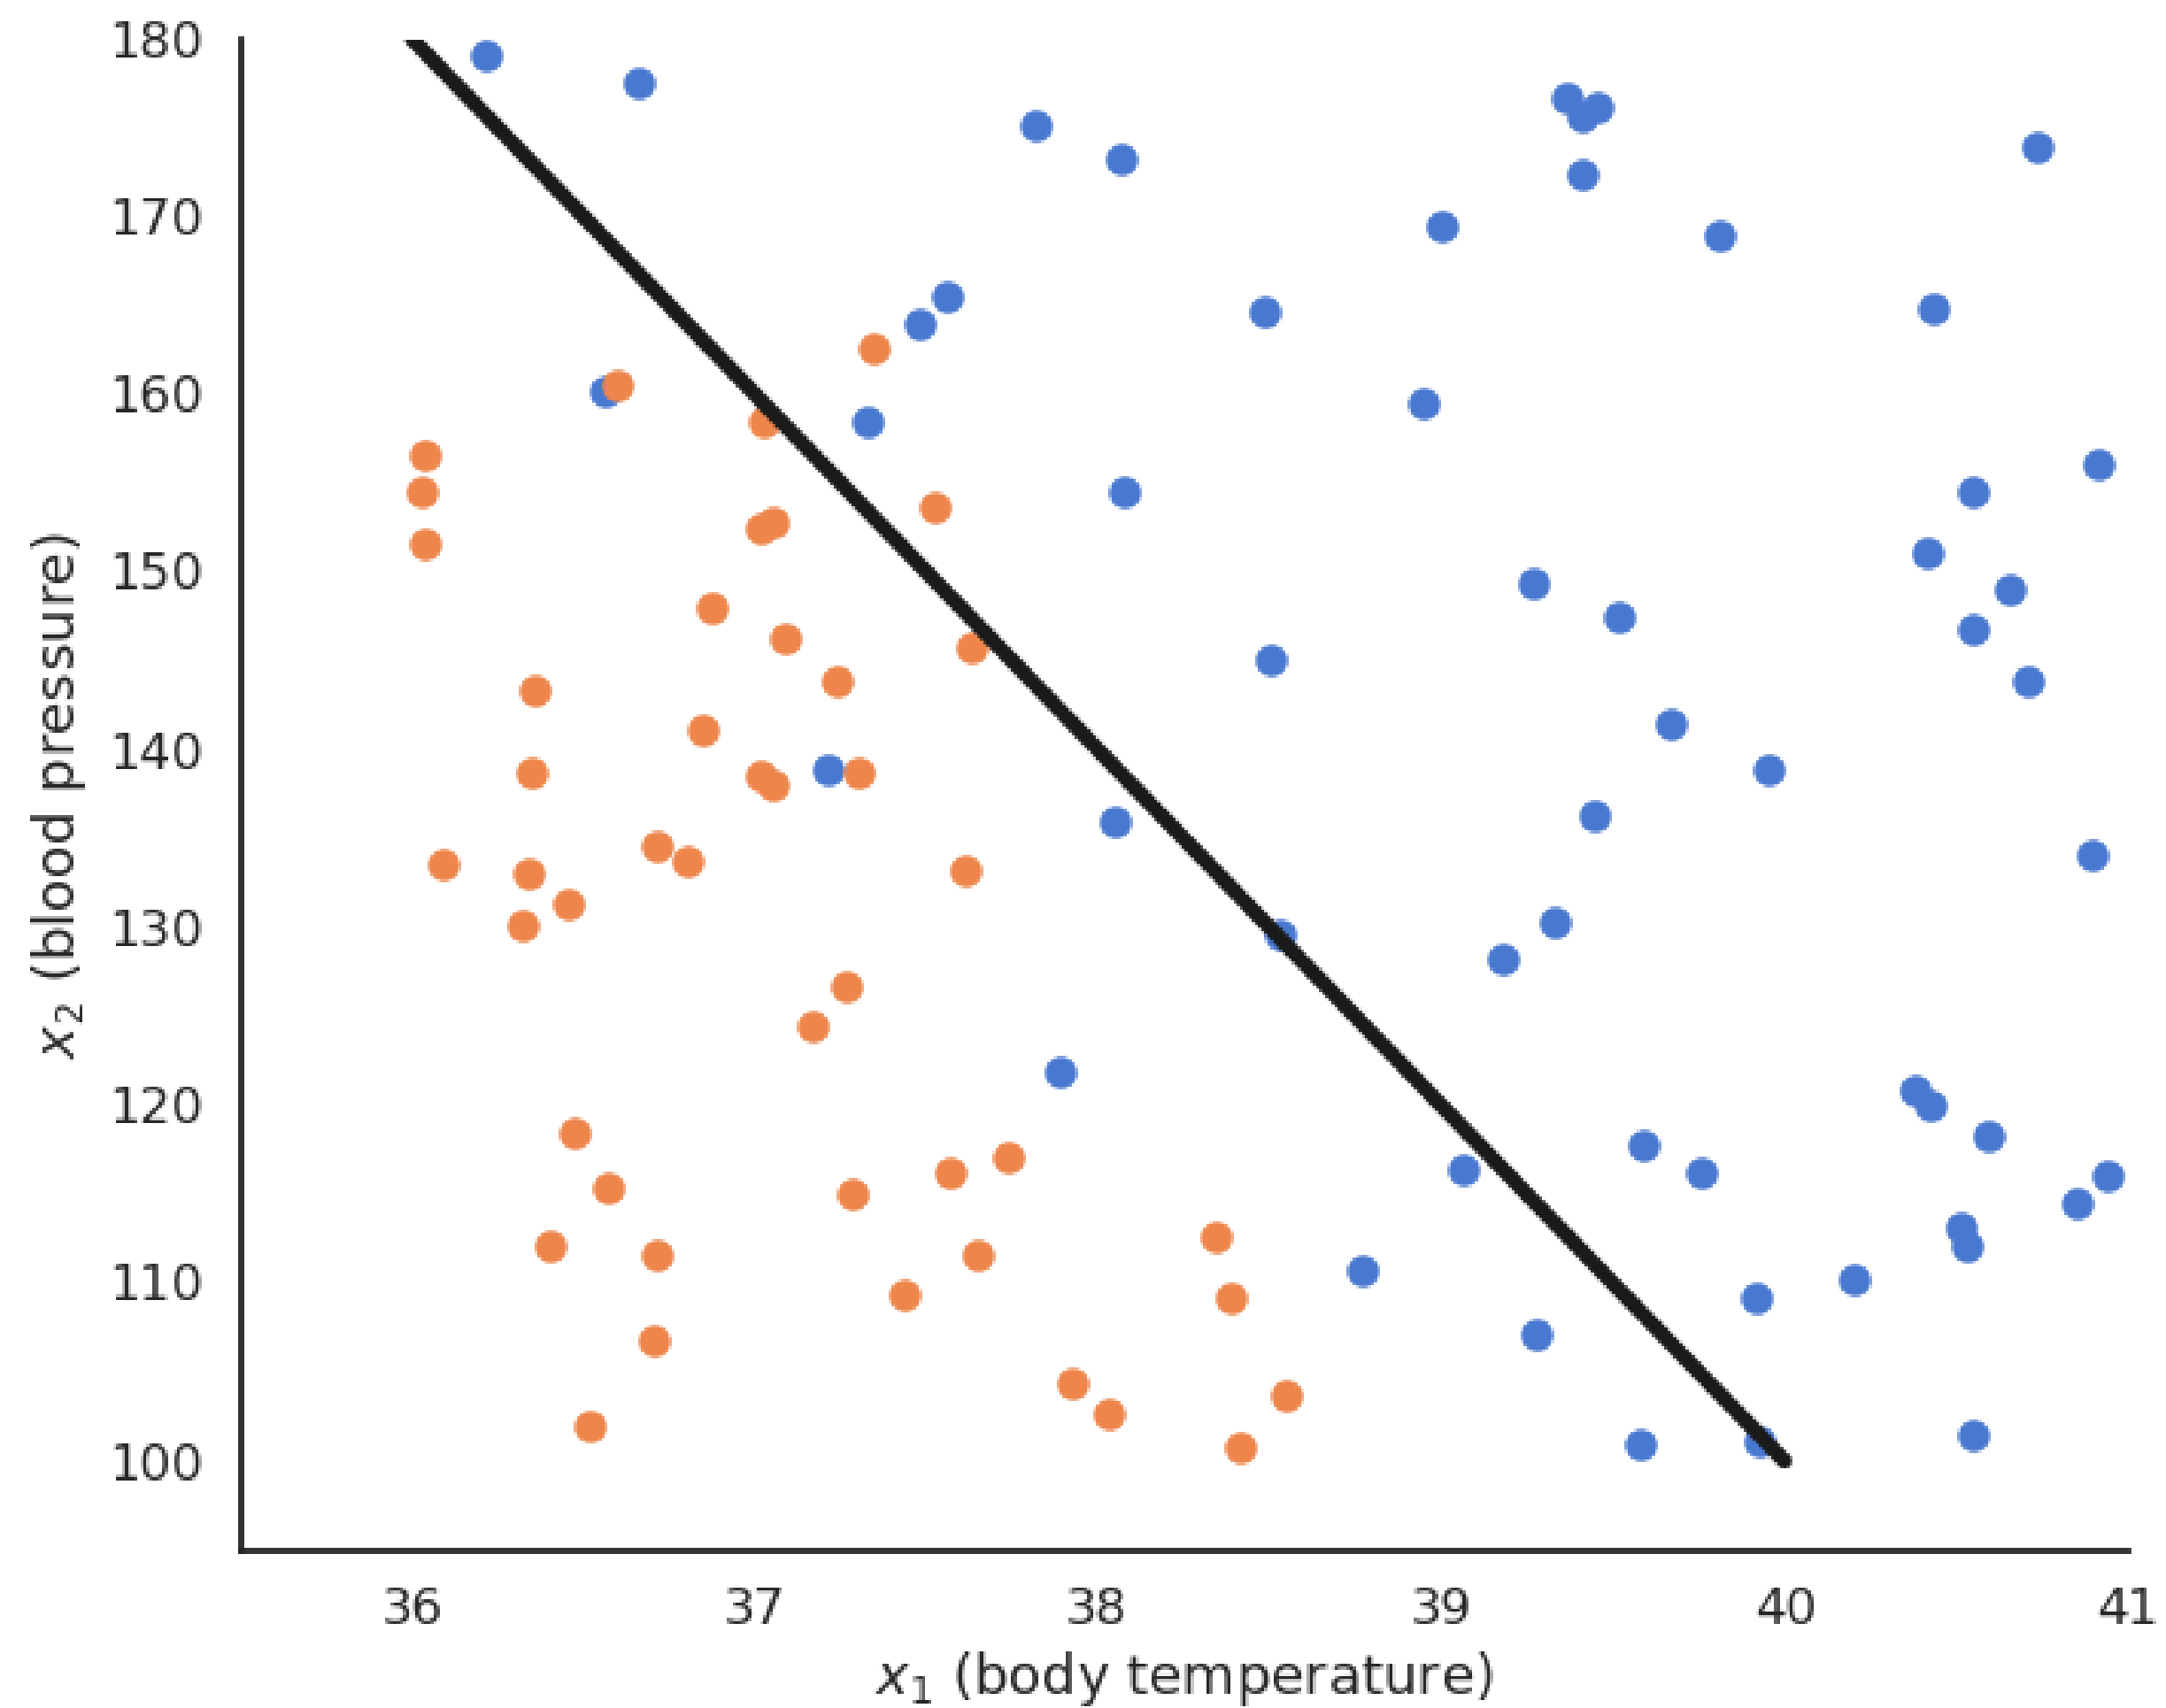
Dogs



Sample of cats & dogs images from Kaggle Dataset

# Binary linear classification


- We want to find the hyperplane  $(\mathbf{w}, b)$  of  $\mathbb{R}^d$  that correctly separates the two classes.



# Binary linear classification

- For a point  $\mathbf{x} \in \mathcal{D}$ ,  $\langle \mathbf{w} \cdot \mathbf{x} \rangle + b$  is the projection of  $\mathbf{x}$  onto the hyperplane  $(\mathbf{w}, b)$ .
  - If  $\langle \mathbf{w} \cdot \mathbf{x} \rangle + b > 0$ , the point is above the hyperplane.
  - If  $\langle \mathbf{w} \cdot \mathbf{x} \rangle + b < 0$ , the point is below the hyperplane.
  - If  $\langle \mathbf{w} \cdot \mathbf{x} \rangle + b = 0$ , the point is on the hyperplane.
- By looking at the **sign** of  $\langle \mathbf{w} \cdot \mathbf{x} \rangle + b$ , we can predict the class of the input:

$$\text{sign}(\langle \mathbf{w} \cdot \mathbf{x} \rangle + b) = \begin{cases} +1 & \text{if } \langle \mathbf{w} \cdot \mathbf{x} \rangle + b \geq 0 \\ -1 & \text{if } \langle \mathbf{w} \cdot \mathbf{x} \rangle + b < 0 \end{cases}$$

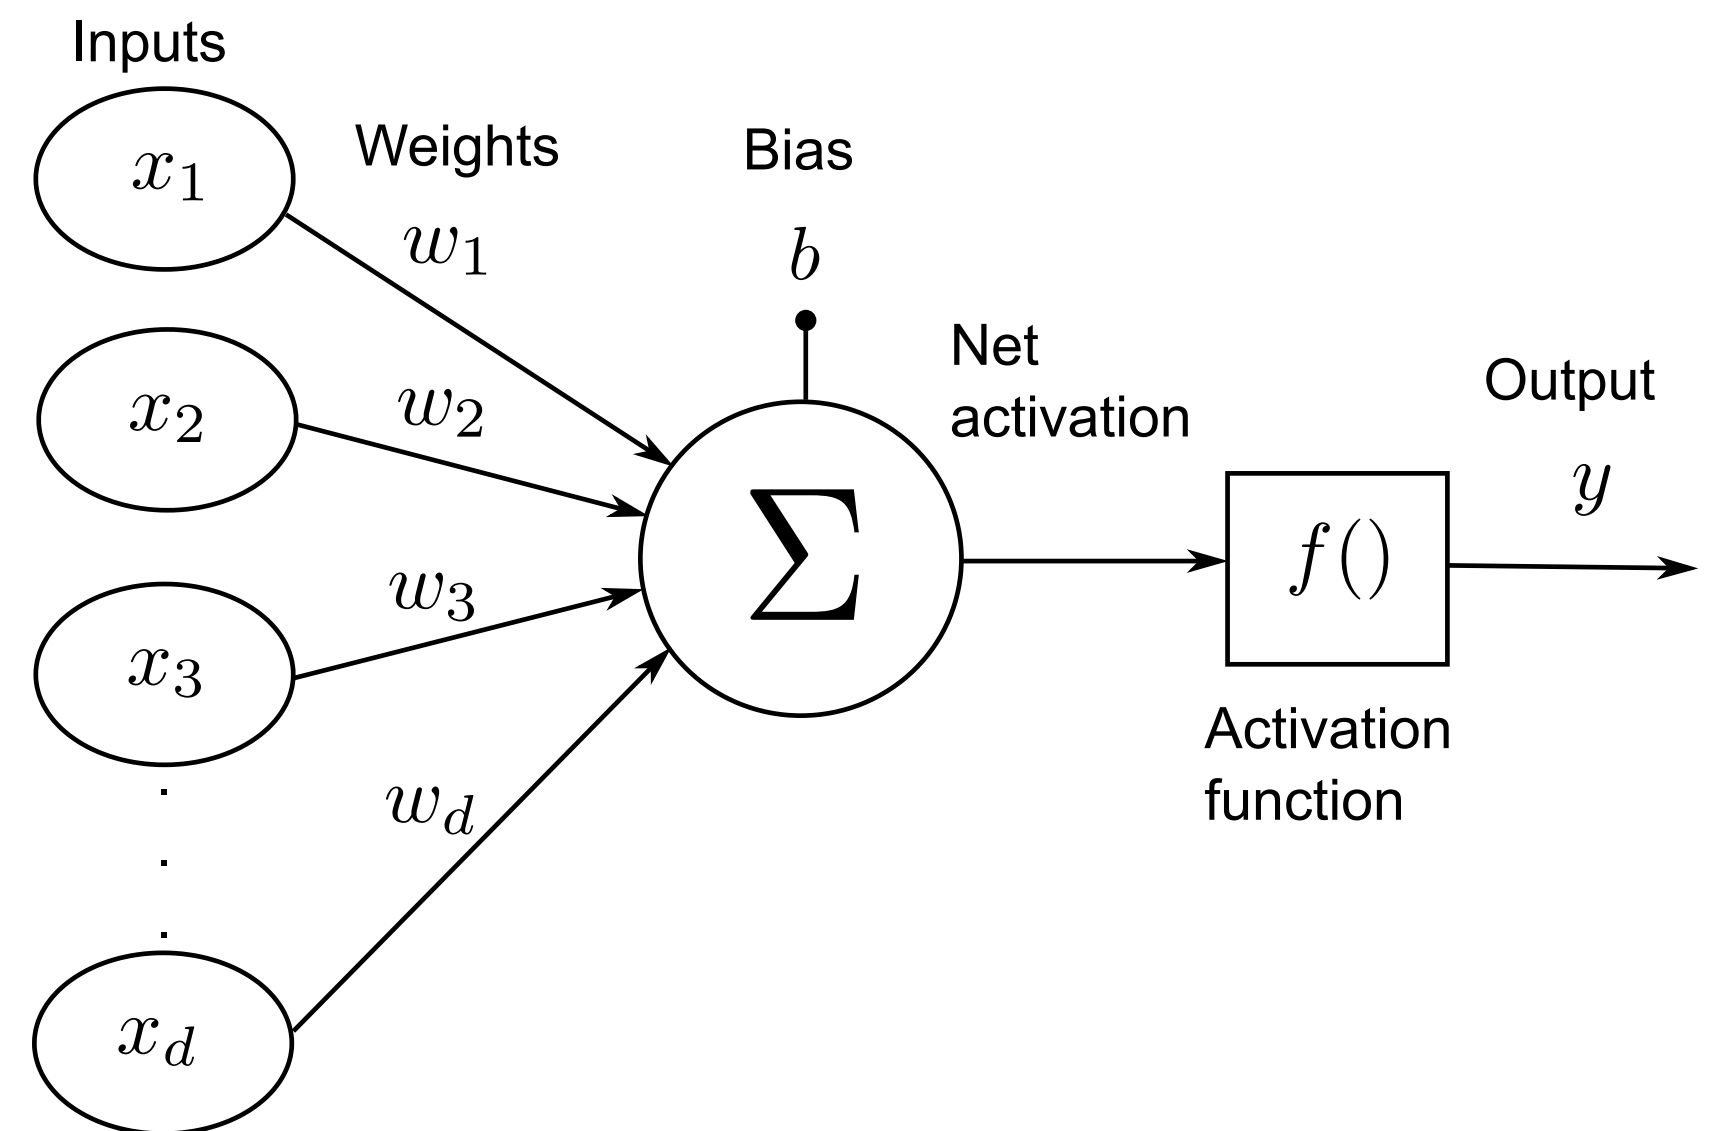
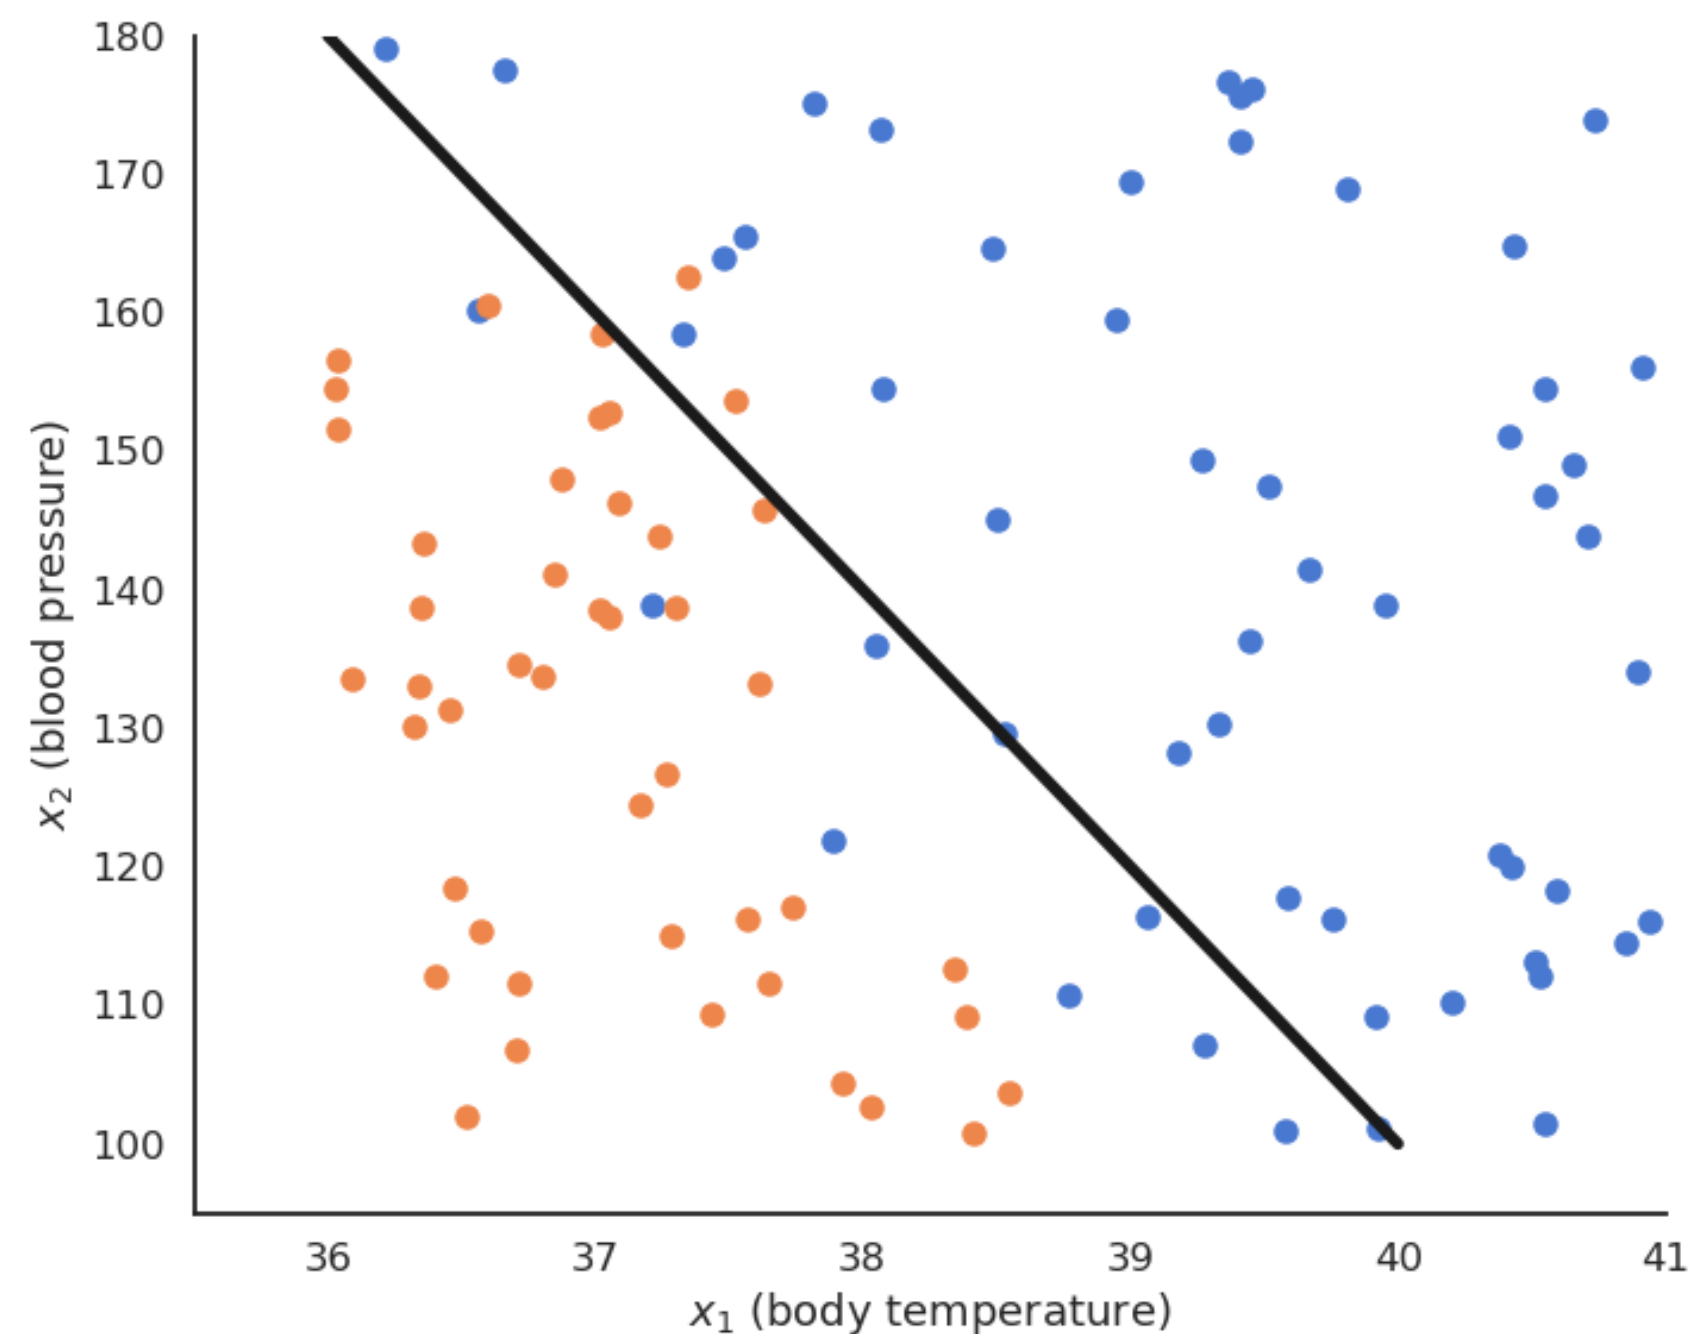




# Binary linear classification

- Binary linear classification can be made by a single **artificial neuron** using the sign transfer function.

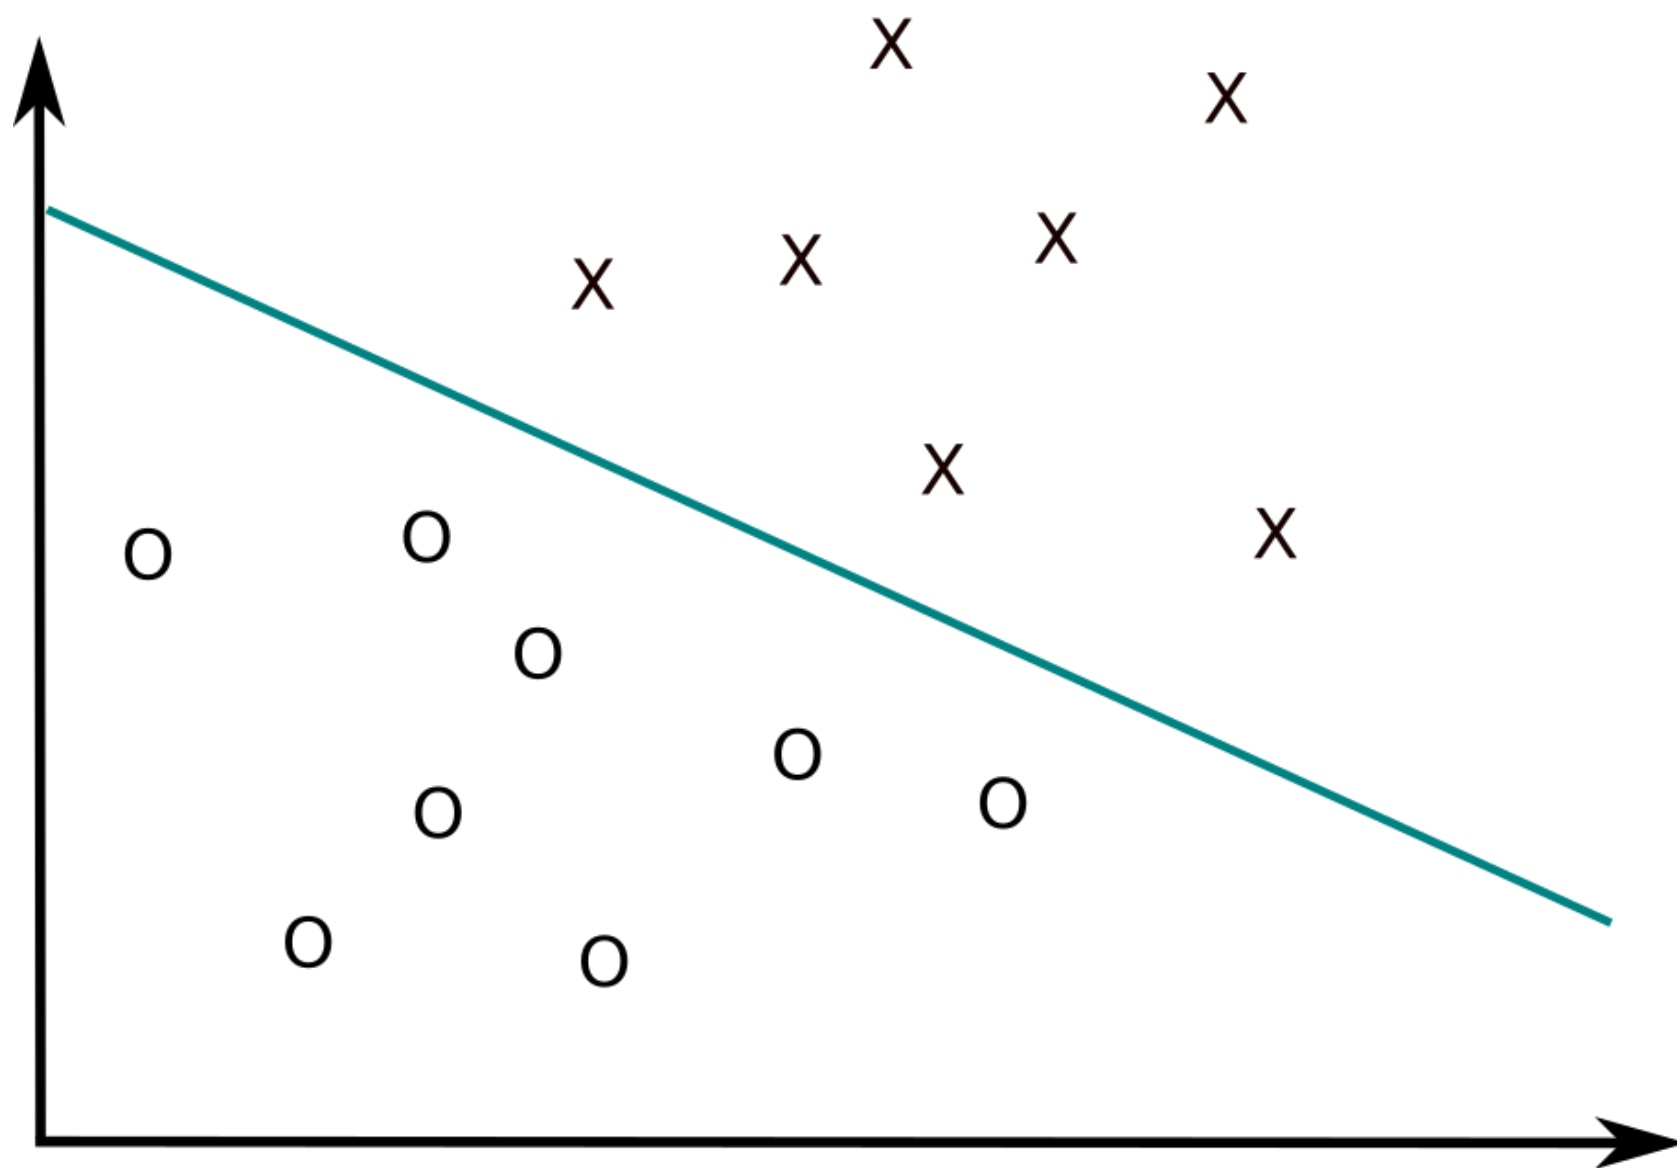
$$y = f_{\mathbf{w},b}(\mathbf{x}) = \text{sign}(\langle \mathbf{w} \cdot \mathbf{x} \rangle + b) = \text{sign}\left(\sum_{j=1}^d w_j x_j + b\right)$$



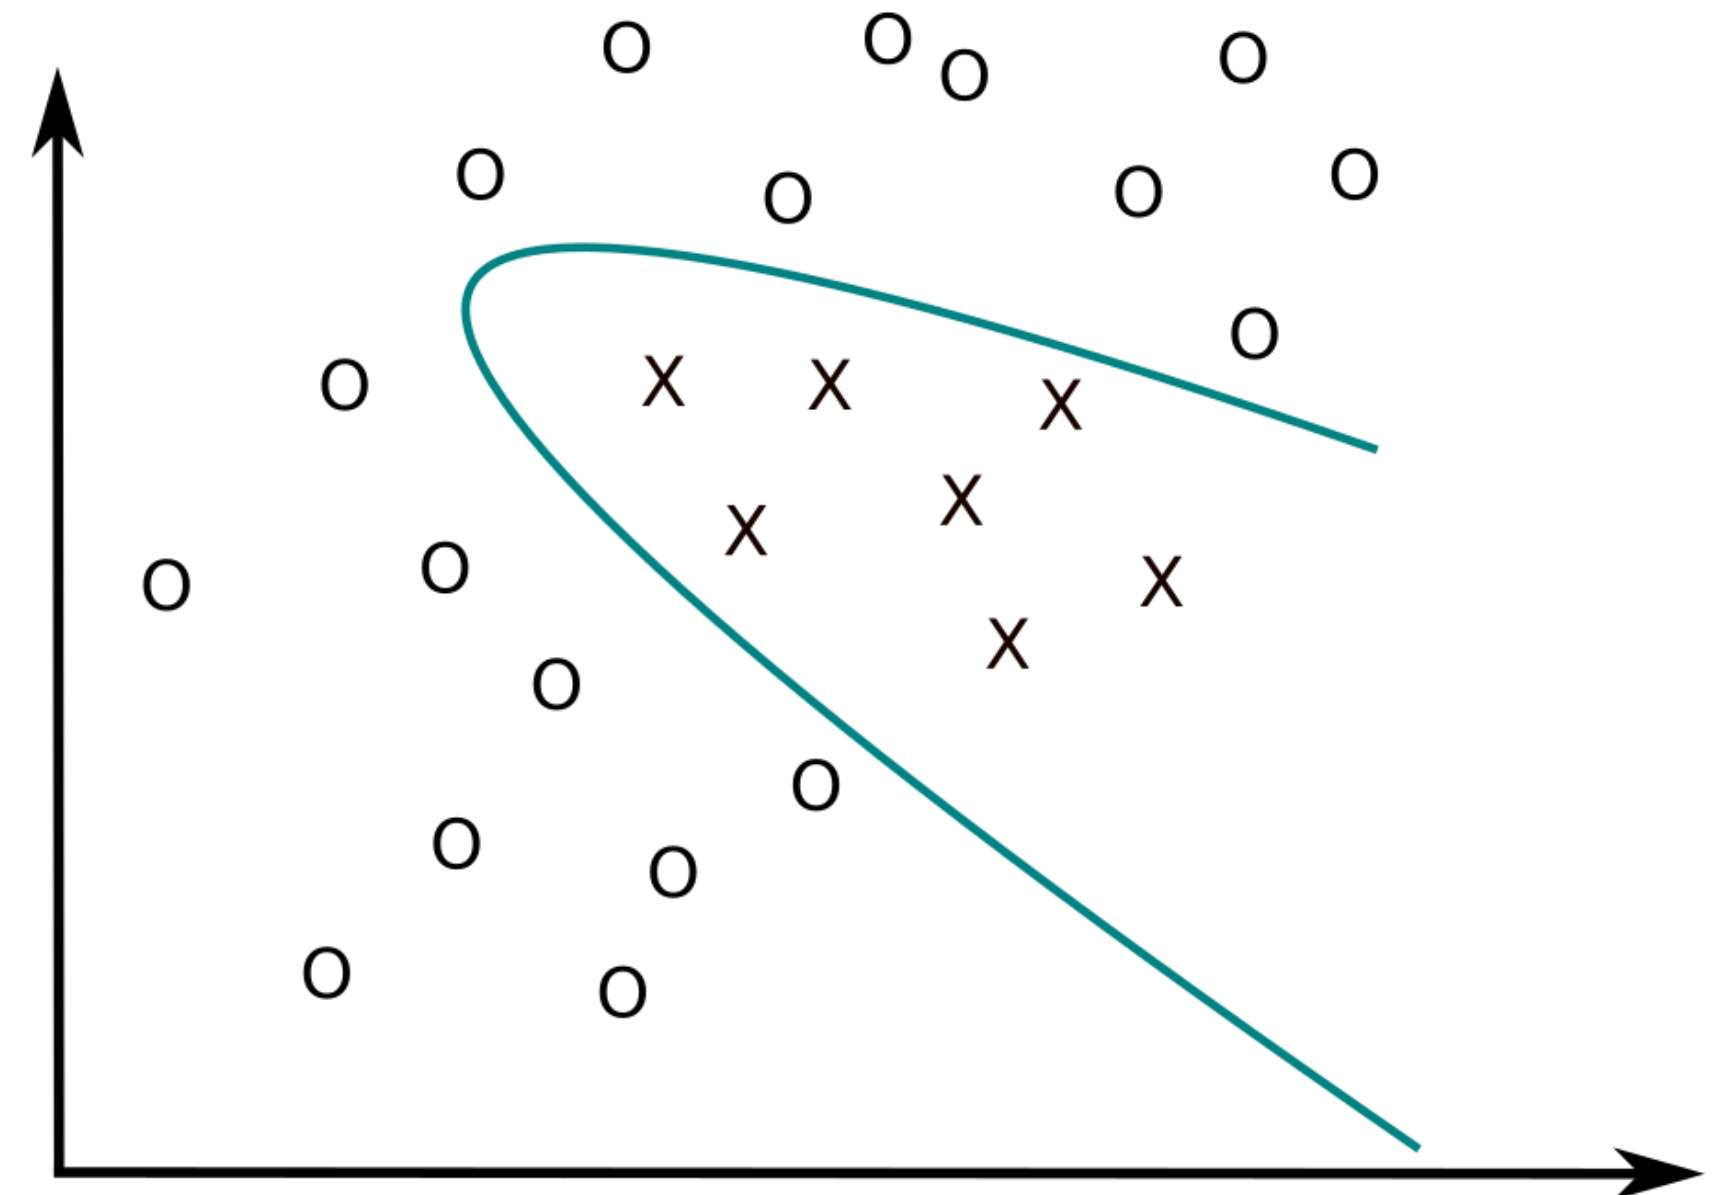
- $\mathbf{w}$  is the weight vector and  $b$  is the bias.

# Linearly separable datasets

Linearly separable



Non-linearly separable



- Linear classification is the process of finding an hyperplane ( $\mathbf{w}, b$ ) that correctly separates the two classes.
- If such an hyperplane can be found, the training set is said **linearly separable**.
- Otherwise, the problem is **non-linearly separable** and other methods have to be applied (MLP, SVM...).



# Linear classification as an optimization problem

- The Perceptron algorithm tries to find the weights and biases minimizing the **mean square error** (*mse*) or **quadratic loss**:

$$\mathcal{L}(\mathbf{w}, b) = \mathbb{E}_{\mathcal{D}}[(t_i - y_i)^2] \approx \frac{1}{N} \sum_{i=1}^N (t_i - y_i)^2$$

- When the prediction  $y_i$  is the same as the data  $t_i$  for all examples in the training set (perfect classification), the mse is minimal and equal to 0.
- We can apply gradient descent to find this minimum.

$$\begin{cases} \Delta \mathbf{w} = -\eta \nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}, b) \\ \Delta b = -\eta \nabla_b \mathcal{L}(\mathbf{w}, b) \end{cases}$$

# Linear classification as an optimization problem

- Let's search for the partial derivative of the quadratic error function with respect to the weight vector:

$$\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}, b) = \nabla_{\mathbf{w}} \frac{1}{N} \sum_{i=1}^N (t_i - y_i)^2 = \frac{1}{N} \sum_{i=1}^N \nabla_{\mathbf{w}} (t_i - y_i)^2 = \frac{1}{N} \sum_{i=1}^N \nabla_{\mathbf{w}} l_i(\mathbf{w}, b)$$

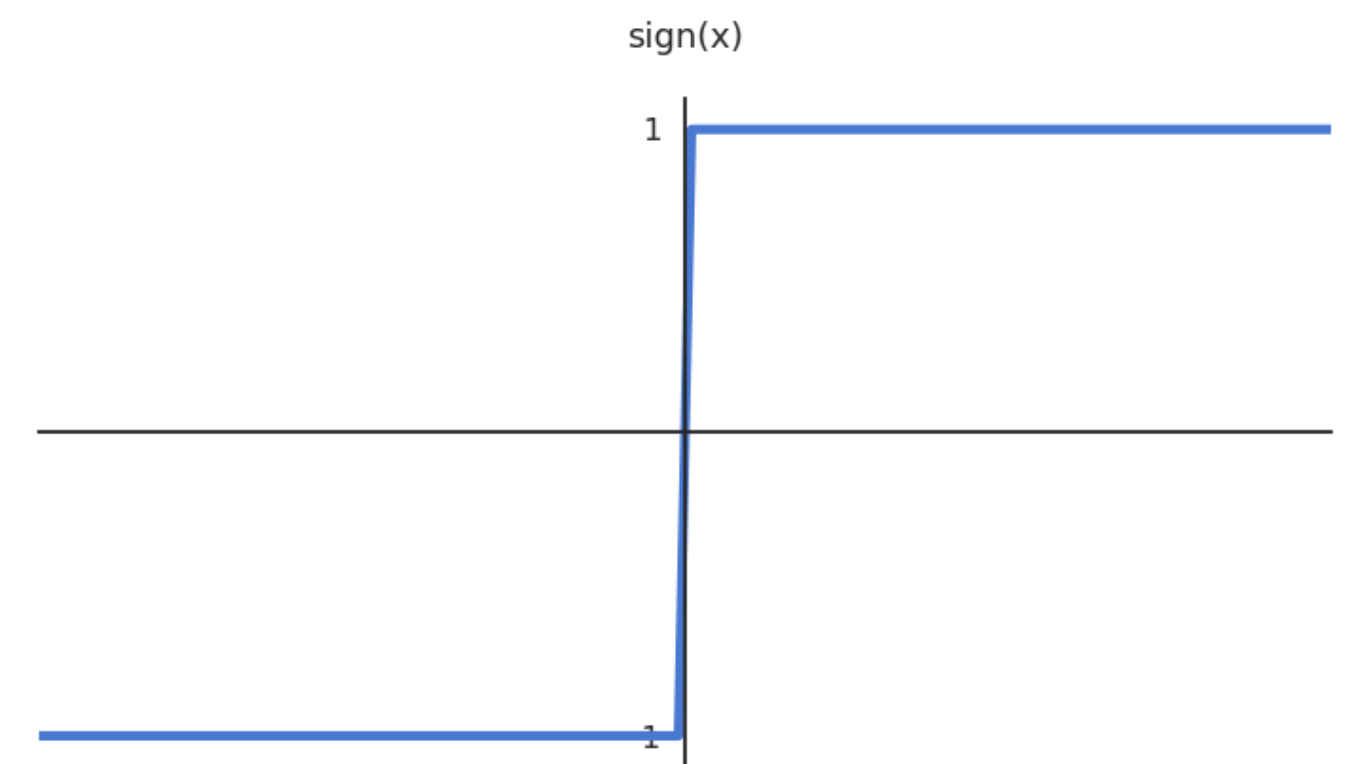
- Everything is similar to linear regression until we get:

$$\nabla_{\mathbf{w}} l_i(\mathbf{w}, b) = -2 (t_i - y_i) \nabla_{\mathbf{w}} \text{sign}(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b)$$

- In order to continue with the chain rule, we would need to differentiate  $\text{sign}(x)$ .

$$\nabla_{\mathbf{w}} l_i(\mathbf{w}, b) = -2 (t_i - y_i) \text{sign}'(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \mathbf{x}_i$$

- But the sign function is **not** differentiable...



# Linear classification as an optimization problem

- We will simply pretend that the  $\text{sign}()$  function is linear, with a derivative of 1:

$$\nabla_{\mathbf{w}} l_i(\mathbf{w}, b) = -2(t_i - y_i) \mathbf{x}_i$$

- The update rule for the weight vector  $\mathbf{w}$  and the bias  $b$  is therefore the same as in linear regression:

$$\begin{cases} \Delta \mathbf{w} = \eta \frac{1}{N} \sum_{i=1}^N (t_i - y_i) \mathbf{x}_i \\ \Delta b = \eta \frac{1}{N} \sum_{i=1}^N (t_i - y_i) \end{cases}$$

# Batch version of linear classification

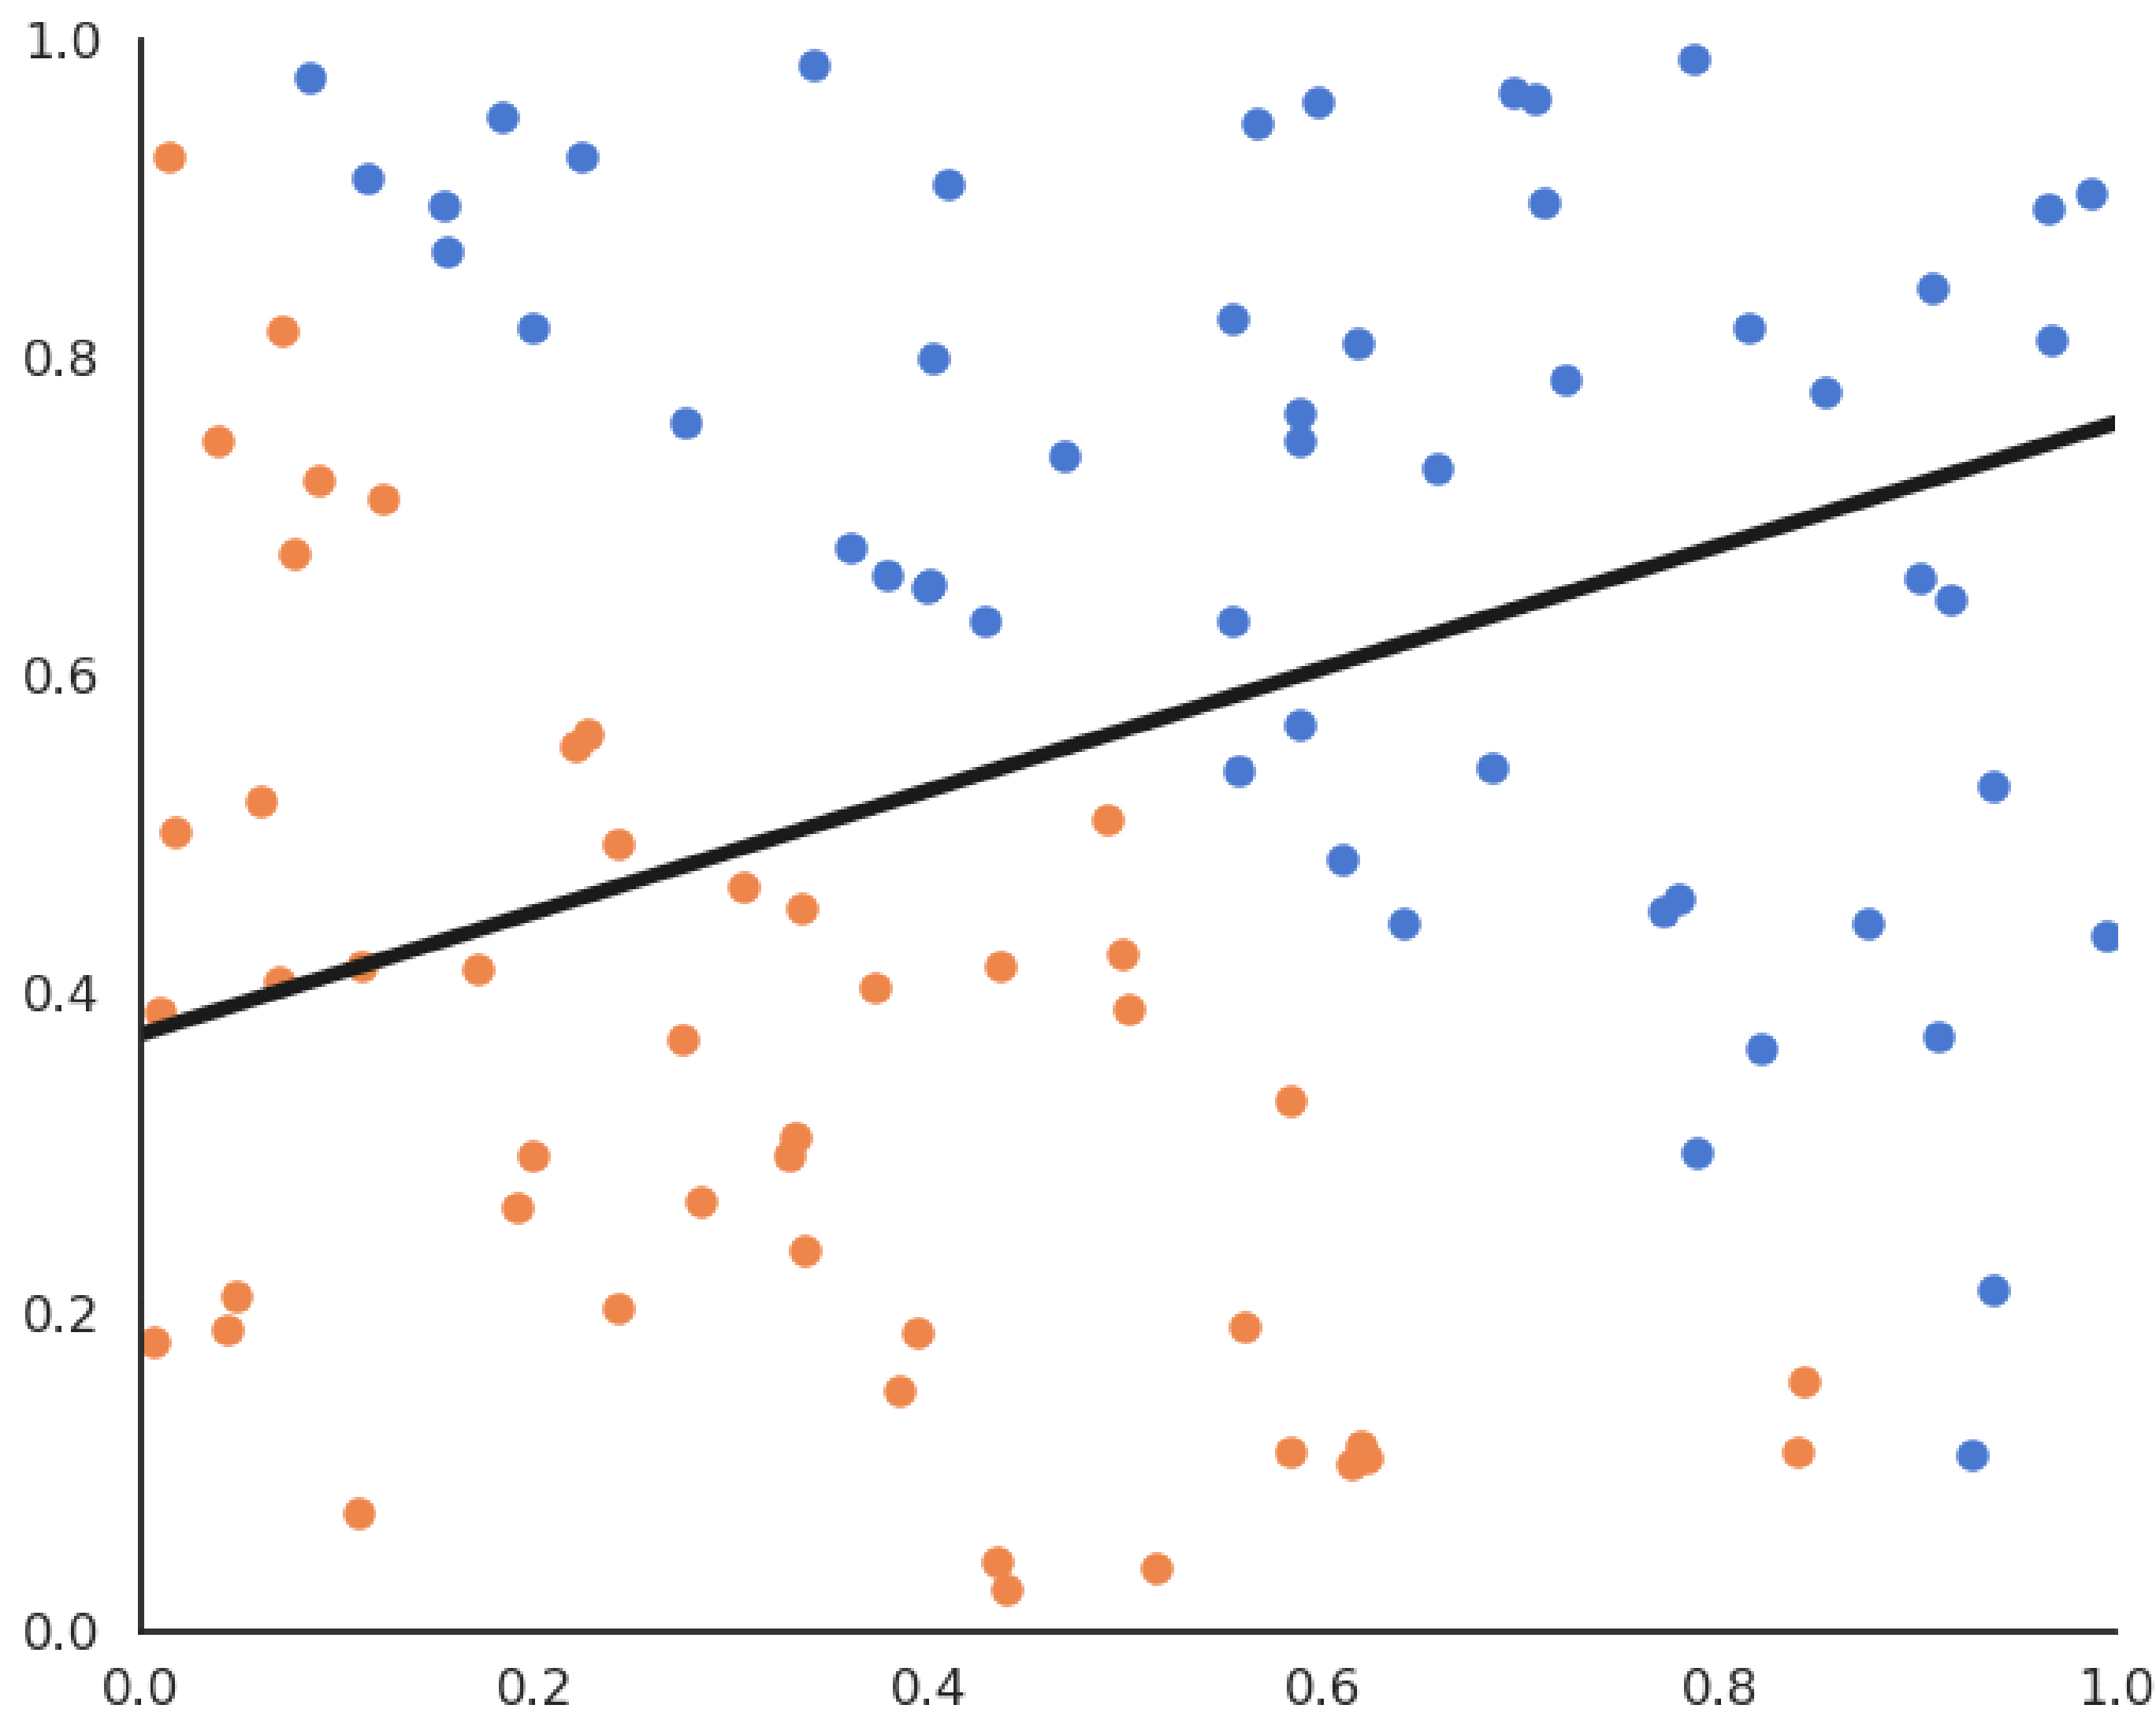
- By applying gradient descent on the quadratic error function, one obtains the following algorithm:



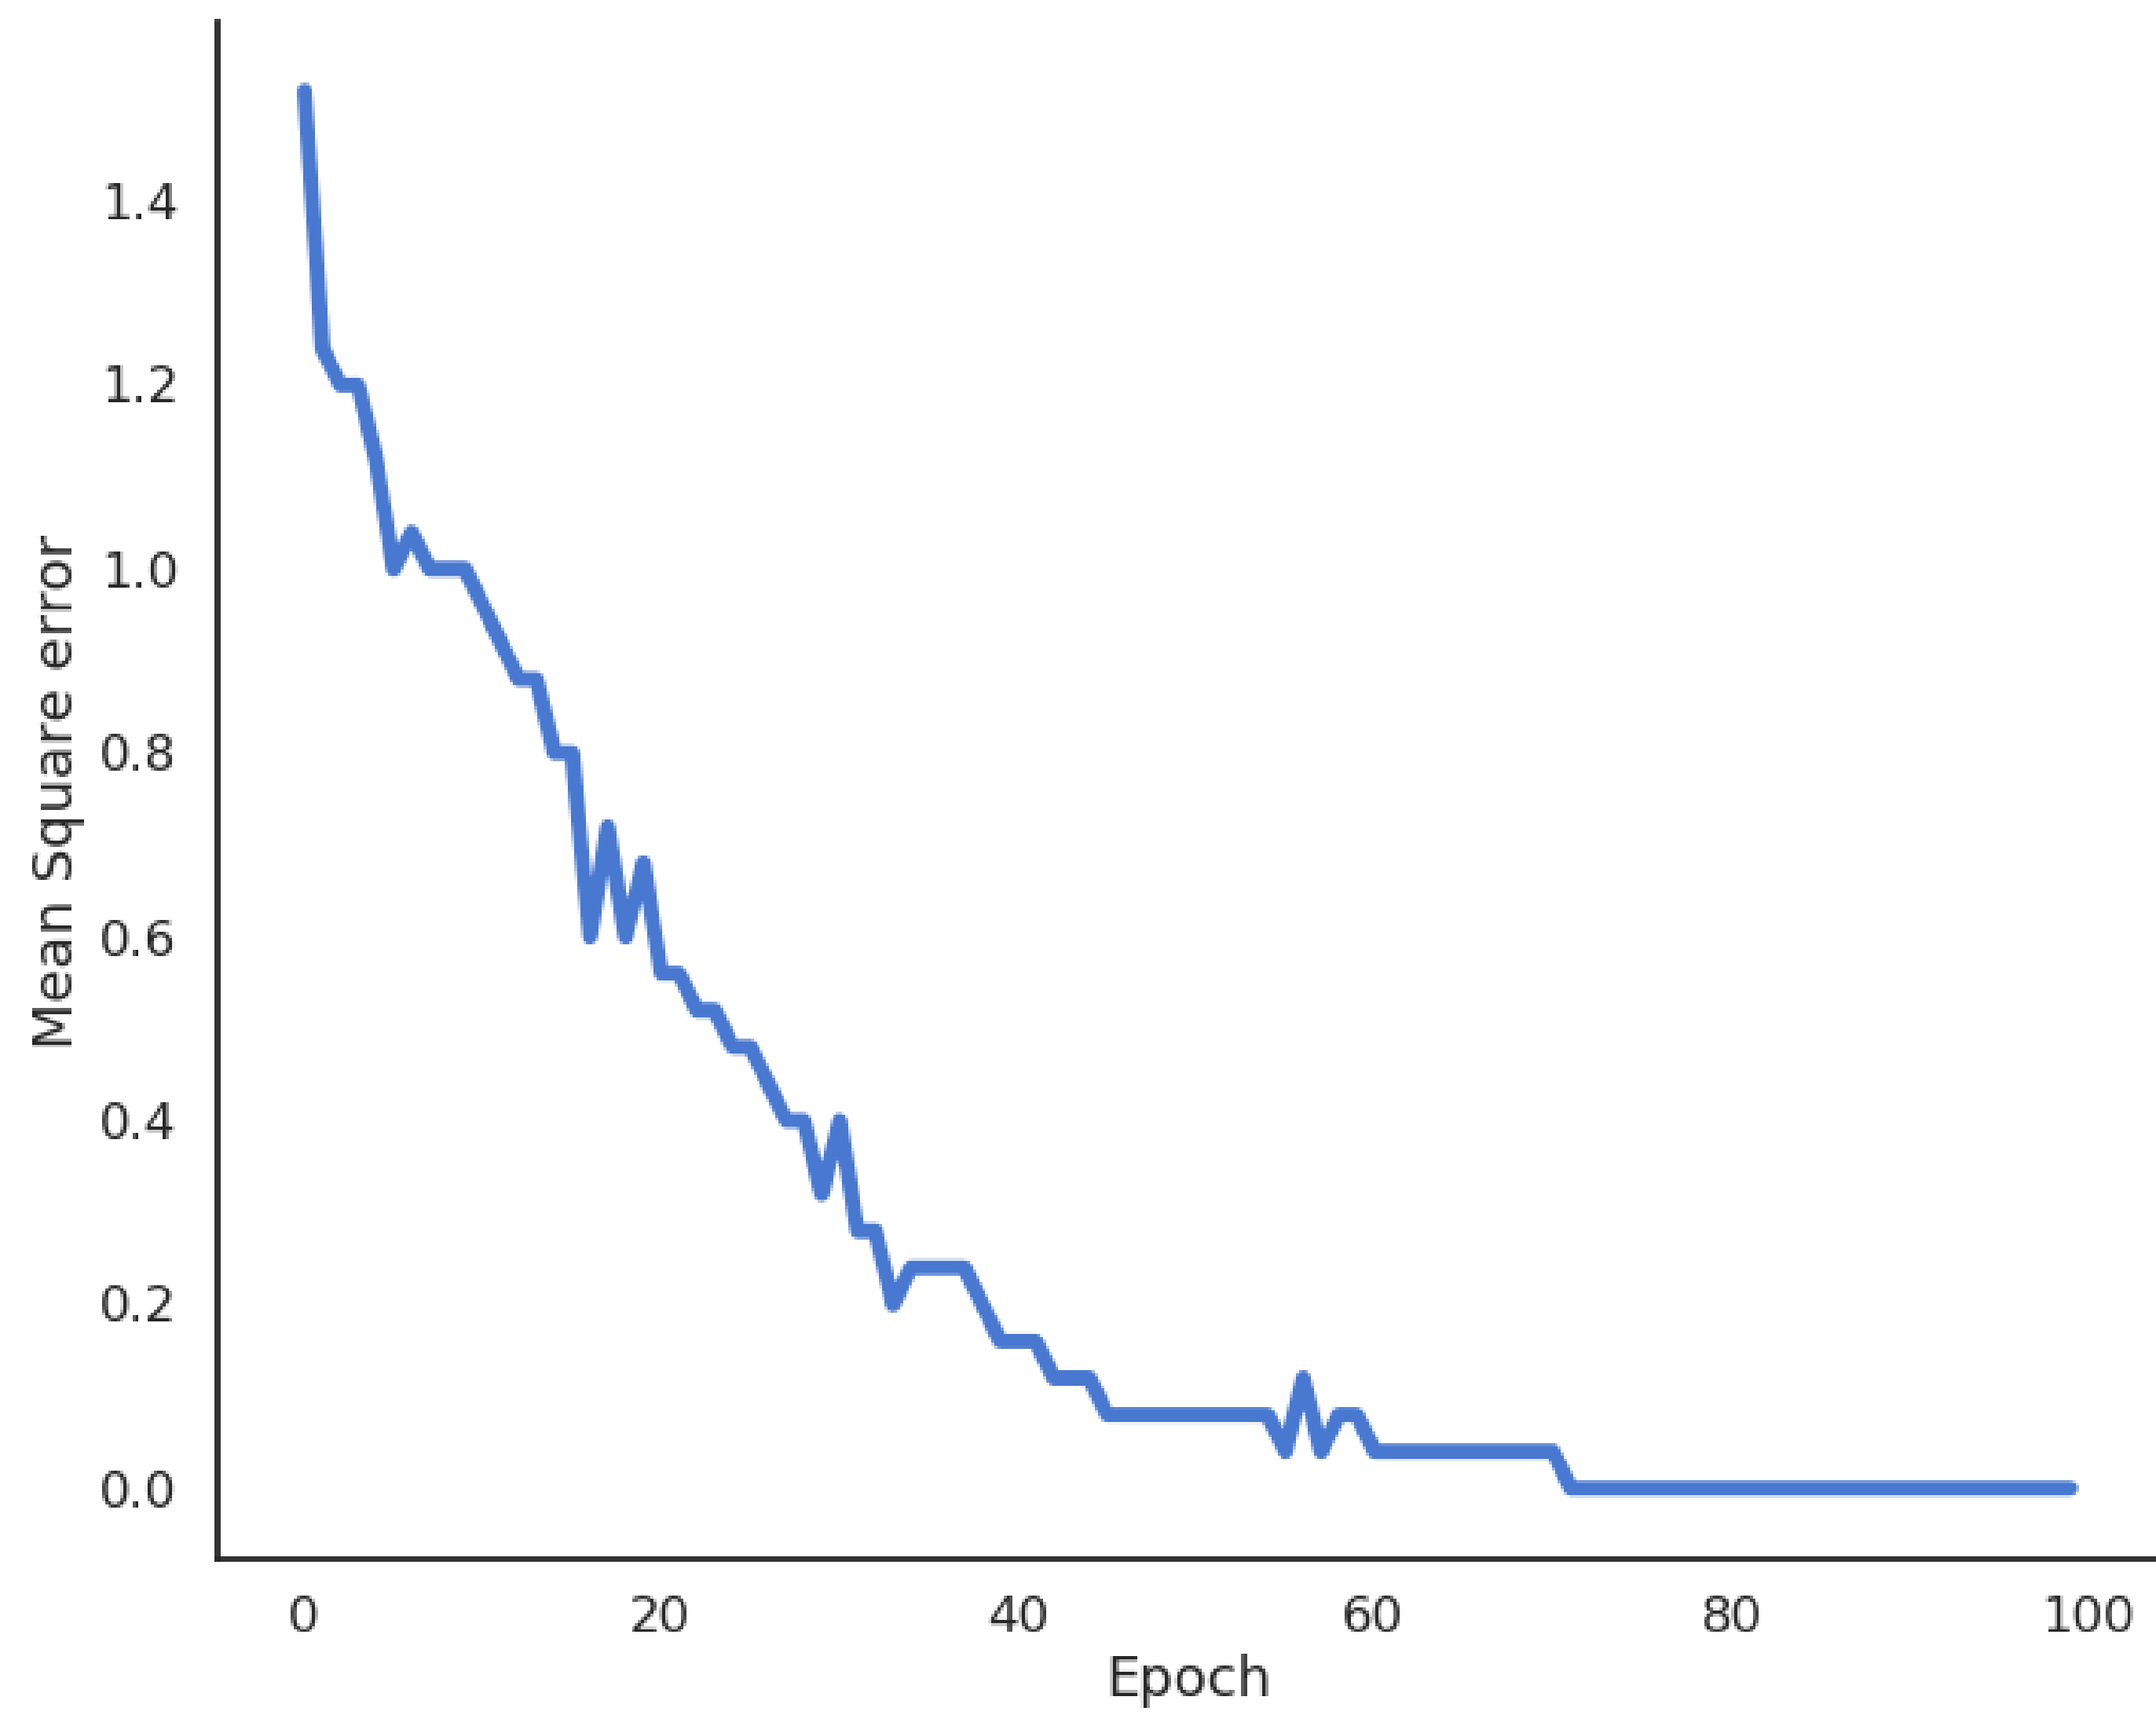
## Batch linear classification

- **for**  $M$  epochs:
  - $\mathbf{dw} = 0$        $db = 0$
  - **for** each sample  $(\mathbf{x}_i, t_i)$ :
    - $y_i = \text{sign}(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b)$
    - $\mathbf{dw} = \mathbf{dw} + (t_i - y_i) \mathbf{x}_i$
    - $db = db + (t_i - y_i)$
  - $\Delta \mathbf{w} = \eta \frac{1}{N} \mathbf{dw}$
  - $\Delta b = \eta \frac{1}{N} db$
- This is called the **batch** version of the Perceptron algorithm.
- If the data is linearly separable and  $\eta$  is well chosen, it converges to the minimum of the mean square error.

# Linear classification: batch version



# Linear classification: batch version





# Online version of linear classification : the Perceptron algorithm

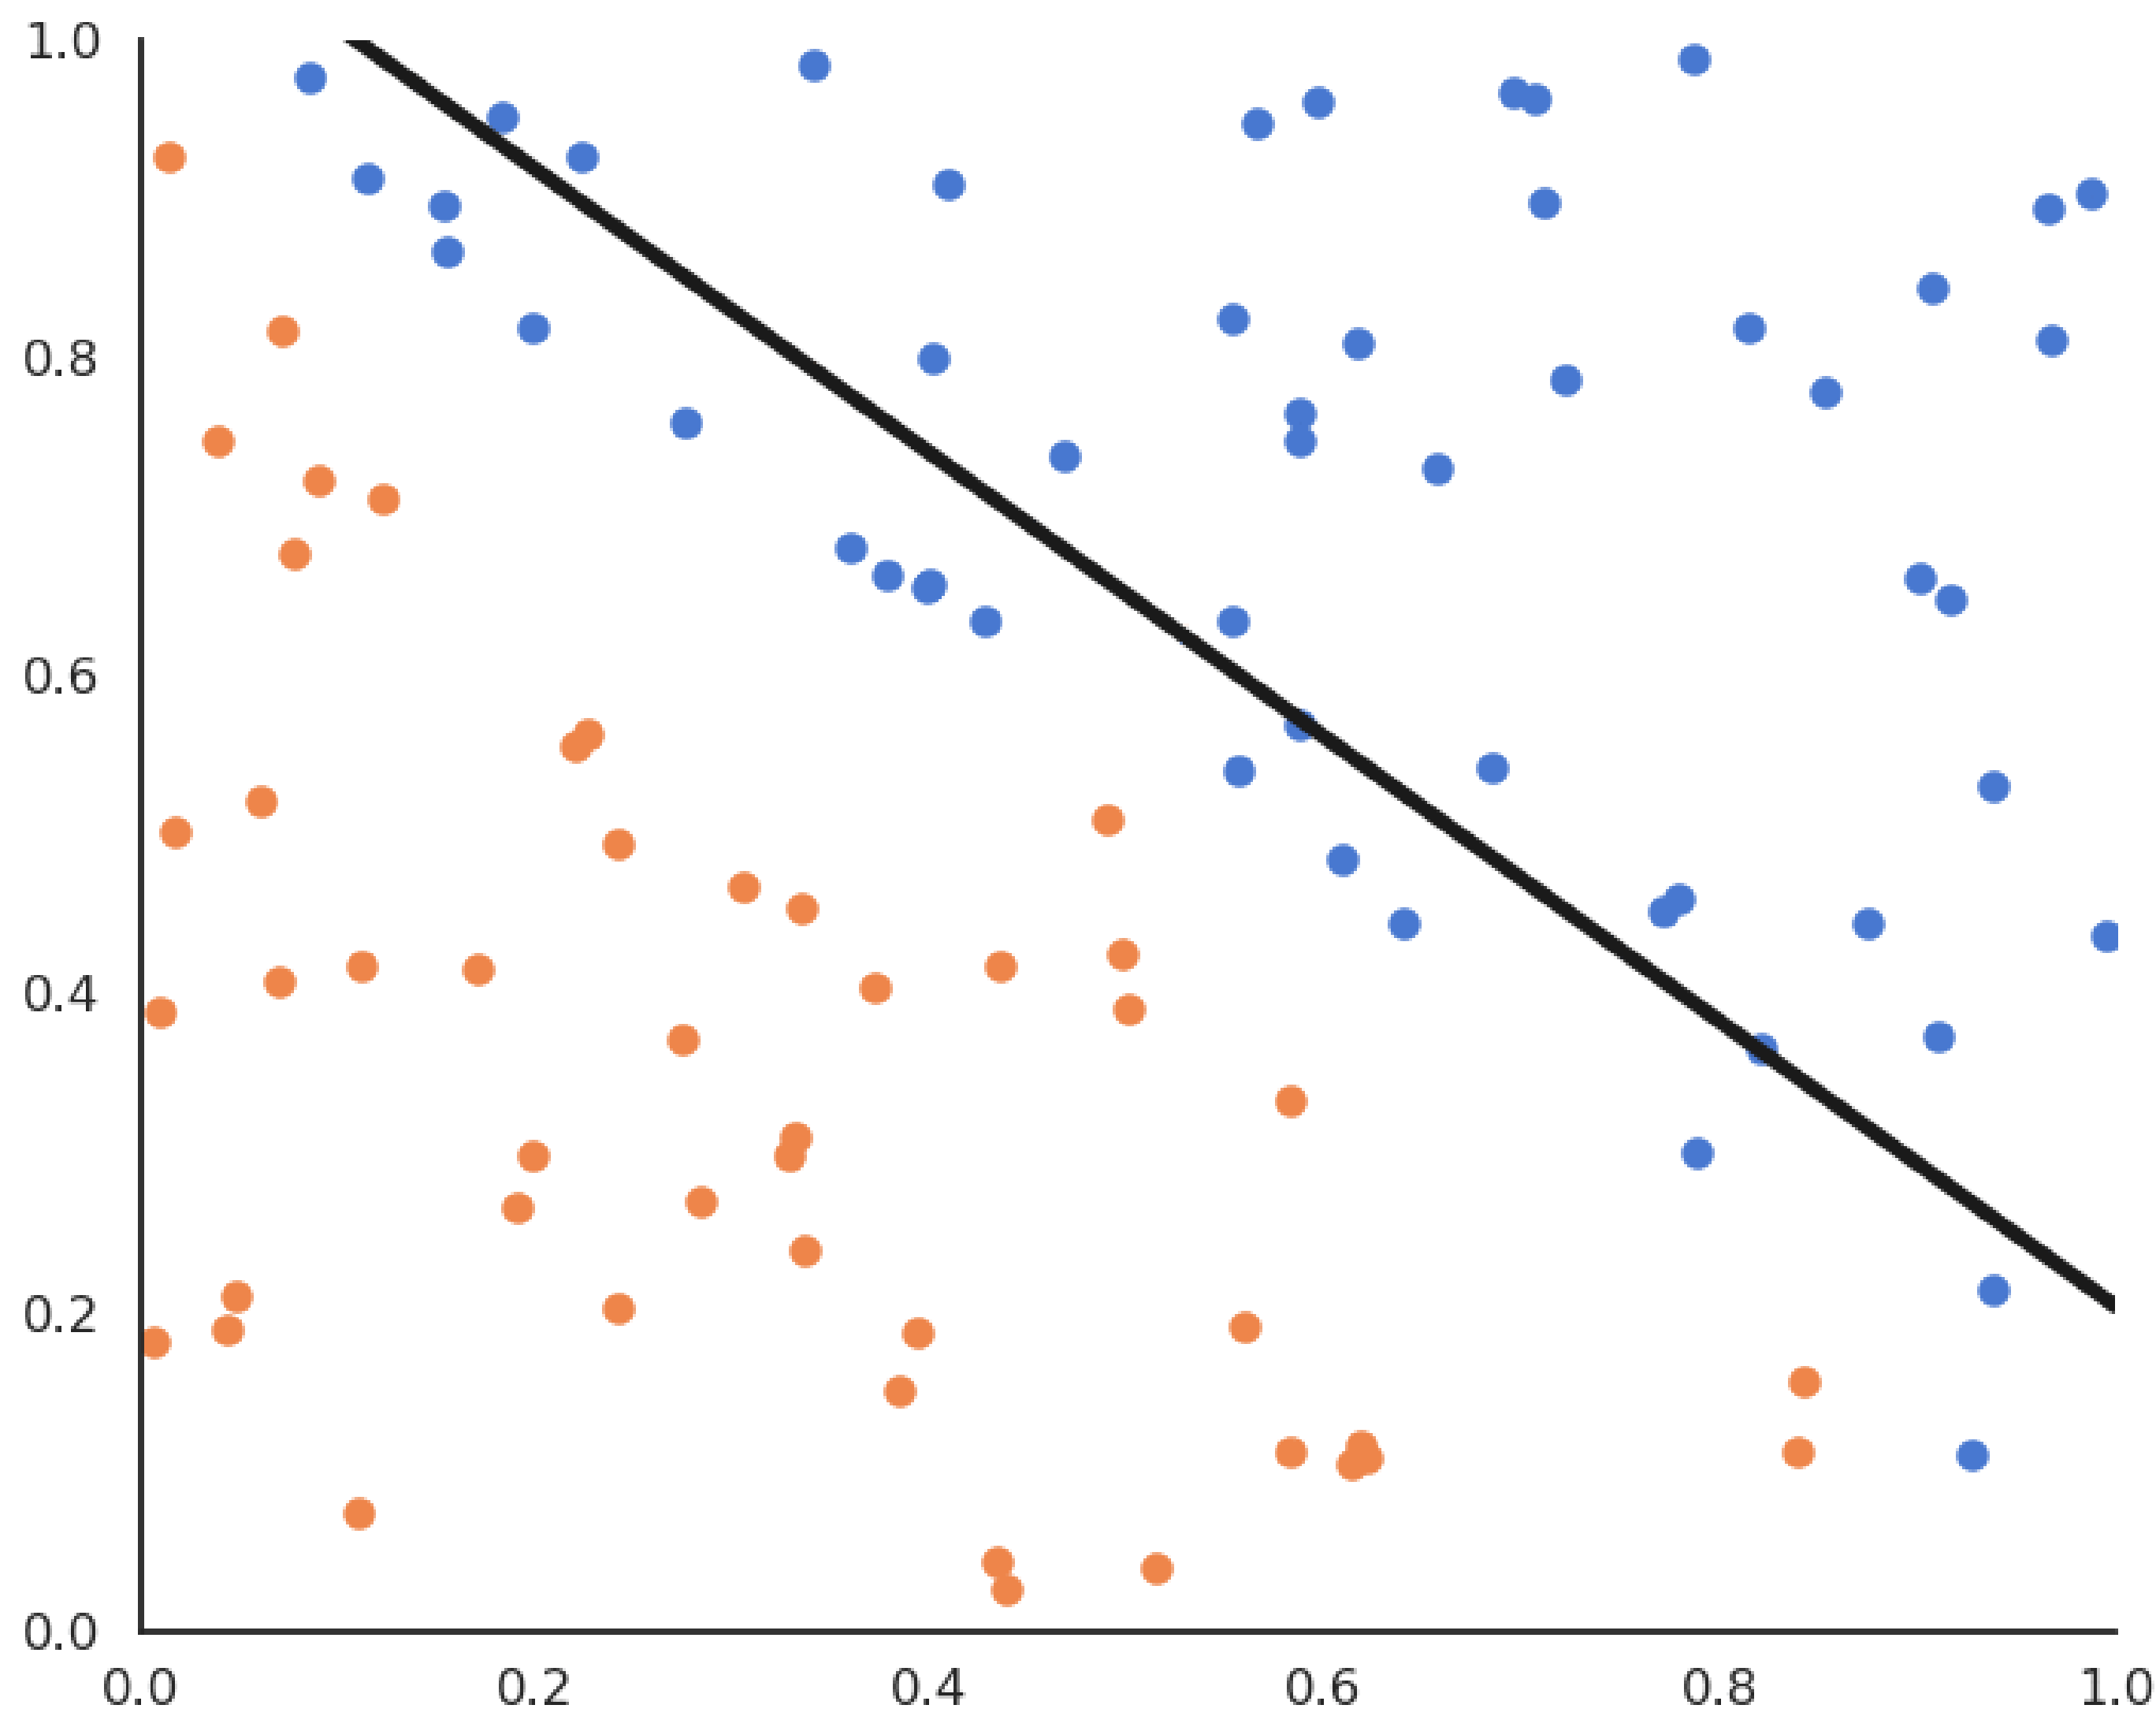
- The **Perceptron algorithm** was invented by the psychologist Frank Rosenblatt in 1958. It was the first algorithmic neural network able to learn linear classification.



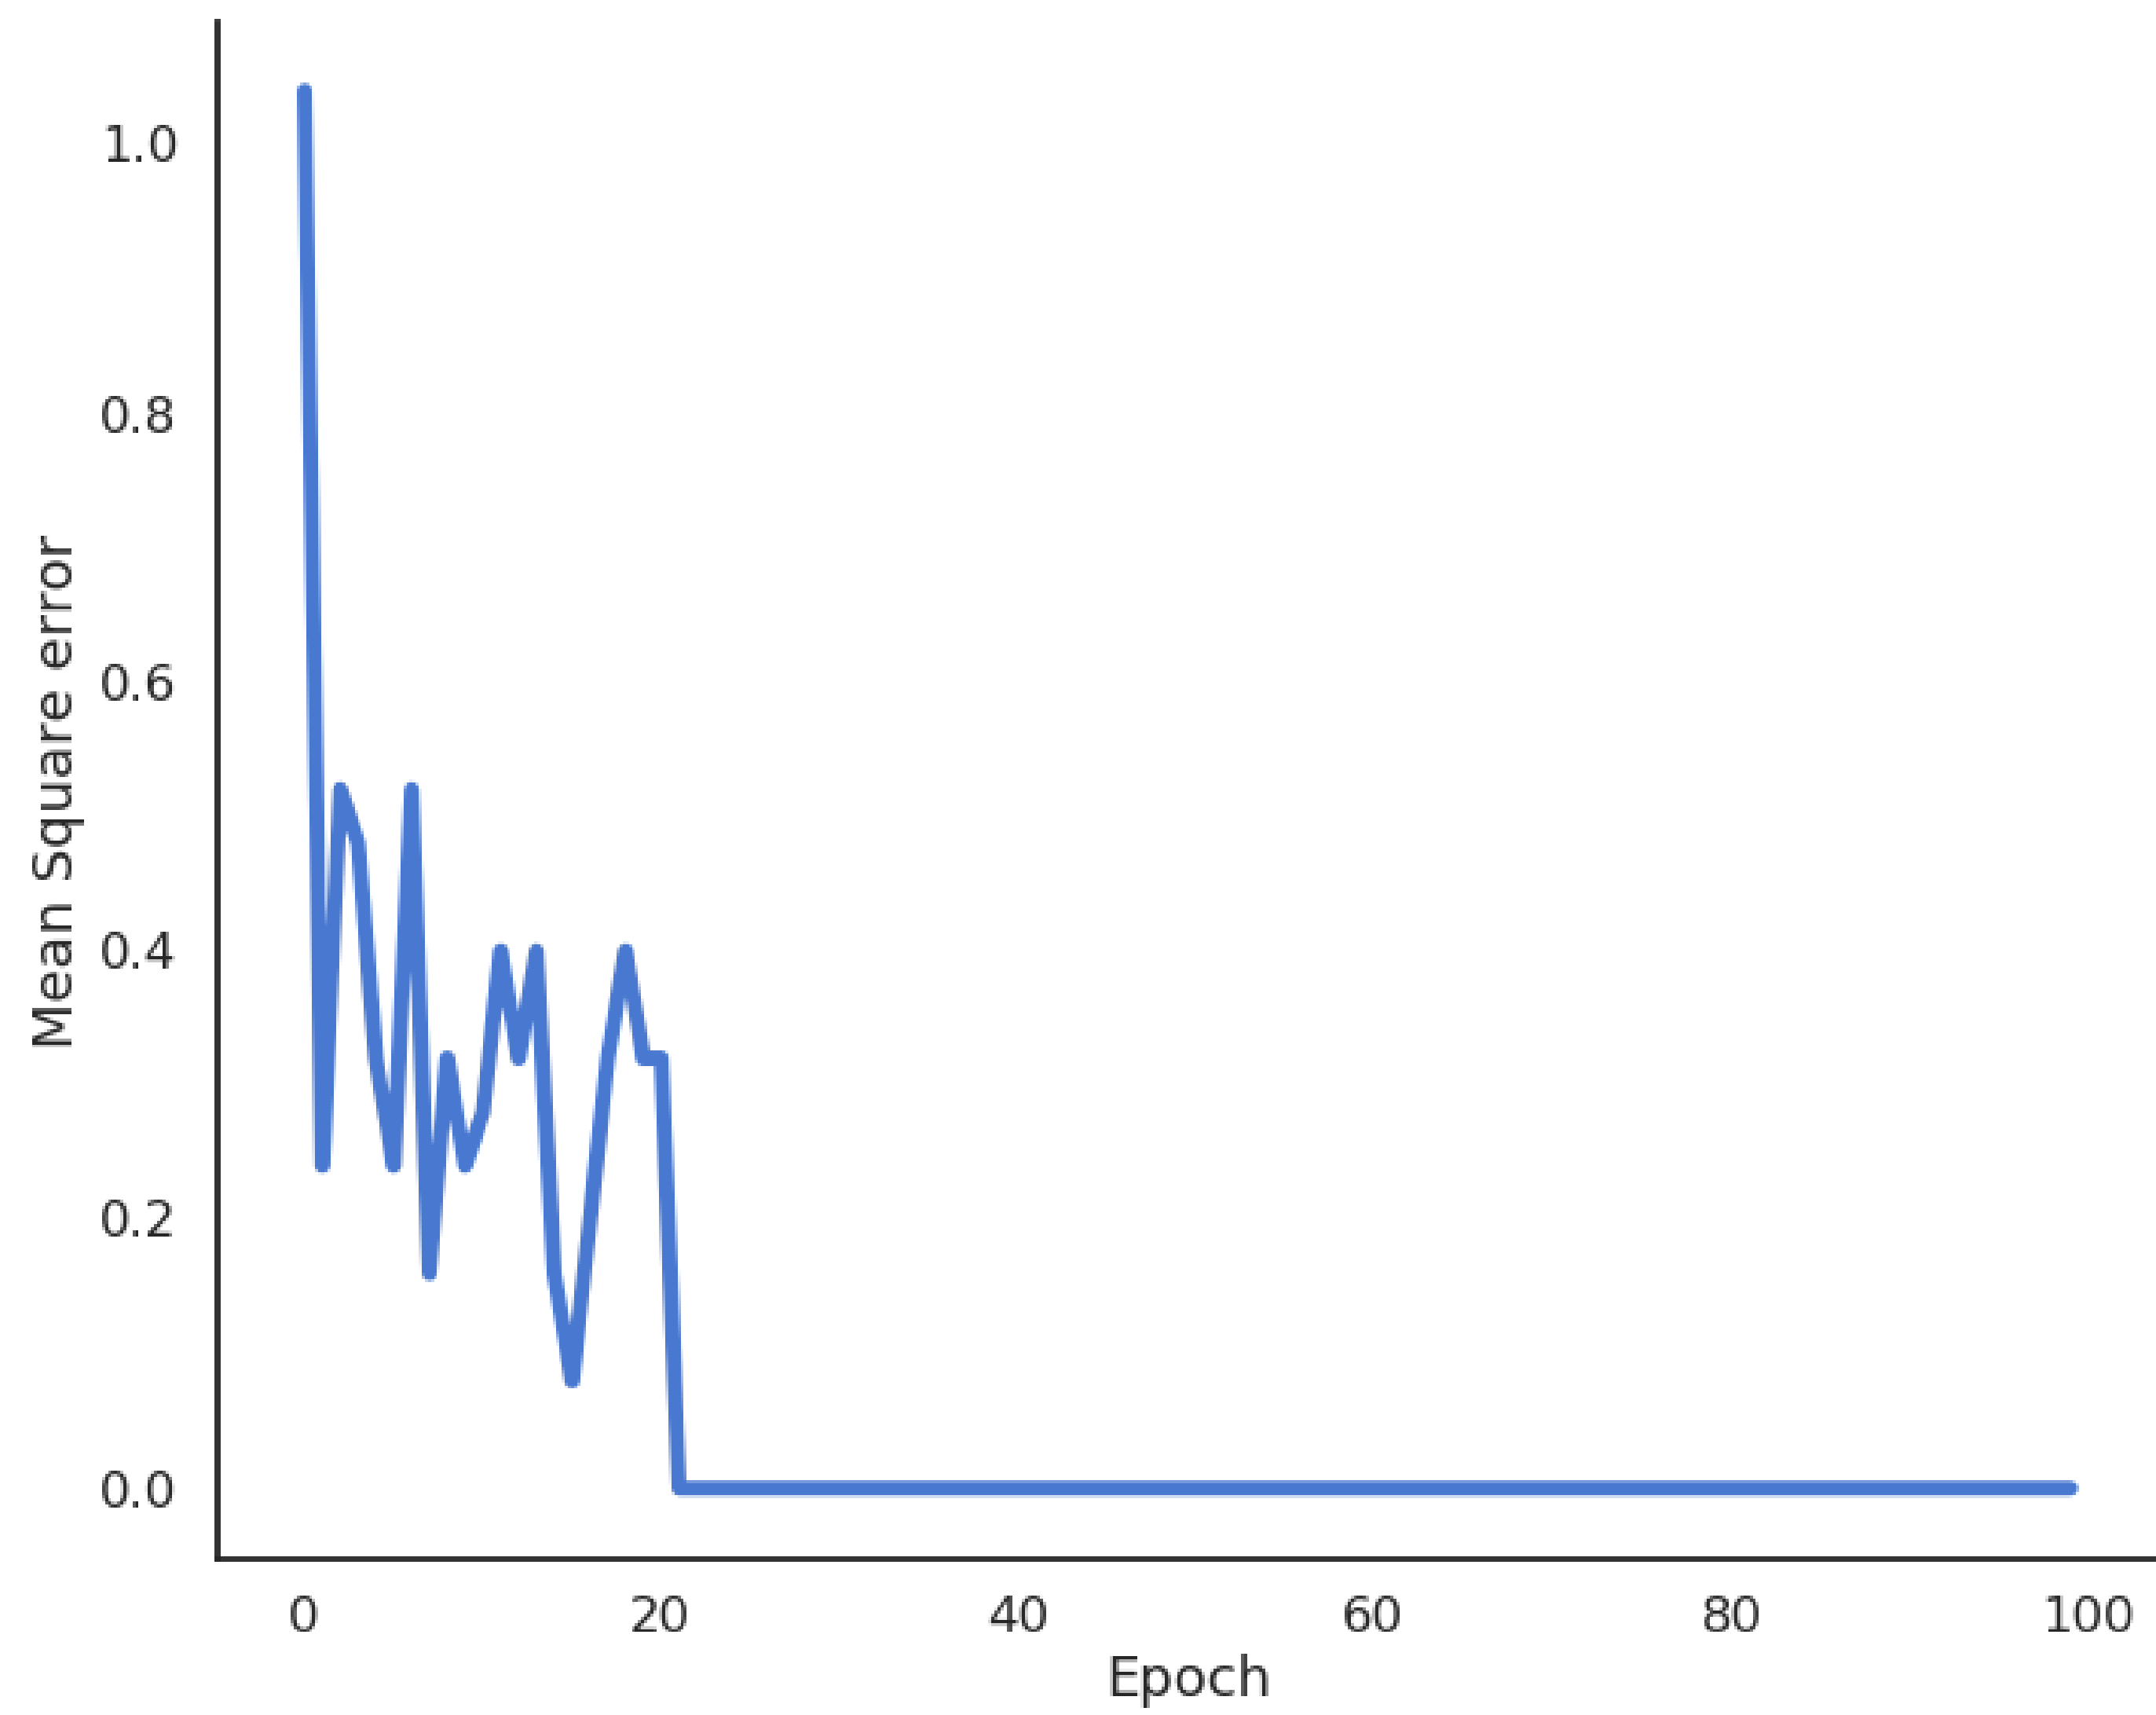
## Perceptron algorithm

- **for**  $M$  epochs:
  - **for** each sample  $(\mathbf{x}_i, t_i)$ :
    - $y_i = \text{sign}(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b)$
    - $\Delta \mathbf{w} = \eta (t_i - y_i) \mathbf{x}_i$
    - $\Delta b = \eta (t_i - y_i)$
- This algorithm iterates over all examples of the training set and applies the **delta learning rule** to each of them immediately, not at the end on the whole training set.
- One could check whether there are still classification errors on the training set at the end of each epoch and stop the algorithm.
- The delta learning rule depends on the learning rate  $\eta$ , the error made by the prediction  $(t_i - y_i)$  and the input  $\mathbf{x}_i$ .

# Linear classification: online version



# Linear classification: online version



# Batch vs. Online learning

- The mean square error is defined as the **expectation** over the data:

$$\mathcal{L}(\mathbf{w}, b) = \mathbb{E}_{\mathcal{D}}[(t_i - y_i)^2]$$

- Batch learning** uses the whole training set as samples to estimate the mse:
- Online learning** uses a single sample to estimate the mse:

$$\mathcal{L}(\mathbf{w}, b) \approx \frac{1}{N} \sum_{i=1}^N (t_i - y_i)^2$$

$$\Delta \mathbf{w} = \eta \frac{1}{N} \sum_{i=1}^N (t_i - y_i) \mathbf{x}_i$$

$$\mathcal{L}(\mathbf{w}, b) \approx (t_i - y_i)^2$$

$$\Delta \mathbf{w} = \eta (t_i - y_i) \mathbf{x}_i$$

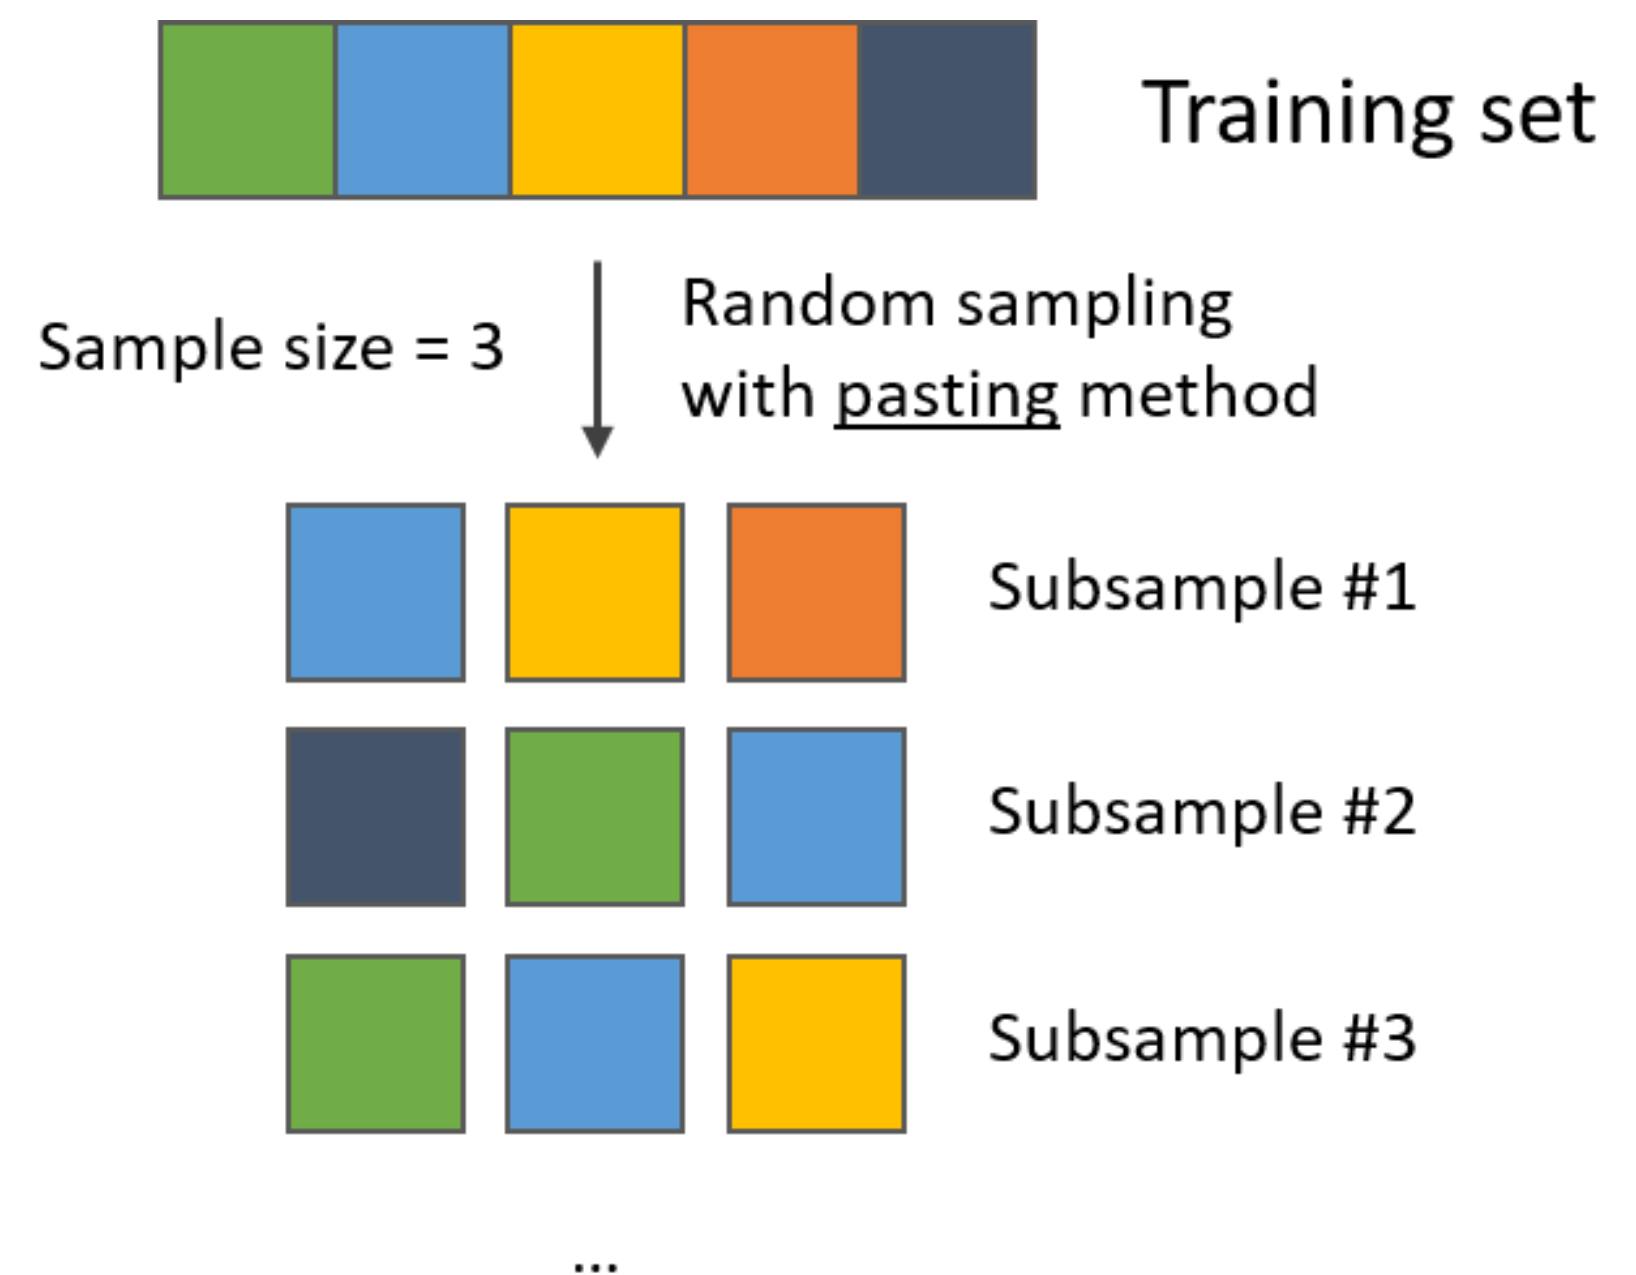
- Batch learning has less bias (central limit theorem) and is less sensible to noise in the data, but is very slow.
- Online learning converges faster, but can be instable and overfits (high variance).

# Stochastic Gradient Descent - SGD

- In practice, we use a trade-off between batch and online learning called **Stochastic Gradient Descent (SGD)** or **Minibatch Gradient Descent**.
- The training set is randomly split at each epoch into small chunks of data (a **minibatch**, usually 32 or 64 examples) and the batch learning rule is applied on each chunk.

$$\Delta \mathbf{w} = \eta \frac{1}{K} \sum_{i=1}^K (t_i - y_i) \mathbf{x}_i$$

- If the **batch size** is well chosen, SGD is as stable as batch learning and as fast as online learning.
- The minibatches are randomly selected at each epoch (i.i.d).
- Online learning is a stochastic gradient descent with a batch size of 1.



## 2 - Maximum Likelihood Estimation

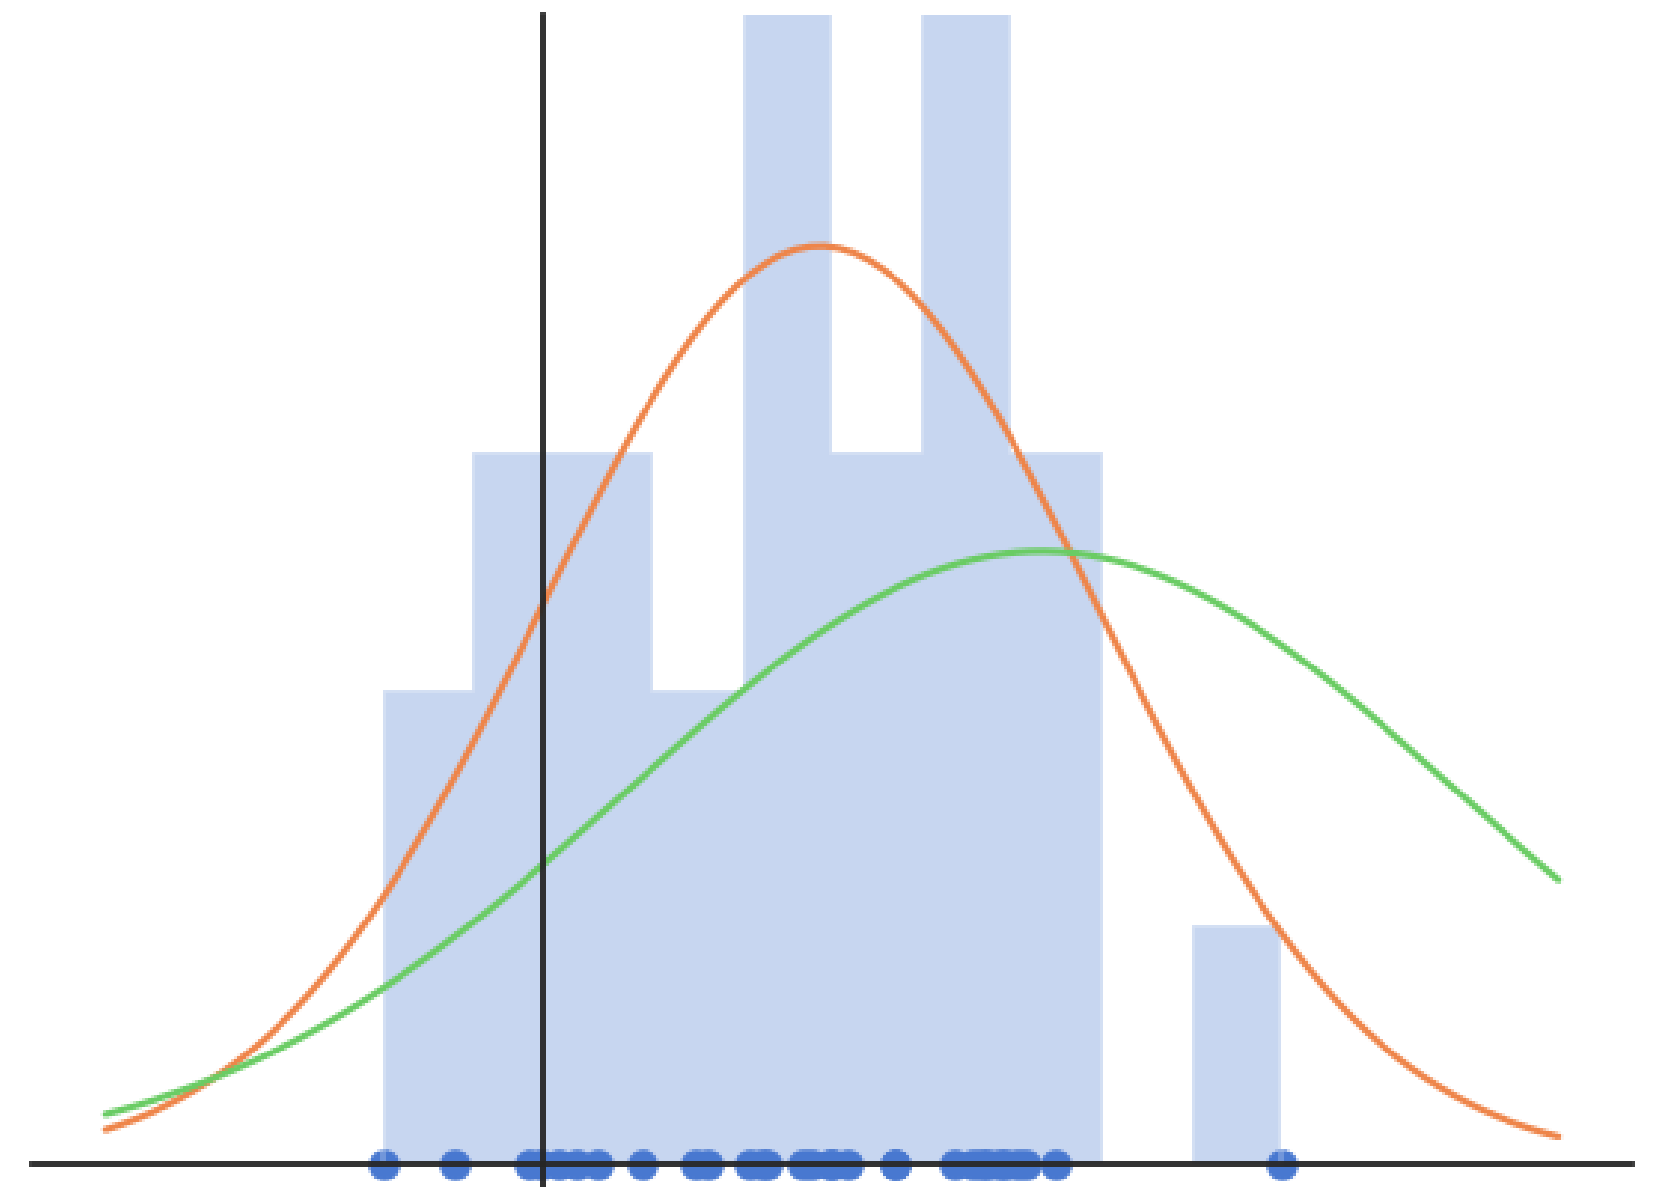


# Maximum Likelihood Estimation

- Let's consider  $N$  **samples**  $\{x_i\}_{i=1}^N$  independently taken from a **normal distribution**  $X$ .
- The probability density function (pdf) of a normal distribution is:

$$f(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp -\frac{(x - \mu)^2}{2\sigma^2}$$

where  $\mu$  is the mean of the distribution and  $\sigma$  its standard deviation.




- The problem is to find the values of  $\mu$  and  $\sigma$  which explain best the observations  $\{x_i\}_{i=1}^N$ .

# Maximum Likelihood Estimation

- The idea of MLE is to maximize the joint density function for all observations. This function is expressed by the **likelihood function**:

$$L(\mu, \sigma) = P(\mathbf{x}; \mu, \sigma) = \prod_{i=1}^N f(x_i; \mu, \sigma)$$

- When the pdf takes high values for all samples, it is quite likely that the samples come from this particular distribution.



- The likelihood function reflects how well the parameters  $\mu$  and  $\sigma$  explain the observations  $\{x_i\}_{i=1}^N$ .
- Note: the samples must be i.i.d. so that the likelihood is a product.

# Maximum Likelihood Estimation

- We therefore search for the values  $\mu$  and  $\sigma$  which **maximize** the likelihood function.

$$\max_{\mu, \sigma} L(\mu, \sigma) = \prod_{i=1}^N f(x_i; \mu, \sigma)$$

- For the normal distribution, the likelihood function is:

$$\begin{aligned} L(\mu, \sigma) &= \prod_{i=1}^N f(x_i; \mu, \sigma) \\ &= \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp -\frac{(x_i - \mu)^2}{2\sigma^2} \\ &= \left(\frac{1}{\sqrt{2\pi\sigma^2}}\right)^N \prod_{i=1}^N \exp -\frac{(x_i - \mu)^2}{2\sigma^2} \\ &= \left(\frac{1}{\sqrt{2\pi\sigma^2}}\right)^N \exp -\frac{\sum_{i=1}^N (x_i - \mu)^2}{2\sigma^2} \end{aligned}$$

# Maximum Likelihood Estimation

- To find the maximum of  $L(\mu, \sigma)$ , we need to search where the gradient is equal to zero:

$$\begin{cases} \frac{\partial L(\mu, \sigma)}{\partial \mu} = 0 \\ \frac{\partial L(\mu, \sigma)}{\partial \sigma} = 0 \end{cases}$$

- The likelihood function is complex to differentiate, so we consider its logarithm  $l(\mu, \sigma) = \log(L(\mu, \sigma))$  which has a maximum for the same value of  $(\mu, \sigma)$  as the log function is monotonic.

$$\begin{aligned} l(\mu, \sigma) &= \log(L(\mu, \sigma)) \\ &= \log \left( \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right)^N \exp - \frac{\sum_{i=1}^N (x_i - \mu)^2}{2\sigma^2} \right) \\ &= -\frac{N}{2} \log(2\pi\sigma^2) - \frac{\sum_{i=1}^N (x_i - \mu)^2}{2\sigma^2} \end{aligned}$$

- $l(\mu, \sigma)$  is called the **log-likelihood** function.

# Maximum Likelihood Estimation

$$l(\mu, \sigma) = -\frac{N}{2} \log(2\pi\sigma^2) - \frac{\sum_{i=1}^N (x_i - \mu)^2}{2\sigma^2}$$

- The maximum of the log-likelihood function respects:

$$\begin{aligned}\frac{\partial l(\mu, \sigma)}{\partial \mu} &= \frac{\sum_{i=1}^N (x_i - \mu)}{\sigma^2} = 0 \\ \frac{\partial l(\mu, \sigma)}{\partial \sigma} &= -\frac{N}{2} \frac{4\pi\sigma}{2\pi\sigma^2} + \frac{\sum_{i=1}^N (x_i - \mu)^2}{\sigma^3} \\ &= -\frac{N}{\sigma} + \frac{\sum_{i=1}^N (x_i - \mu)^2}{\sigma^3} = 0\end{aligned}$$

- We obtain:

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i \qquad \sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2$$

# Maximum Likelihood Estimation

- Unsurprisingly, the mean and variance of the normal distribution which best explains the data are the mean and variance of the data...

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i \qquad \sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2$$

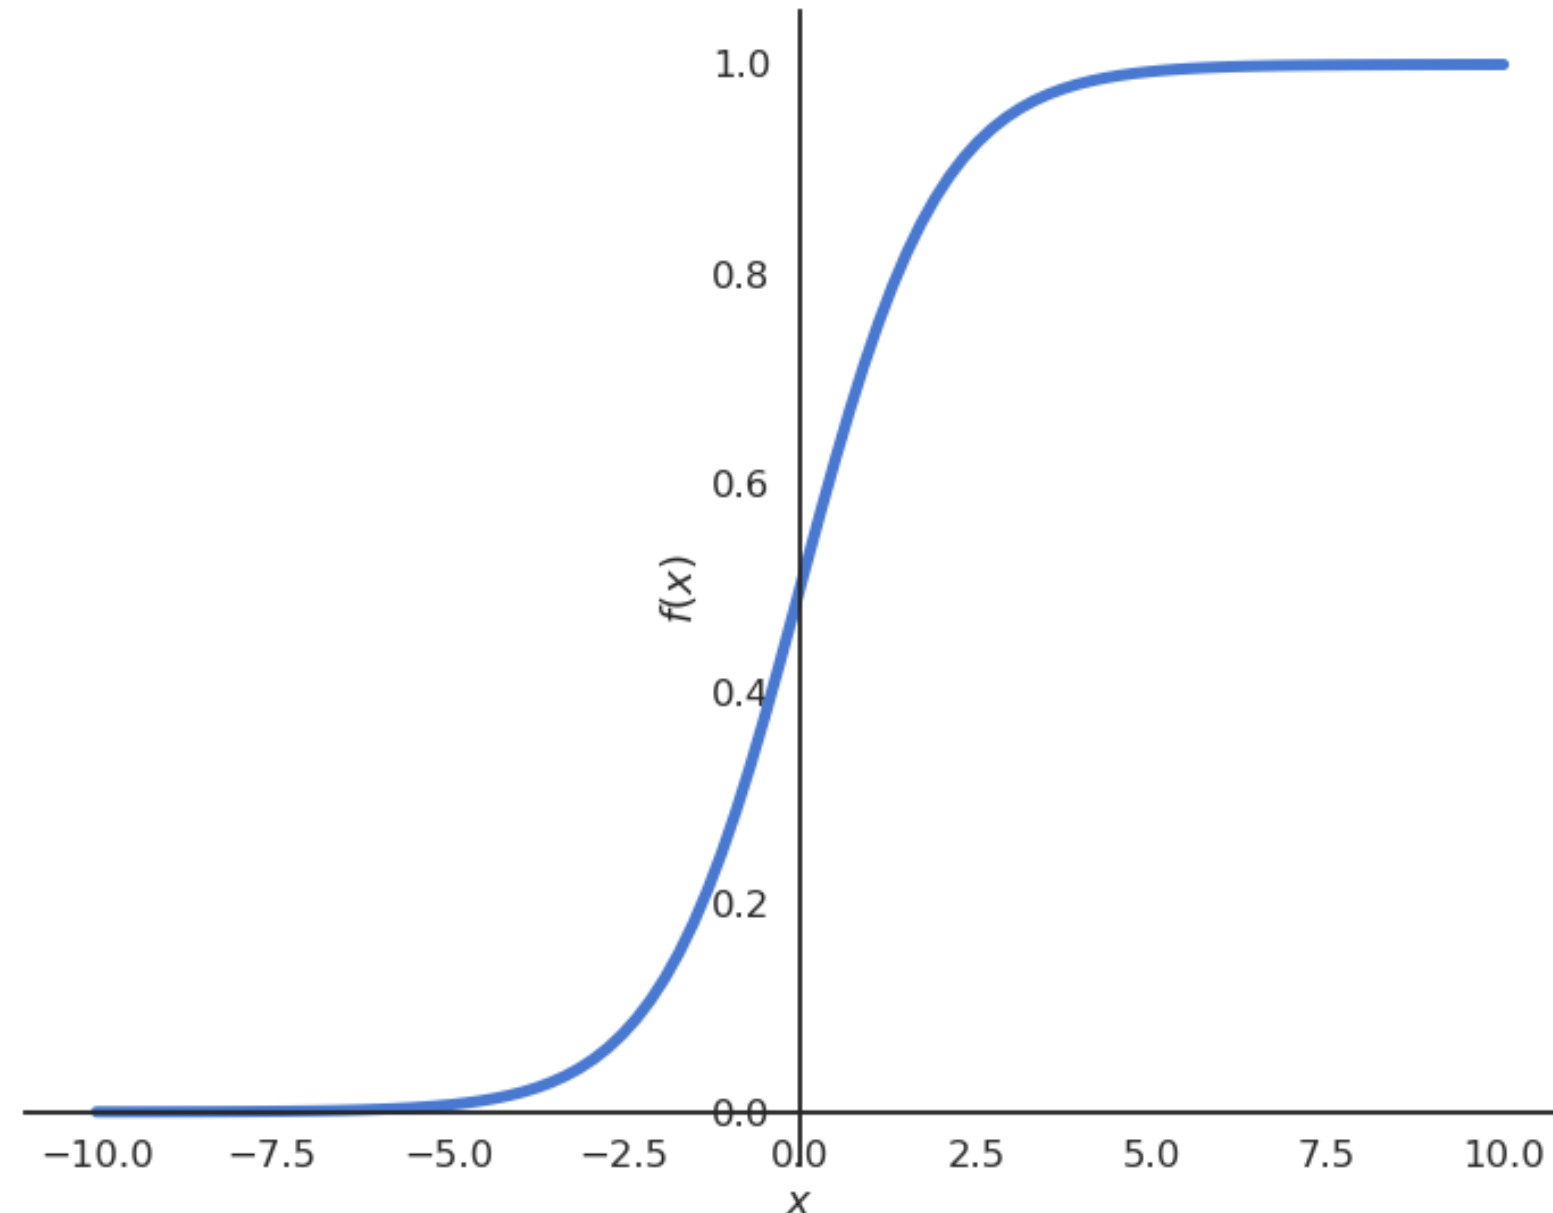
- The same principle can be applied to estimate the parameters of any distribution: normal, exponential, Bernoulli, Poisson, etc...
- When a machine learning method has a probabilistic interpretation (i.e. it outputs probabilities), MLE can be used to find its parameters.
- One can use global optimization like here, or gradient descent to estimate the parameters iteratively.



## **3 - Soft linear classification : Logistic regression**

## Reminder: Logistic regression

- We want to perform a regression, but where the targets  $t_i$  are bounded between 0 and 1.



- We can use a logistic function instead of a linear function in order to transform the net activation into an output:

$$y = \sigma(wx + b) = \frac{1}{1 + \exp(-wx - b)}$$

## Use of logistic regression for soft classification

- Logistic regression can be used in binary classification if we consider  $y = \sigma(w x + b)$  as the probability that the example belongs to the positive class ( $t = 1$ ).

$$P(t = 1|x; w, b) = y; \quad P(t = 0|x; w, b) = 1 - y$$

- The output  $t$  therefore comes from a Bernoulli distribution  $\mathcal{B}$  of parameter  $p = y = f_{w,b}(x)$ . The probability mass function (pmf) is:

$$f(t|x; w, b) = y^t (1 - y)^{1-t}$$

- If we consider our training samples  $(x_i, t_i)$  as independently taken from this distribution, our task is:
  - to find the parameterized distribution that best explains the data, which means:
  - to find the parameters  $w$  and  $b$  maximizing the **likelihood** that the samples  $t$  come from a Bernoulli distribution when  $x$ ,  $w$  and  $b$  are given.
- We only need to apply **Maximum Likelihood Estimation** (MLE) on this Bernoulli distribution!

# MLE for logistic regression

- The likelihood function for logistic regression is :

$$\begin{aligned} L(w, b) &= P(t|x; w, b) = \prod_{i=1}^N f(t_i|x_i; w, b) \\ &= \prod_{i=1}^N y_i^{t_i} (1 - y_i)^{1-t_i} \end{aligned}$$

- The likelihood function is quite hard to differentiate, so we take the **log-likelihood** function:

$$\begin{aligned} l(w, b) &= \log L(w, b) \\ &= \sum_{i=1}^N [t_i \log y_i + (1 - t_i) \log(1 - y_i)] \end{aligned}$$

- or even better: the **negative log-likelihood** which will be minimized using gradient descent:

$$\mathcal{L}(w, b) = - \sum_{i=1}^N [t_i \log y_i + (1 - t_i) \log(1 - y_i)]$$

## MLE for logistic regression

- We then search for the minimum of the negative log-likelihood function by computing its gradient (here for a single sample):

$$\begin{aligned}\frac{\partial l_i(w, b)}{\partial w} &= -\frac{\partial}{\partial w} [t_i \log y_i + (1 - t_i) \log(1 - y_i)] \\ &= -t_i \frac{\partial}{\partial w} \log y_i - (1 - t_i) \frac{\partial}{\partial w} \log(1 - y_i) \\ &= -t_i \frac{\frac{\partial}{\partial w} y_i}{y_i} - (1 - t_i) \frac{\frac{\partial}{\partial w} (1 - y_i)}{1 - y_i} \\ &= -t_i \frac{y_i (1 - y_i) x_i}{y_i} + (1 - t_i) \frac{y_i (1 - y_i) x_i}{1 - y_i} \\ &= -(t_i - y_i) x_i\end{aligned}$$

- Same gradient as the linear perceptron, but with a non-linear output function!

# Logistic regression for soft classification

- Logistic regression is a regression method used for classification. It uses a non-linear transfer function  $\sigma(x) = \frac{1}{1+\exp(-x)}$  applied on the net activation:

$$y_i = \sigma(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b)$$

- The continuous output  $y$  is interpreted as the probability of belonging to the positive class.

$$P(t_i = 1 | \mathbf{x}_i; \mathbf{w}, b) = y_i; \quad P(t_i = 0 | \mathbf{x}_i; \mathbf{w}, b) = 1 - y_i$$

- We minimize the **negative log-likelihood** loss function using gradient descent:

$$\mathcal{L}(\mathbf{w}, b) = - \sum_{i=1}^N [t_i \log y_i + (1 - t_i) \log(1 - y_i)]$$

- We obtain the delta learning rule, using the class as a target and the probability as a prediction:

$$\begin{cases} \Delta \mathbf{w} = \eta (t_i - y_i) \mathbf{x}_i \\ \Delta b = \eta (t_i - y_i) \end{cases}$$

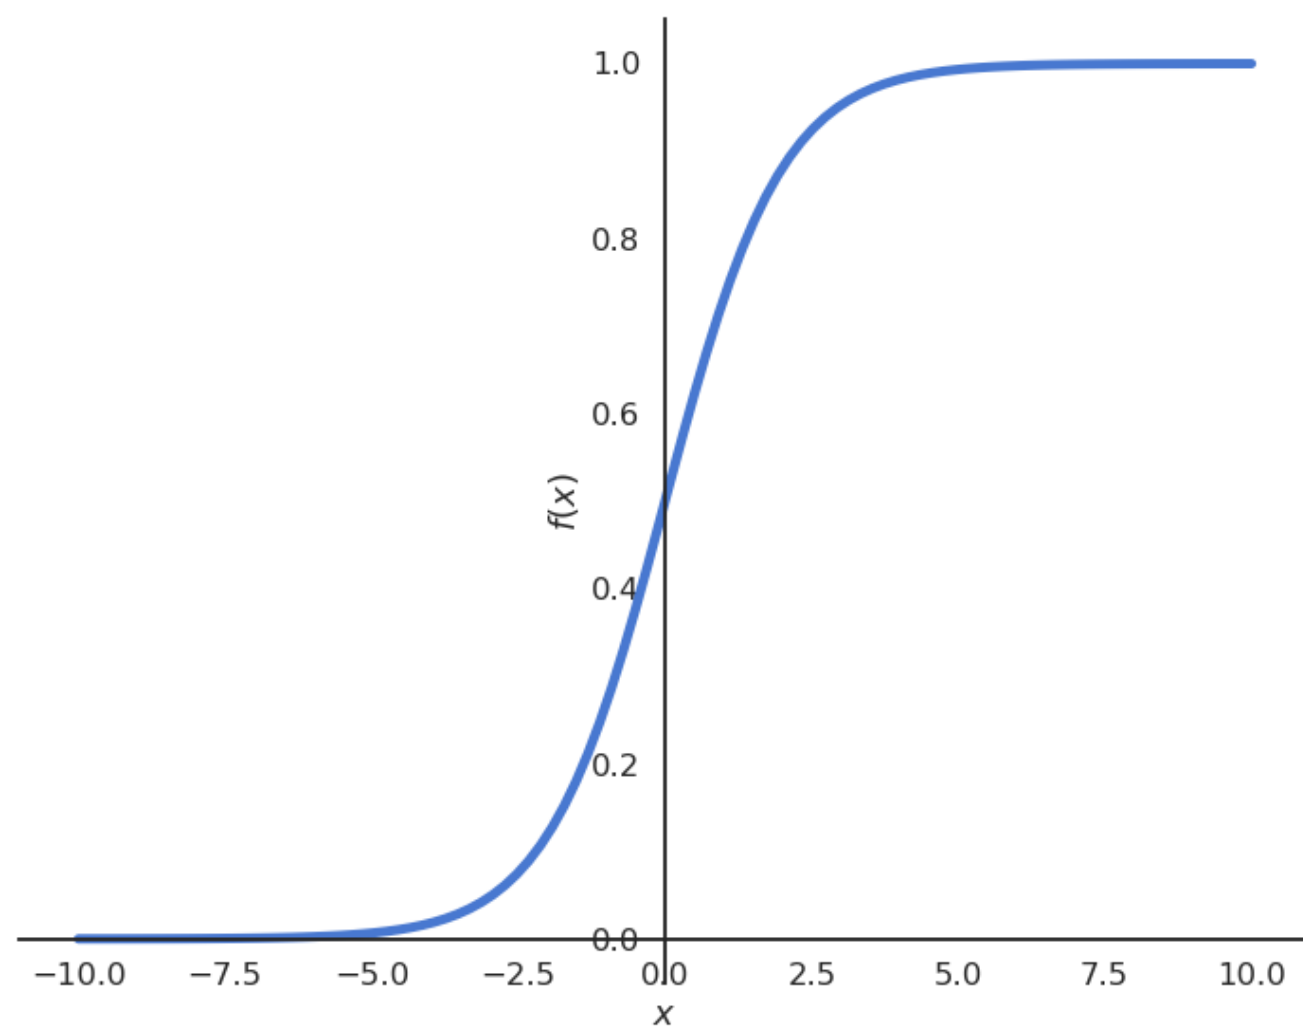


# Logistic regression



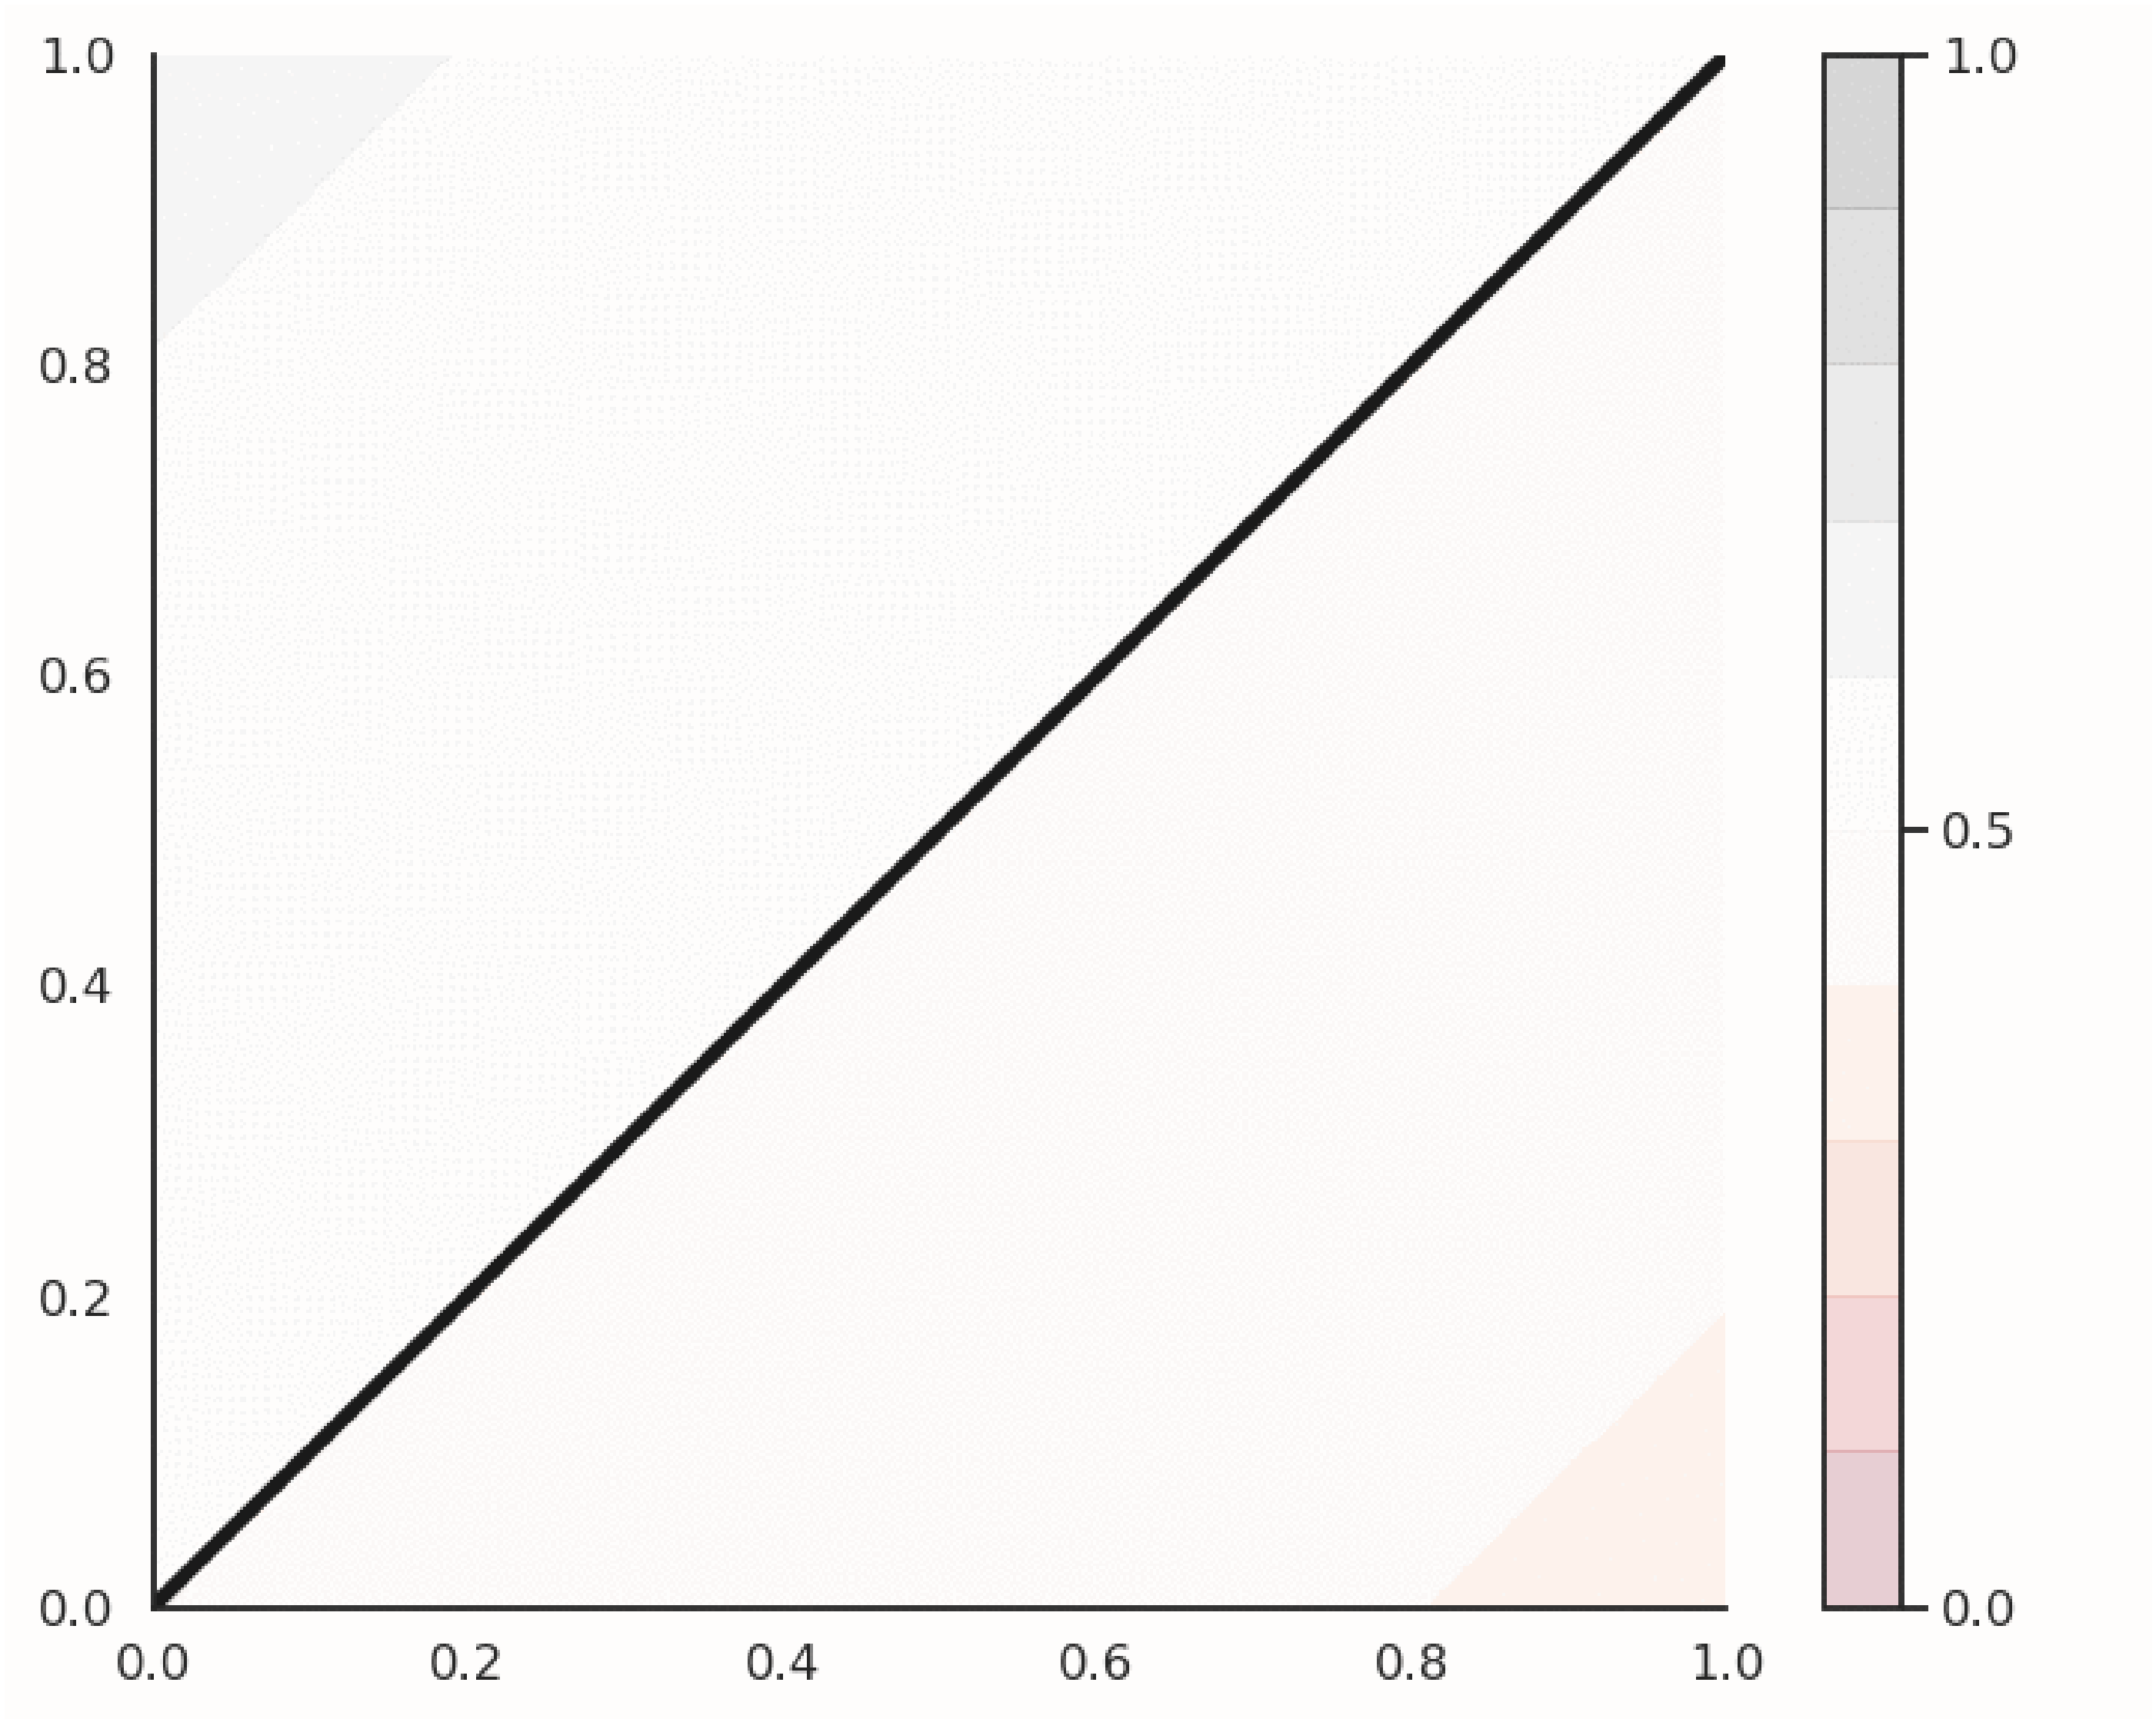
## Logistic regression

- $\mathbf{w} = 0$        $b = 0$
- for  $M$  epochs:
  - for each sample  $(\mathbf{x}_i, t_i)$ :
    - $y_i = \sigma(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b)$
    - $\Delta \mathbf{w} = \eta (t_i - y_i) \mathbf{x}_i$
    - $\Delta b = \eta (t_i - y_i)$

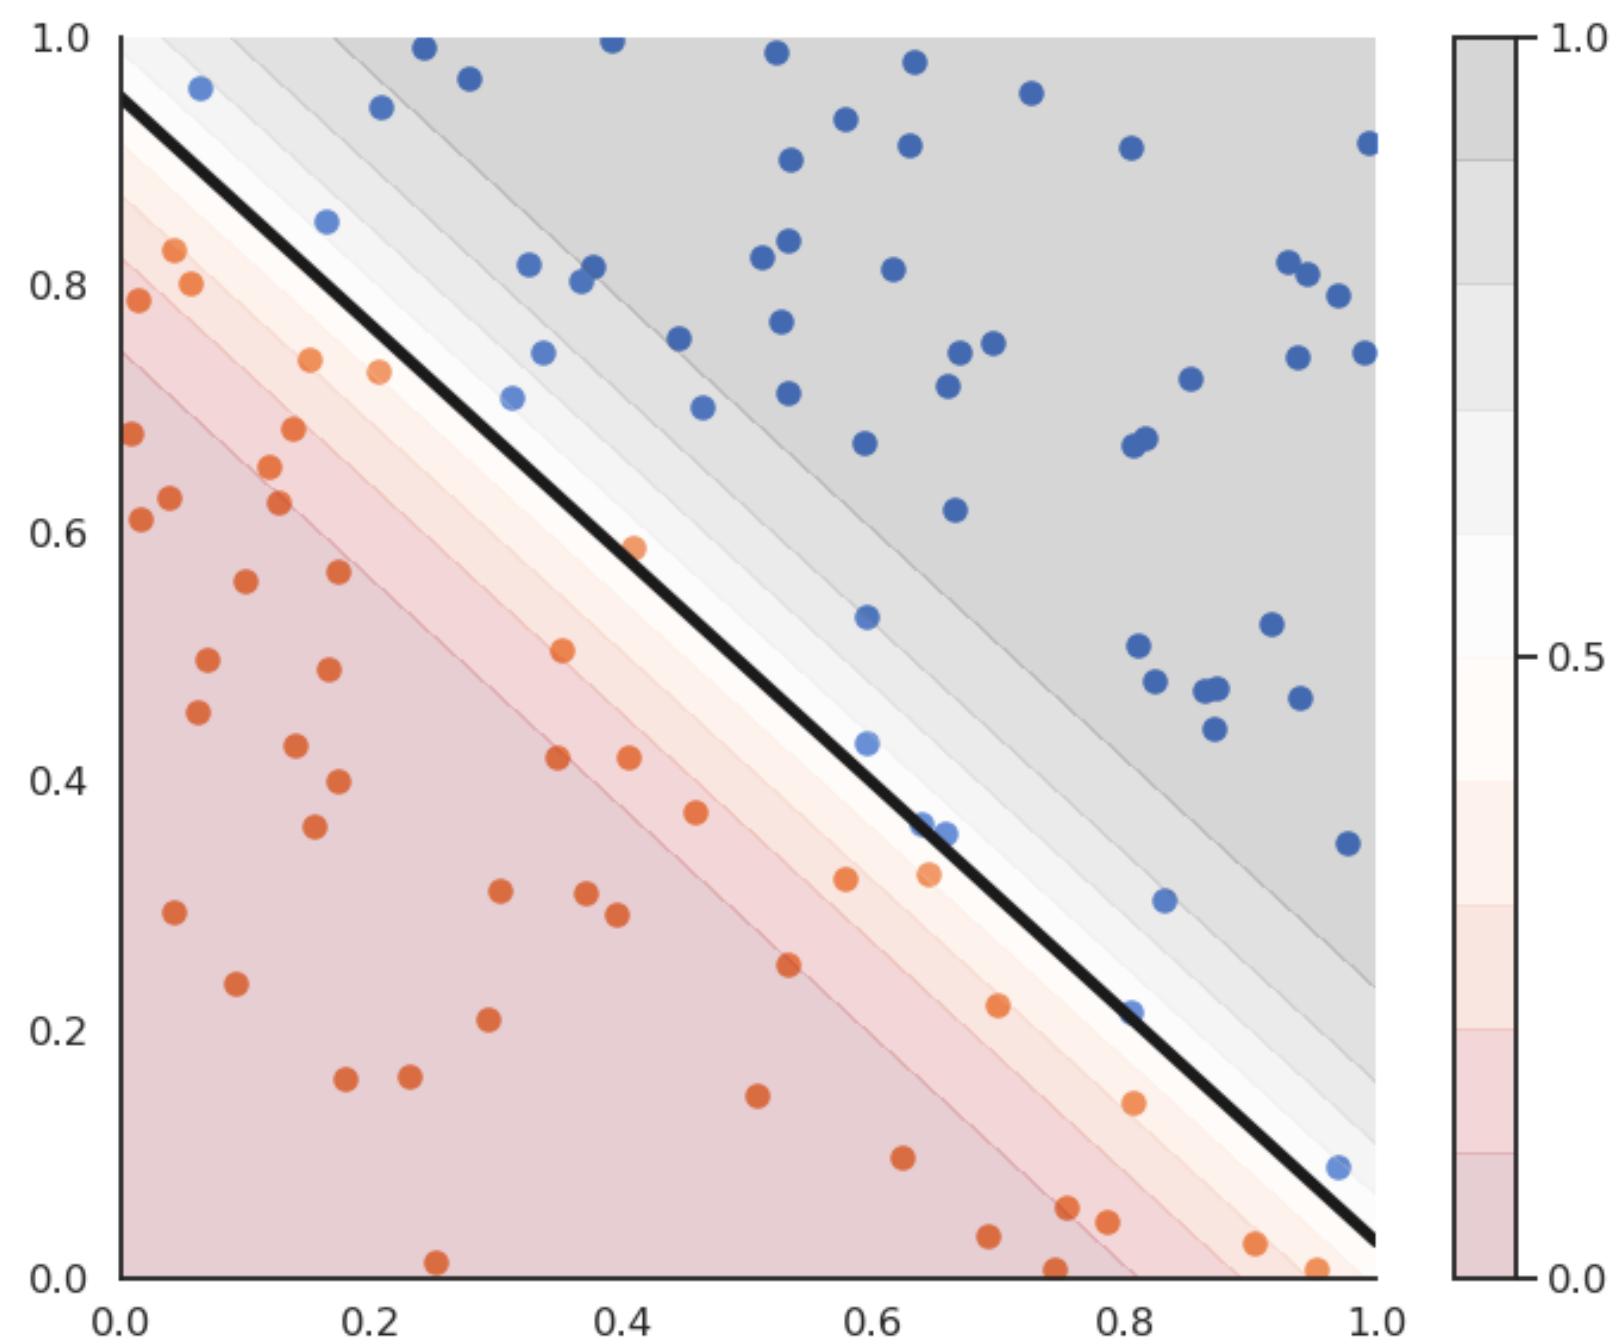


- Logistic regression works just like linear classification, except in the way the prediction is done.
- To know to which class  $\mathbf{x}_i$  belongs, simply draw a random number between 0 and 1:
  - if it is smaller than  $y_i$  (probability  $y_i$ ), it belongs to the positive class.
  - if it is bigger than  $y_i$  (probability  $1 - y_i$ ), it belongs to the negative class.
- Alternatively, you can put a **hard limit** at 0.5:
  - if  $y_i > 0.5$  then the class is positive.
  - if  $y_i < 0.5$  then the class is negative.

# Logistic regression



# Logistic regression and confidence score

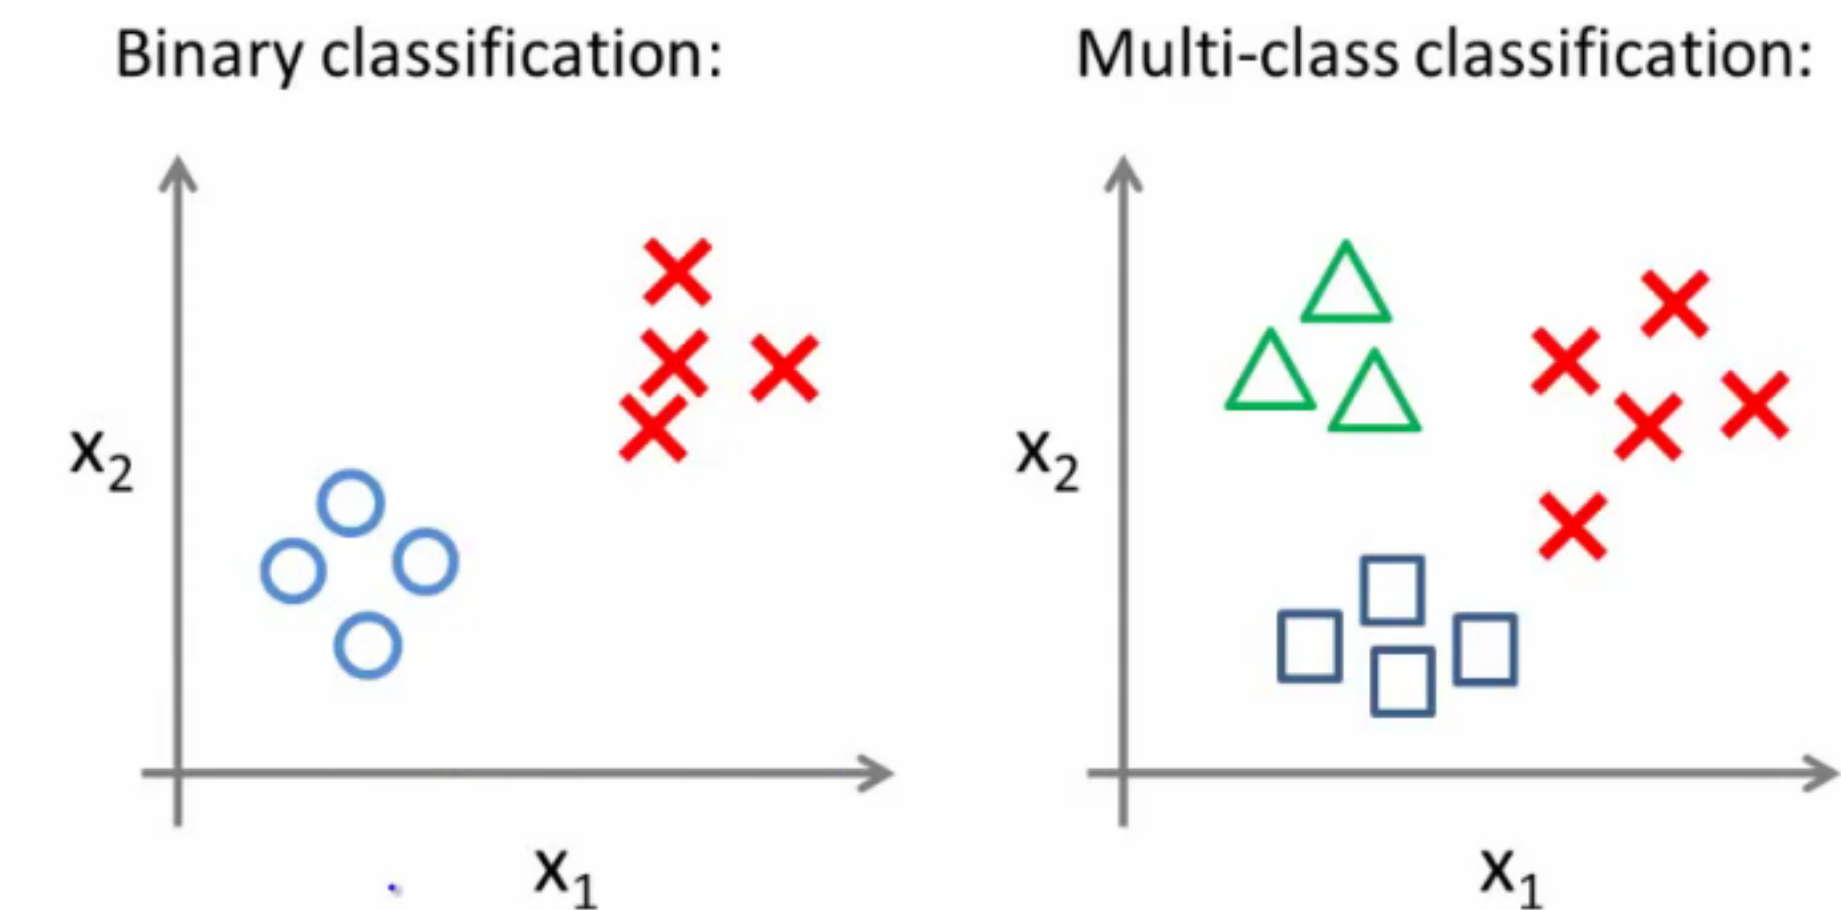


- Logistic regression also provides a **confidence score**:
  - the closer  $y$  is from 0 or 1, the more confident we can be that the classification is correct.
- This is particularly important in **safety critical** applications:
  - If you detect the positive class but with a confidence of 0.51, you should perhaps not trust the prediction.
  - If the confidence score is 0.99, you can probably trust the prediction.

## 4 - Multi-class classification

# Multi-class classification

- Can we perform multi-class classification using the previous methods when  $t \in \{A, B, C\}$  instead of  $t = +1$  or  $-1$ ?



# Multi-class classification

Two main solutions:

- **One-vs-All** (or One-vs-the-rest): one trains simultaneously a binary (linear) classifier for each class. The examples belonging to this class form the positive class, all others are the negative class:
  - A vs. B and C
  - B vs. A and C
  - C vs. A and B

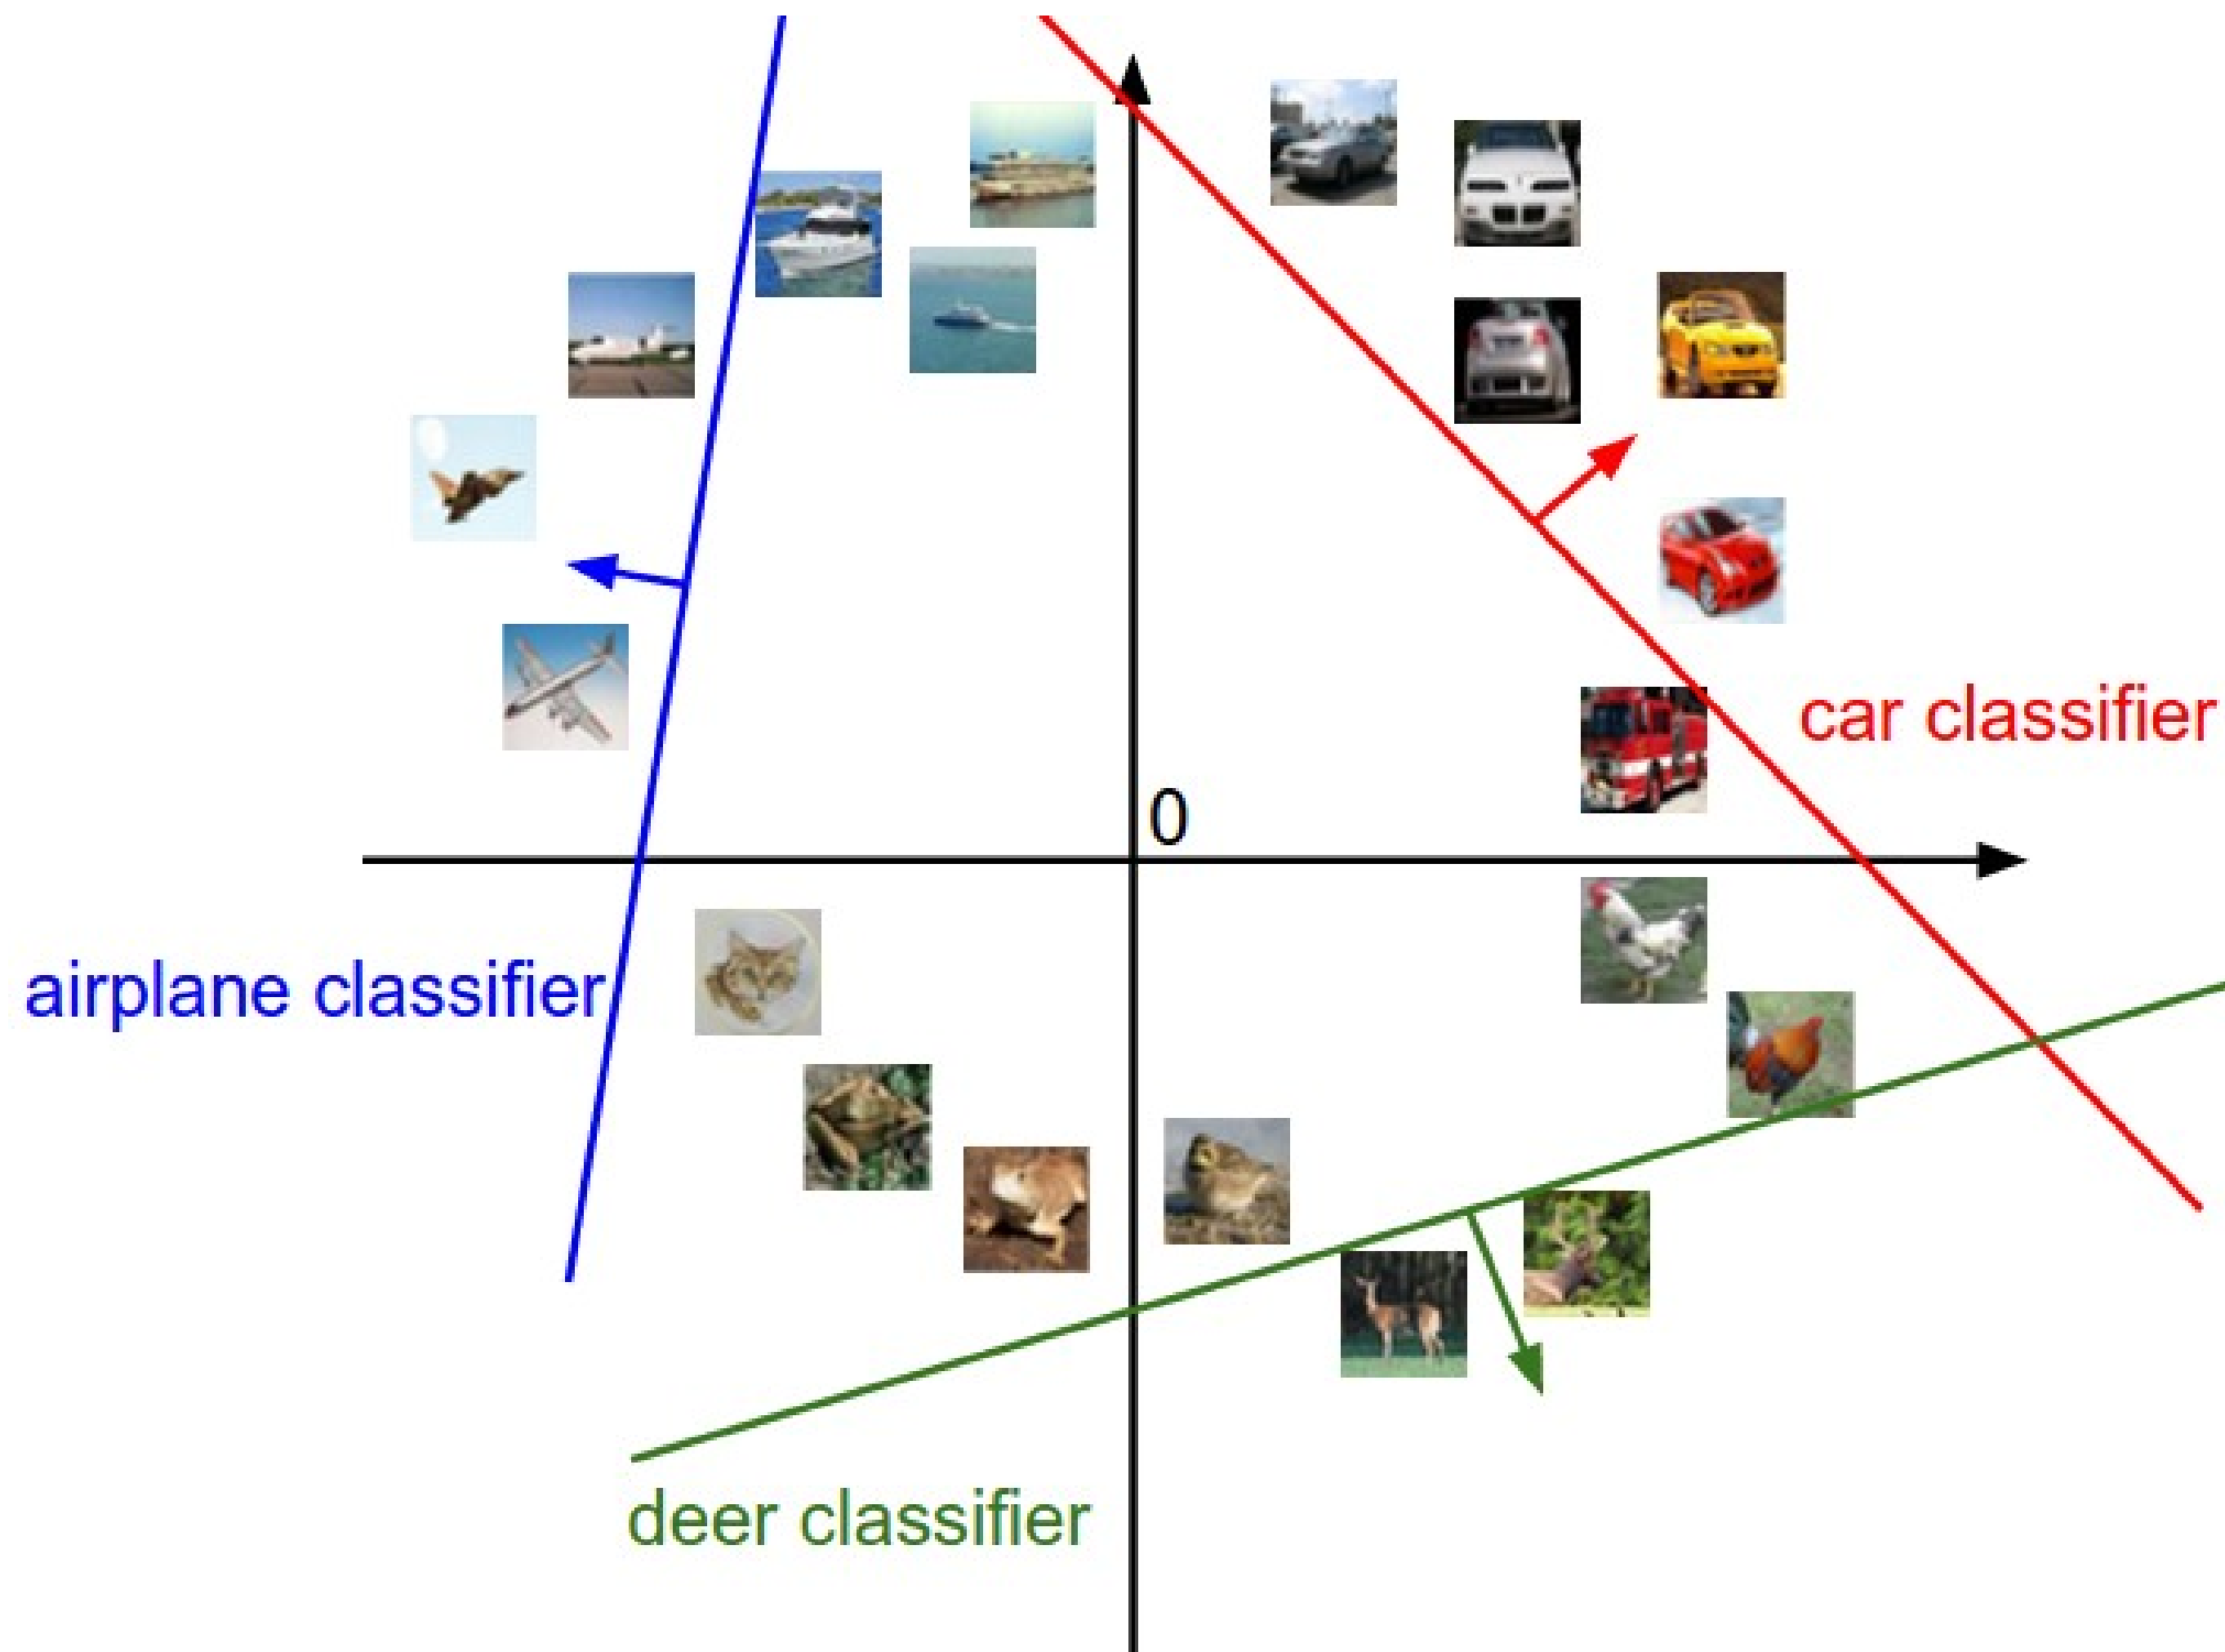
If multiple classes are predicted for a single example, one needs a confidence level for each classifier saying how sure it is of its prediction.

- **One-vs-One**: one trains a classifier for each pair of class:
  - A vs. B
  - B vs. C
  - C vs. A

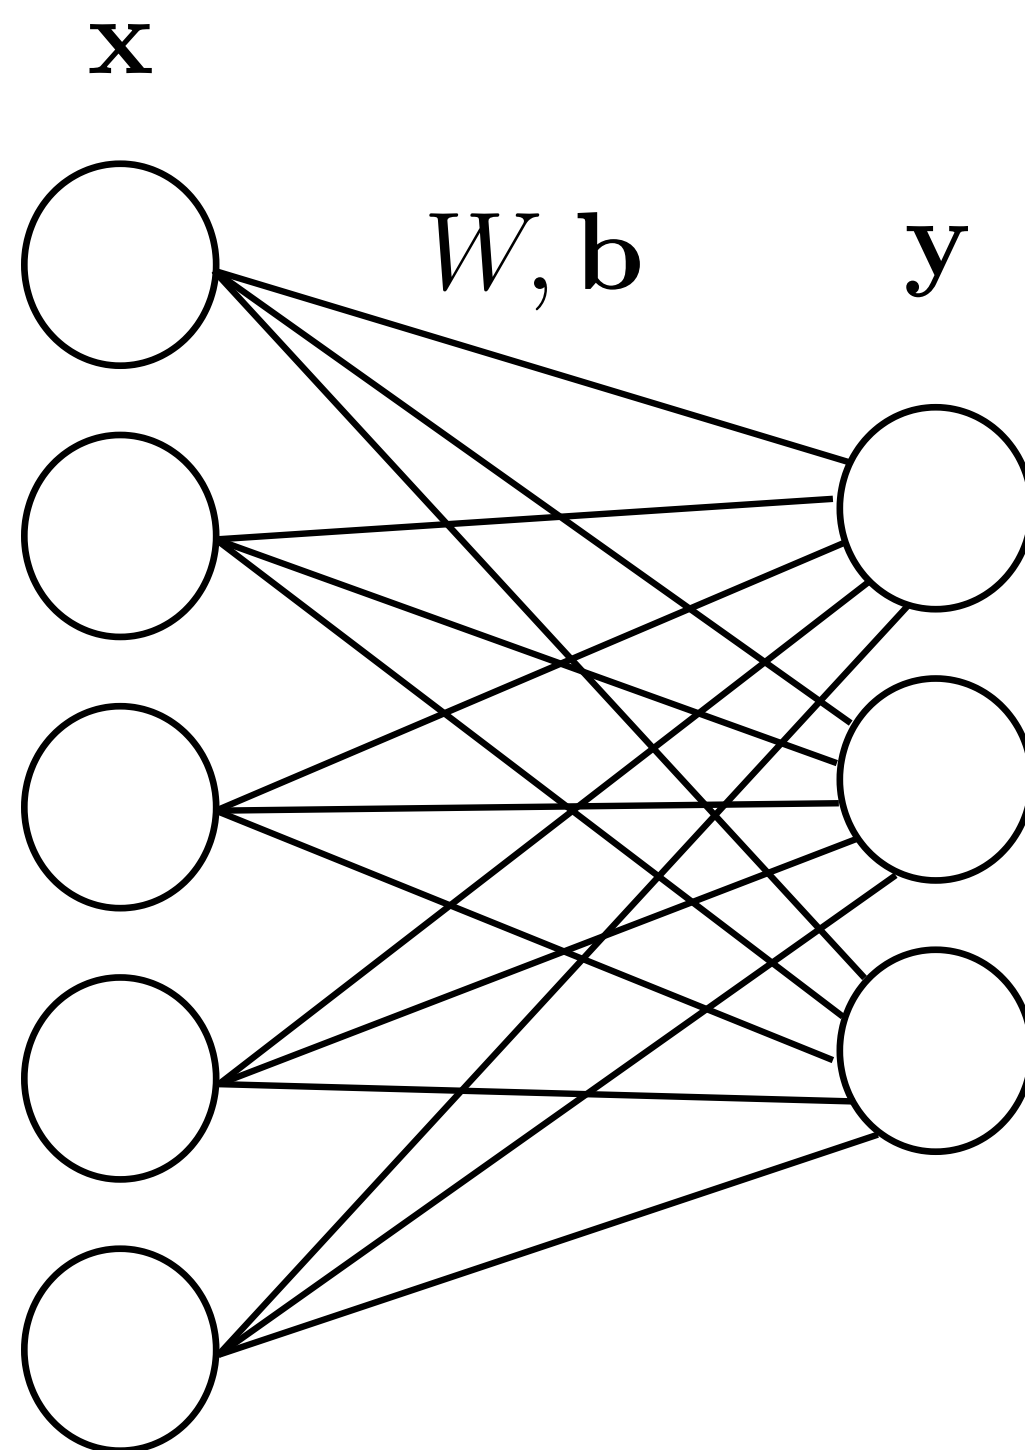
A majority vote is then performed to find the correct class.

# Multi-class classification

- Example of **One-vs-All** classification: one binary classifier per class.



# Softmax linear classifier



- Suppose we have  $C$  classes (dog vs. cat vs. ship vs...).
- The One-vs-All scheme involves  $C$  binary classifiers  $(\mathbf{w}_i, b_i)$ , each with a weight vector and a bias, working on the same input  $\mathbf{x}$ .

$$y_i = f(\langle \mathbf{w}_i \cdot \mathbf{x} \rangle + b_i)$$

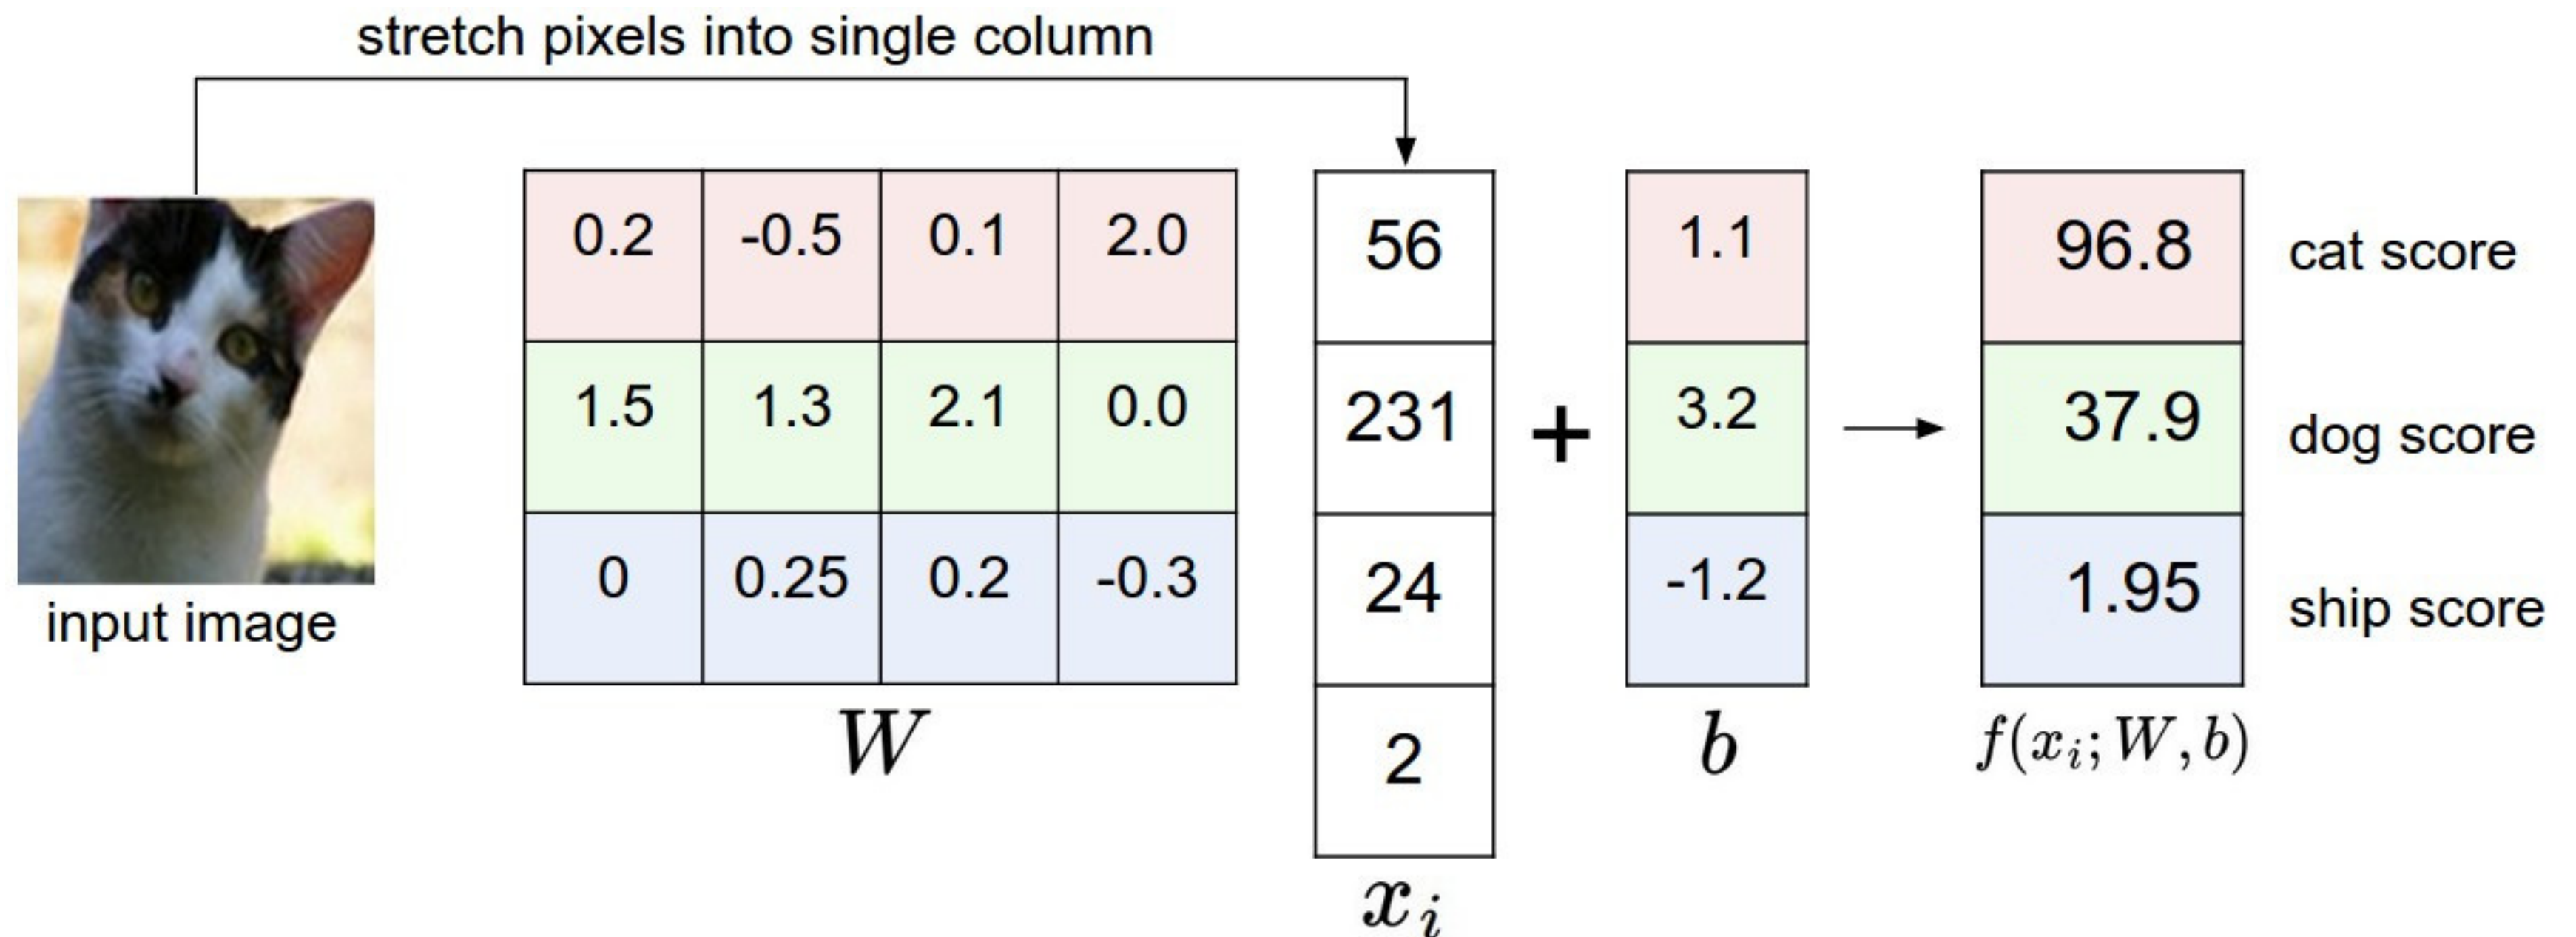
- Putting all neurons together, we obtain a **linear perceptron** similar to multiple linear regression:

$$\mathbf{y} = f(W \times \mathbf{x} + \mathbf{b})$$

- The  $C$  weight vectors form a  $C \times d$  **weight matrix**  $W$ , the biases form a vector  $\mathbf{b}$ .



# Softmax linear classifier



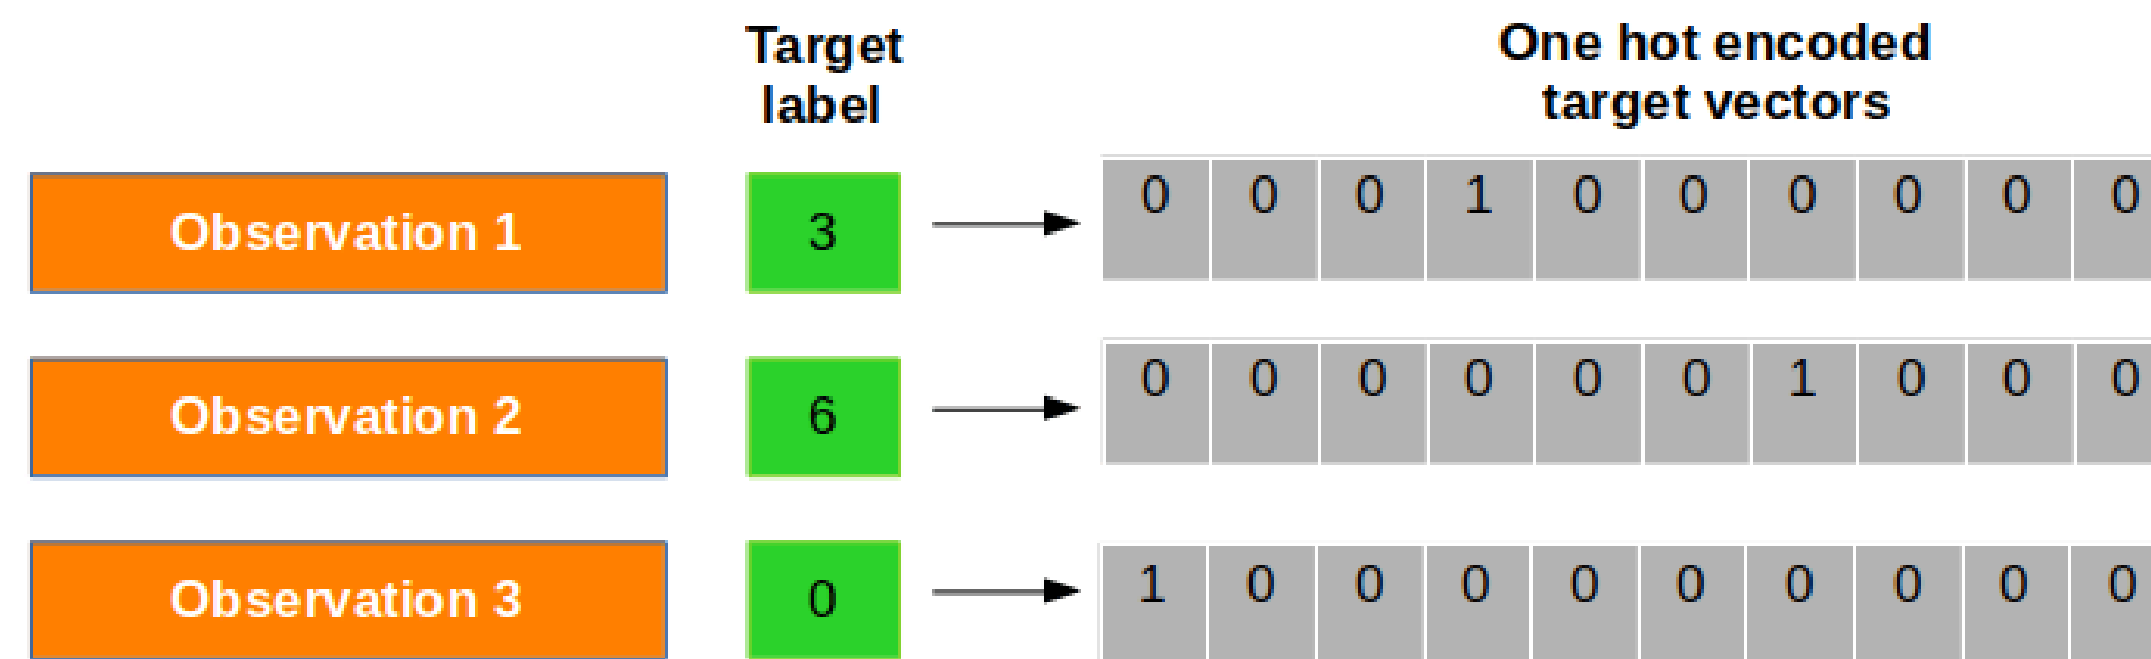
- The net activations form a vector  $\mathbf{z}$ :

$$\mathbf{z} = f_{W,b}(\mathbf{x}) = W \times \mathbf{x} + \mathbf{b}$$

- Each element  $z_j$  of the vector  $\mathbf{z}$  is called the **logit score** of the class:
  - the higher the score, the more likely the input belongs to this class.
- The logit scores are not probabilities, as they can be negative and do not sum to 1.

# One-hot encoding

- How do we represent the ground truth  $\mathbf{t}$  for each neuron?
- The target vector  $\mathbf{t}$  is represented using **one-hot encoding**.



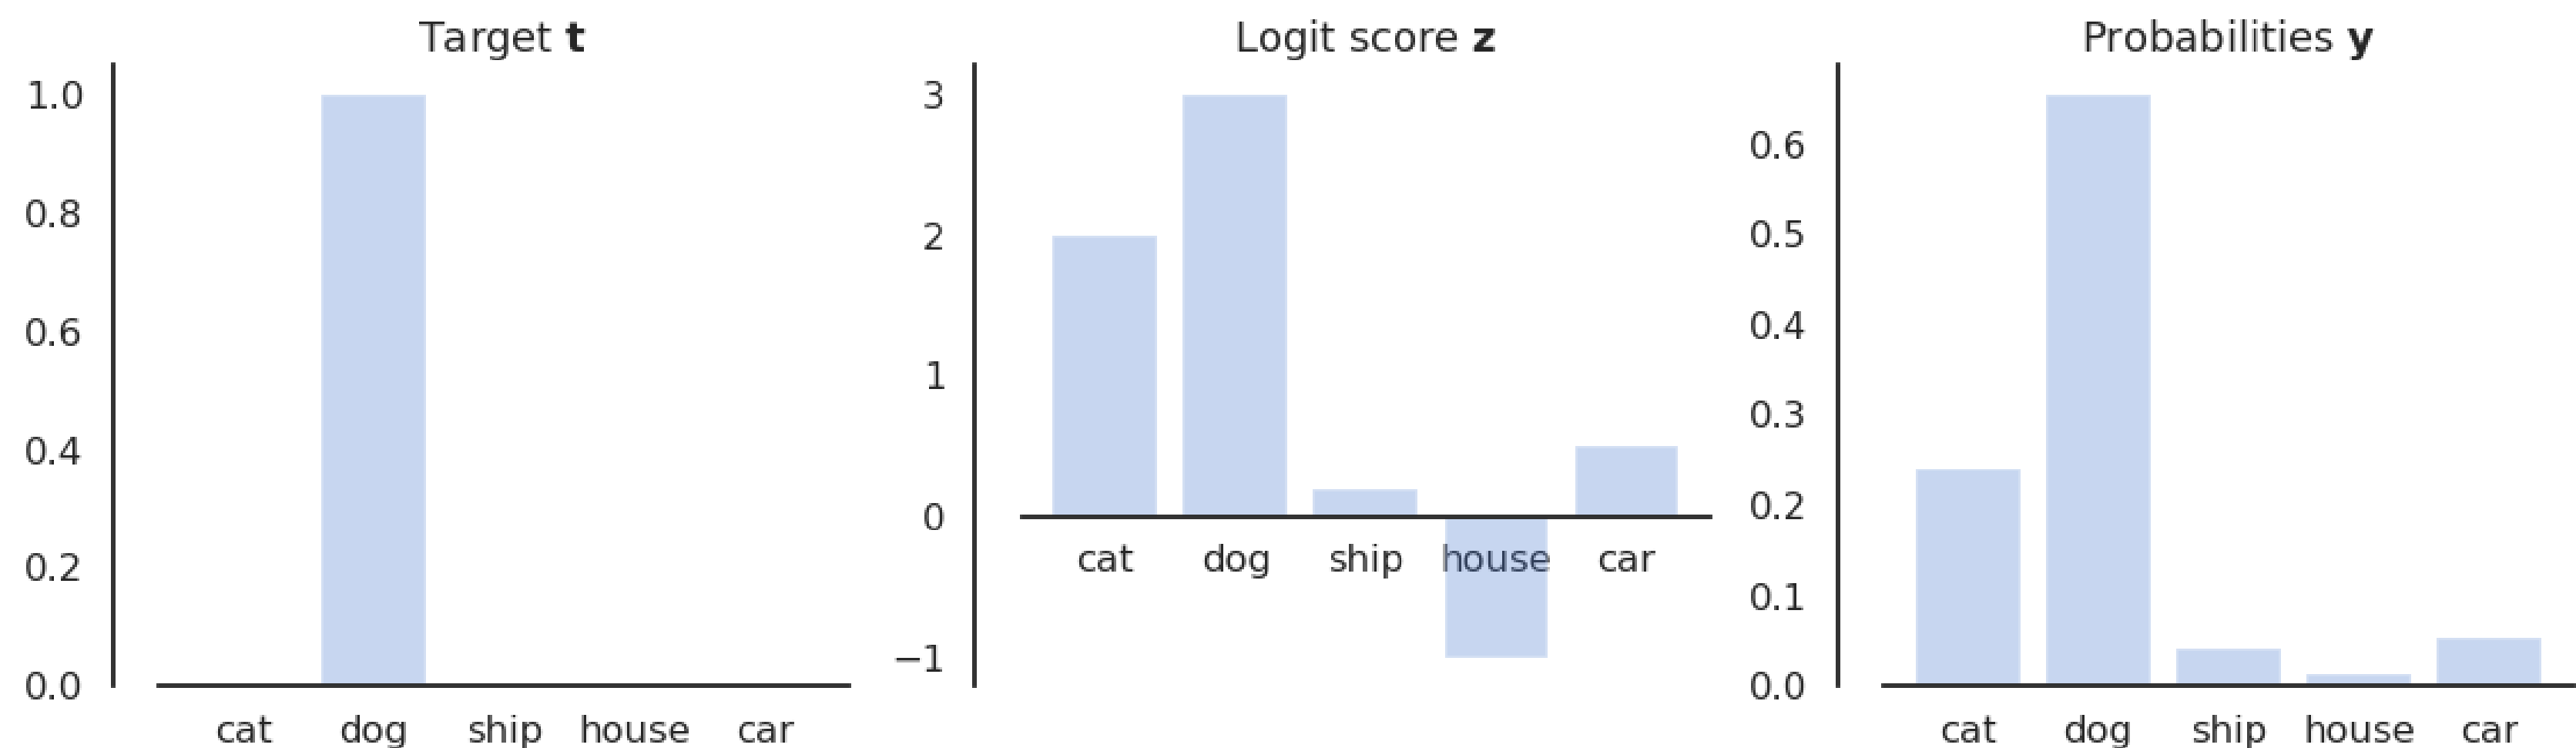
- The binary vector has one element per class: only one element is 1, the others are 0.
- Example:

$$\mathbf{t} = [\text{cat}, \text{dog}, \text{ship}, \text{house}, \text{car}] = [0, 1, 0, 0, 0]$$

# One-hot encoding

- The labels can be seen as a **probability distribution** over the training set, in this case a **multinomial** distribution (a dice with  $C$  sides).
- For a given image  $\mathbf{x}$  (e.g. a picture of a dog), the conditional pmf is defined by the one-hot encoded vector  $\mathbf{t}$ :

$$P(\mathbf{t}|\mathbf{x}) = [P(\text{cat}|\mathbf{x}), P(\text{dog}|\mathbf{x}), P(\text{ship}|\mathbf{x}), P(\text{house}|\mathbf{x}), P(\text{car}|\mathbf{x})] = [0, 1, 0, 0, 0]$$

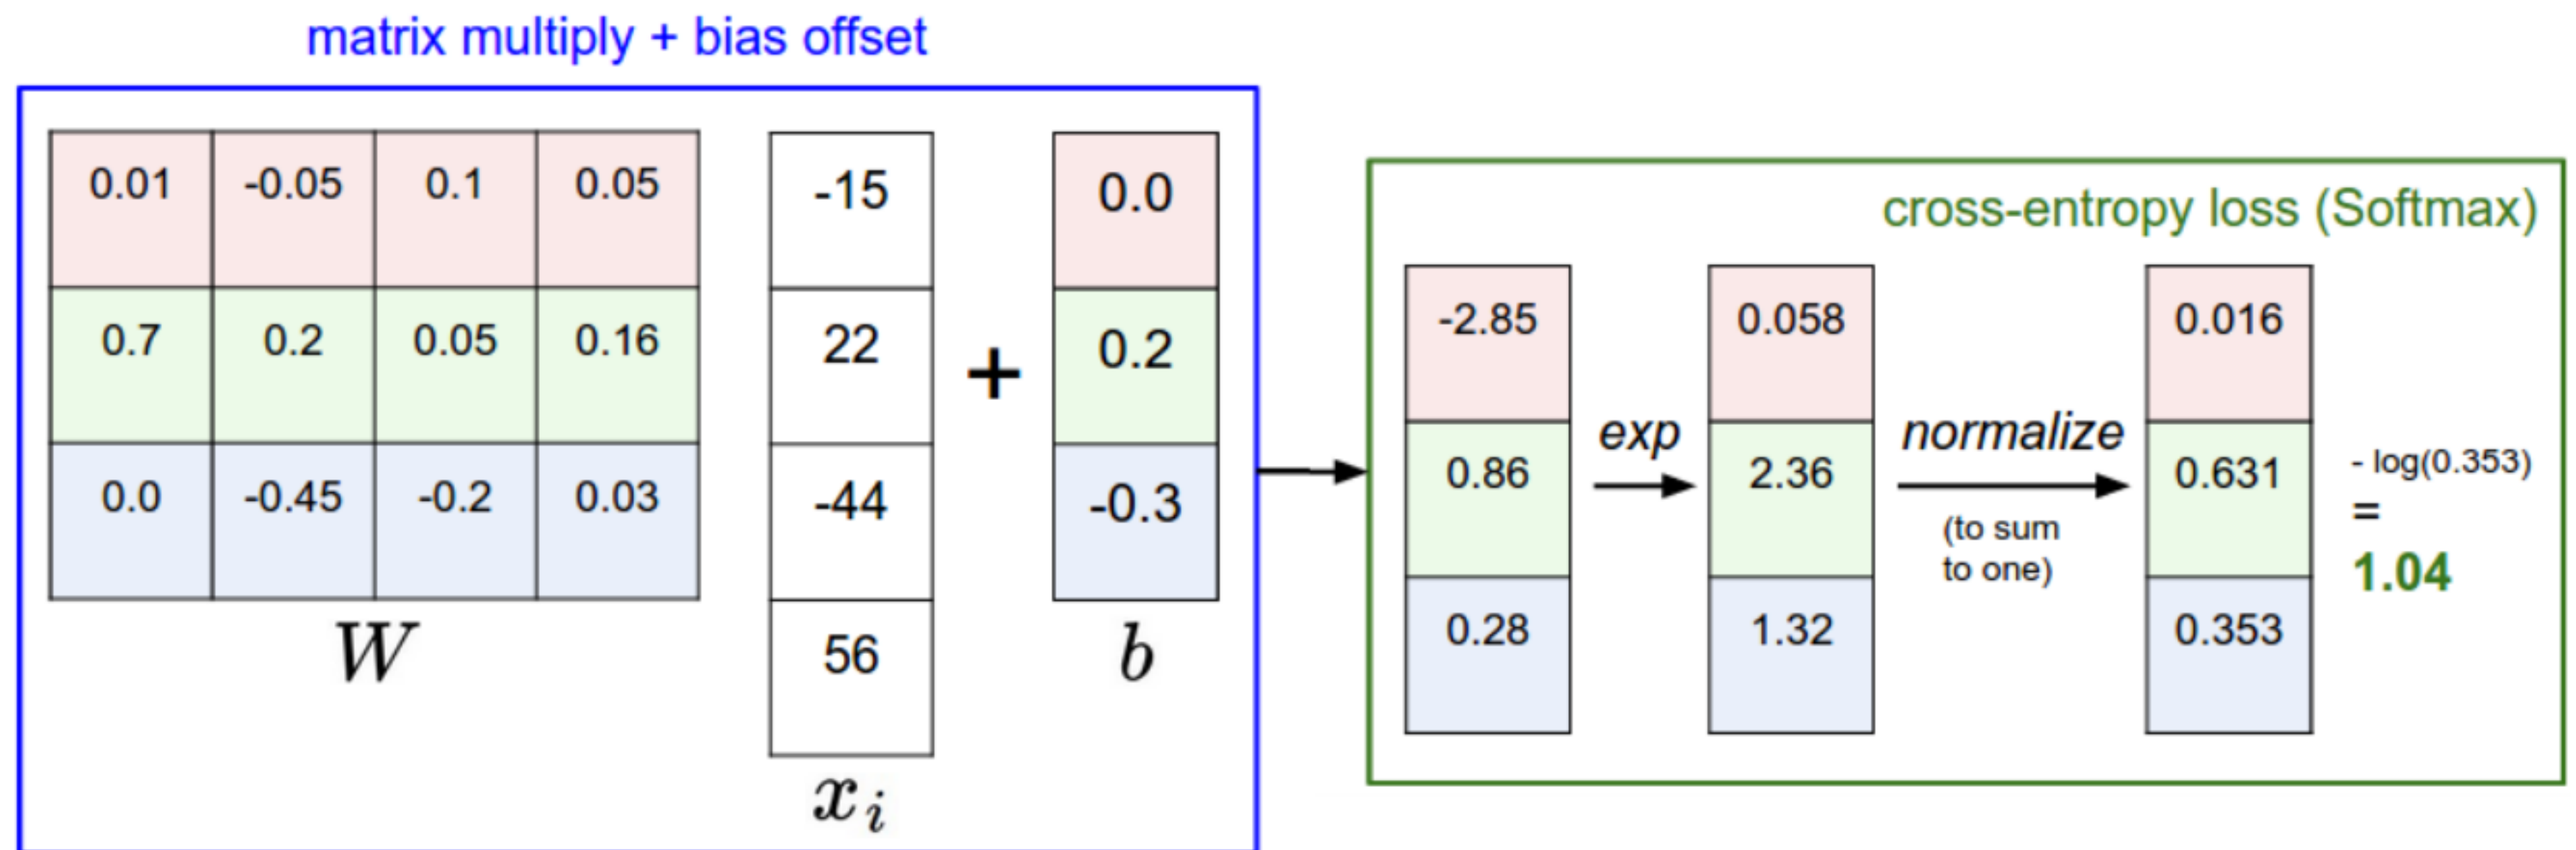


- We need to transform the logit score  $\mathbf{z}$  into a **probability distribution**  $P(\mathbf{y}|\mathbf{x})$  that should be as close as possible from  $P(\mathbf{t}|\mathbf{x})$ .

# Softmax linear classifier

- The **softmax** operator makes sure that the sum of the outputs  $\mathbf{y} = \{y_i\}$  over all classes is 1.

$$y_j = P(\text{class} = j | \mathbf{x}) = \mathcal{S}(z_j) = \frac{\exp(z_j)}{\sum_k \exp(z_k)}$$

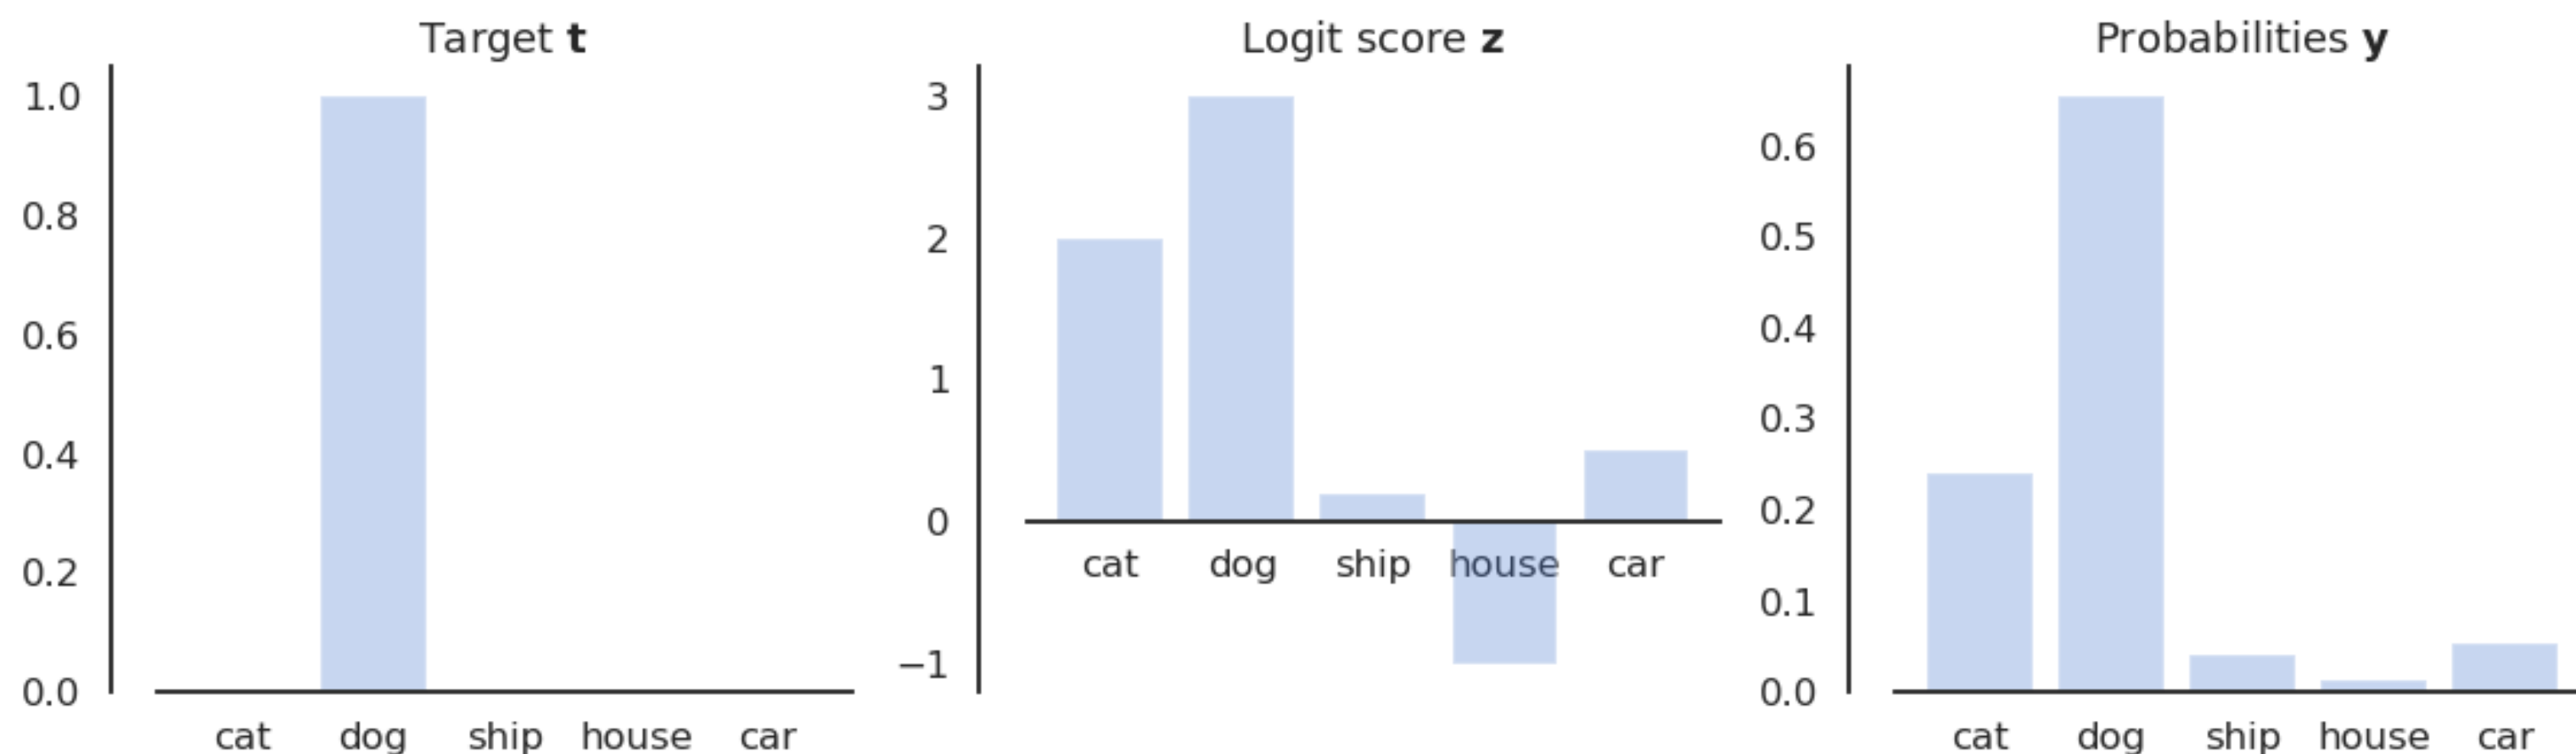


- The higher  $z_j$ , the higher the probability that the example belongs to class  $j$ .
- This is very similar to logistic regression for soft classification, except that we have multiple classes.

# Cross-entropy loss function

- We cannot use the mse as a loss function, as the softmax function would be hard to differentiate:

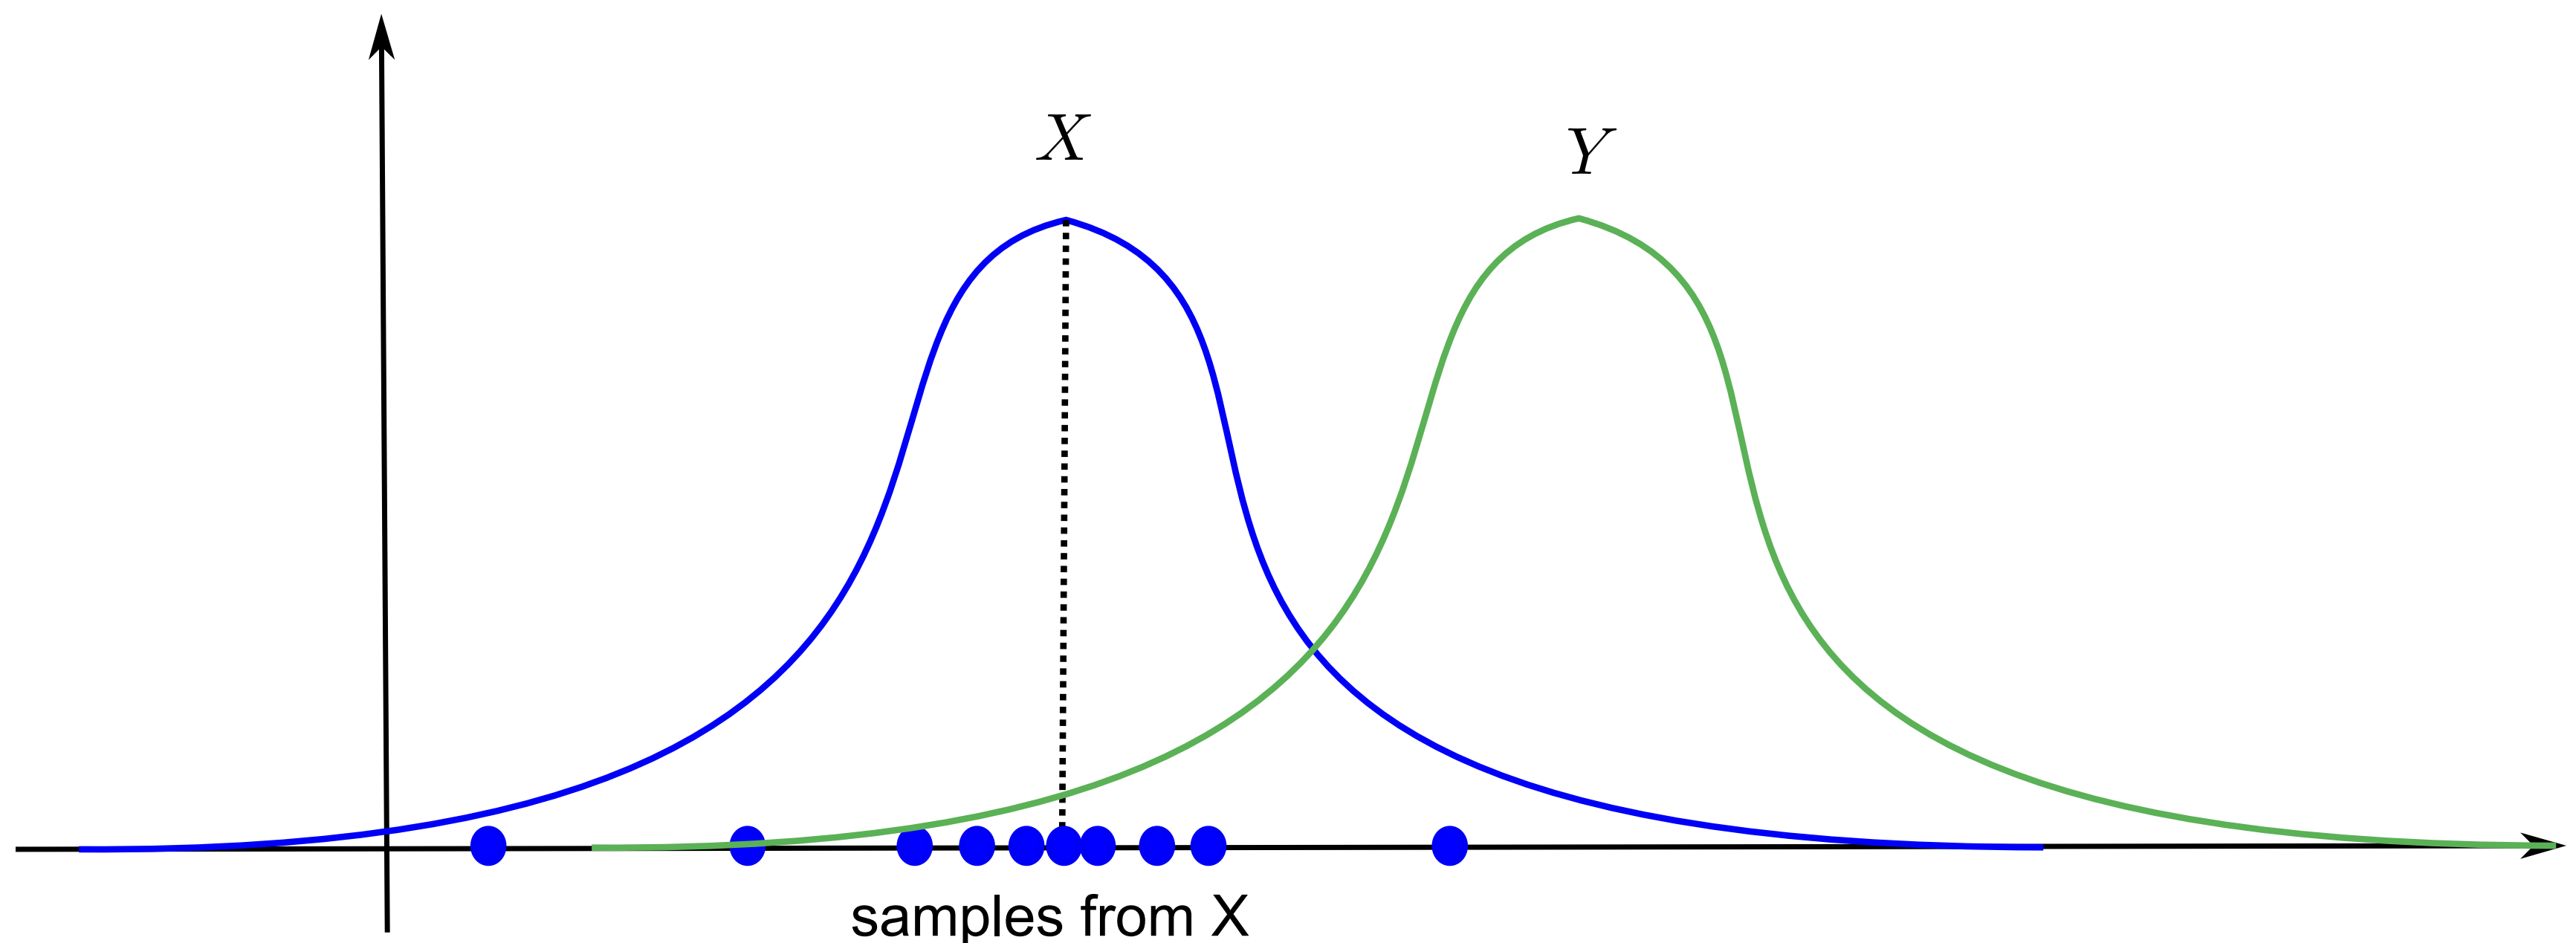
$$\text{mse}(W, \mathbf{b}) = \sum_j (t_j - \frac{\exp(z_j)}{\sum_k \exp(z_k)})^2$$



- We actually want to minimize the statistical distance between two distributions:
  - The model outputs a multinomial probability distribution  $\mathbf{y}$  for an input  $\mathbf{x}$ :  $P(\mathbf{y}|\mathbf{x}; W, \mathbf{b})$ .
  - The one-hot encoded classes also come from a multinomial probability distribution  $P(\mathbf{t}|\mathbf{x})$ .
- We search which parameters  $(W, \mathbf{b})$  make the two distributions  $P(\mathbf{y}|\mathbf{x}; W, \mathbf{b})$  and  $P(\mathbf{t}|\mathbf{x})$  close.

# Cross-entropy loss function

- The training data  $\{\mathbf{x}_i, \mathbf{t}_i\}$  represents samples from  $P(\mathbf{t}|\mathbf{x})$ .
- $P(\mathbf{y}|\mathbf{x}; W, \mathbf{b})$  is a good model of the data when the two distributions are close, i.e. when the **negative log-likelihood** of each sample under the model is small.




- For an input  $\mathbf{x}$ , we minimize the **cross-entropy** between the target distribution and the predicted outputs:

$$l(W, \mathbf{b}) = \mathcal{H}(\mathbf{t}|\mathbf{x}, \mathbf{y}|\mathbf{x}) = \mathbb{E}_{t \sim P(\mathbf{t}|\mathbf{x})} [-\log P(\mathbf{y} = t|\mathbf{x})]$$



# Cross-entropy and negative log-likelihood



- The cross-entropy samples from  $\mathbf{t}|\mathbf{x}$ :

$$l(W, \mathbf{b}) = \mathcal{H}(\mathbf{t}|\mathbf{x}, \mathbf{y}|\mathbf{x}) = \mathbb{E}_{t \sim P(\mathbf{t}|\mathbf{x})} [-\log P(\mathbf{y} = t|\mathbf{x})] = - \sum_{j=1}^C P(t_j|\mathbf{x}) \log P(y_j = t_j|\mathbf{x})$$

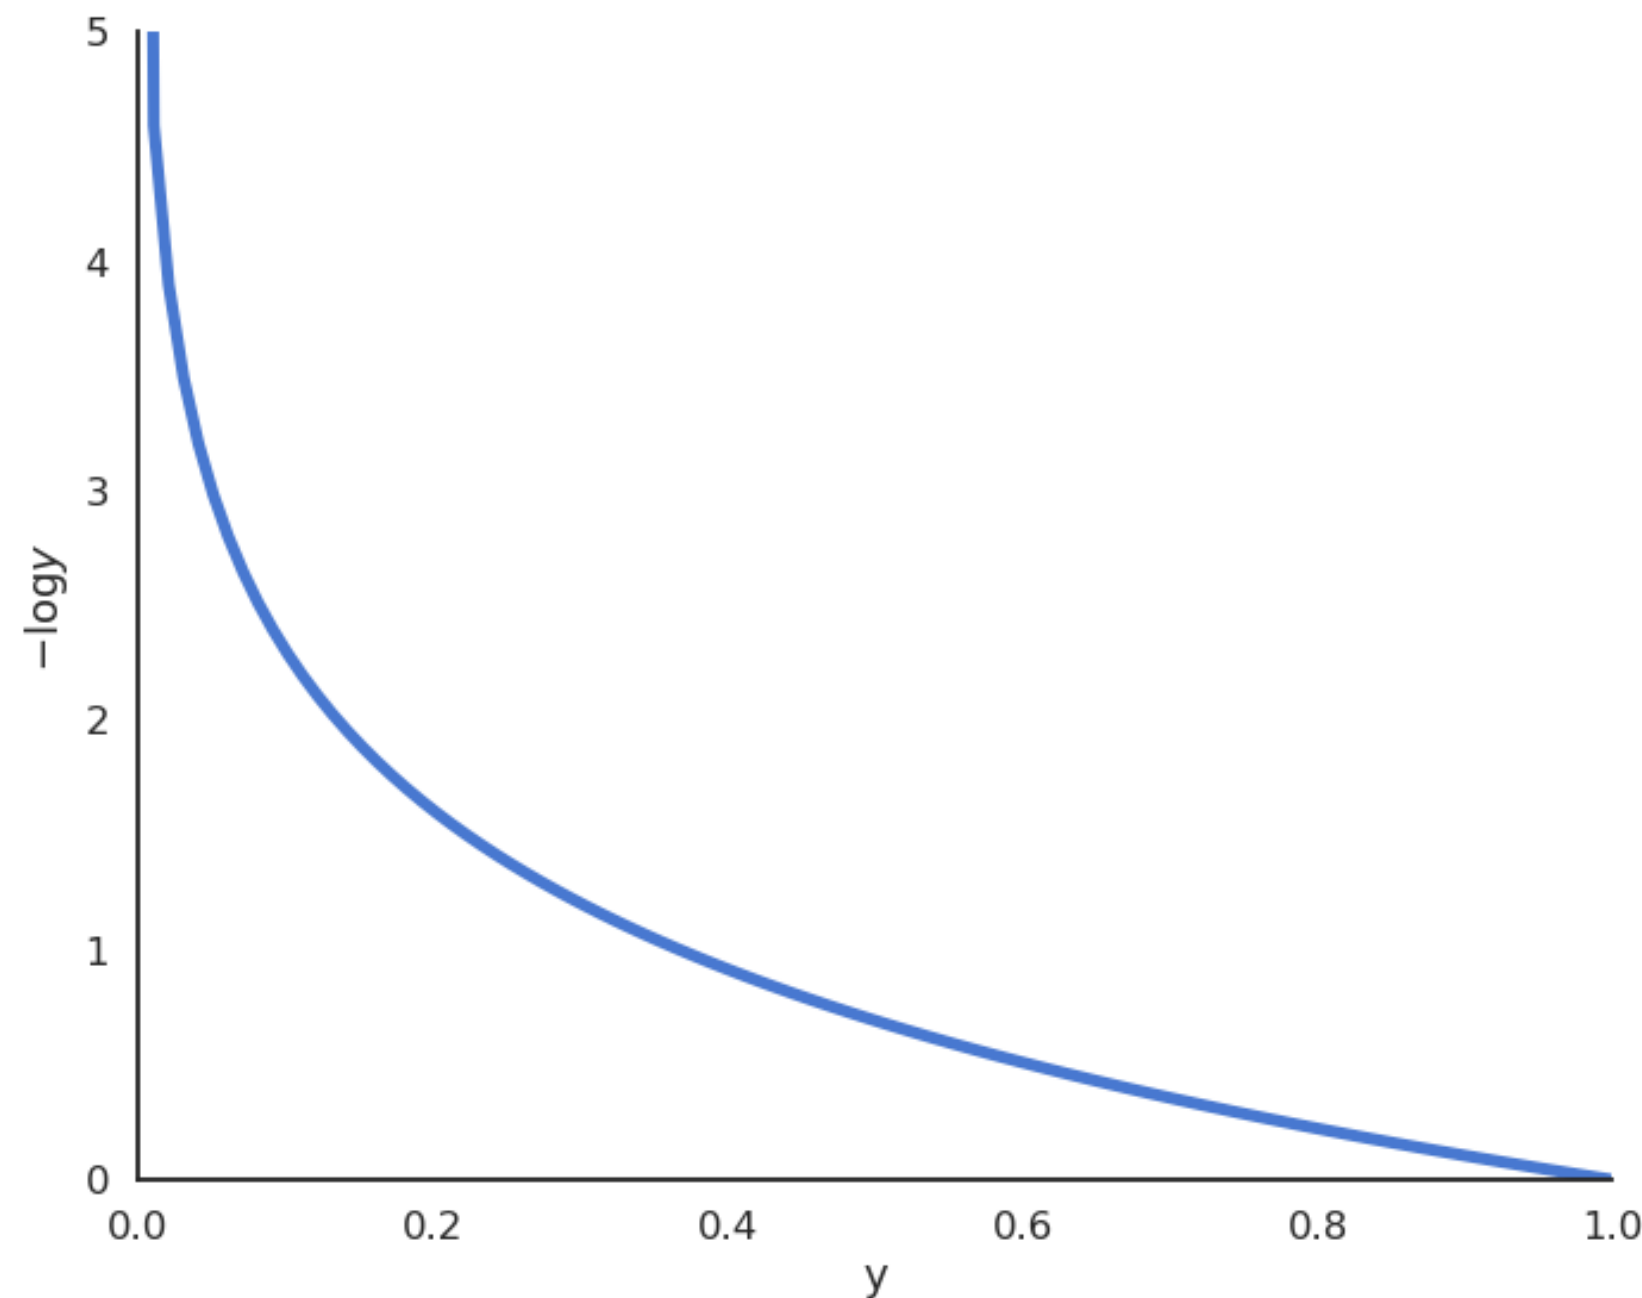
- For a given input  $\mathbf{x}$ ,  $\mathbf{t}$  is non-zero only for the correct class  $t^*$ , as  $\mathbf{t}$  is a one-hot encoded vector  $[0, 1, 0, 0, 0]$ :

$$l(W, \mathbf{b}) = -\log P(\mathbf{y} = t^*|\mathbf{x})$$

- If we note  $j^*$  the index of the correct class  $t^*$ , the cross entropy is simply:

$$l(W, \mathbf{b}) = -\log y_{j^*}$$

# Cross-entropy and negative log-likelihood



- As only one element of  $\mathbf{t}$  is non-zero, the cross-entropy is the same as the **negative log-likelihood** of the prediction for the true label:

$$l(W, \mathbf{b}) = -\log y_{j^*}$$

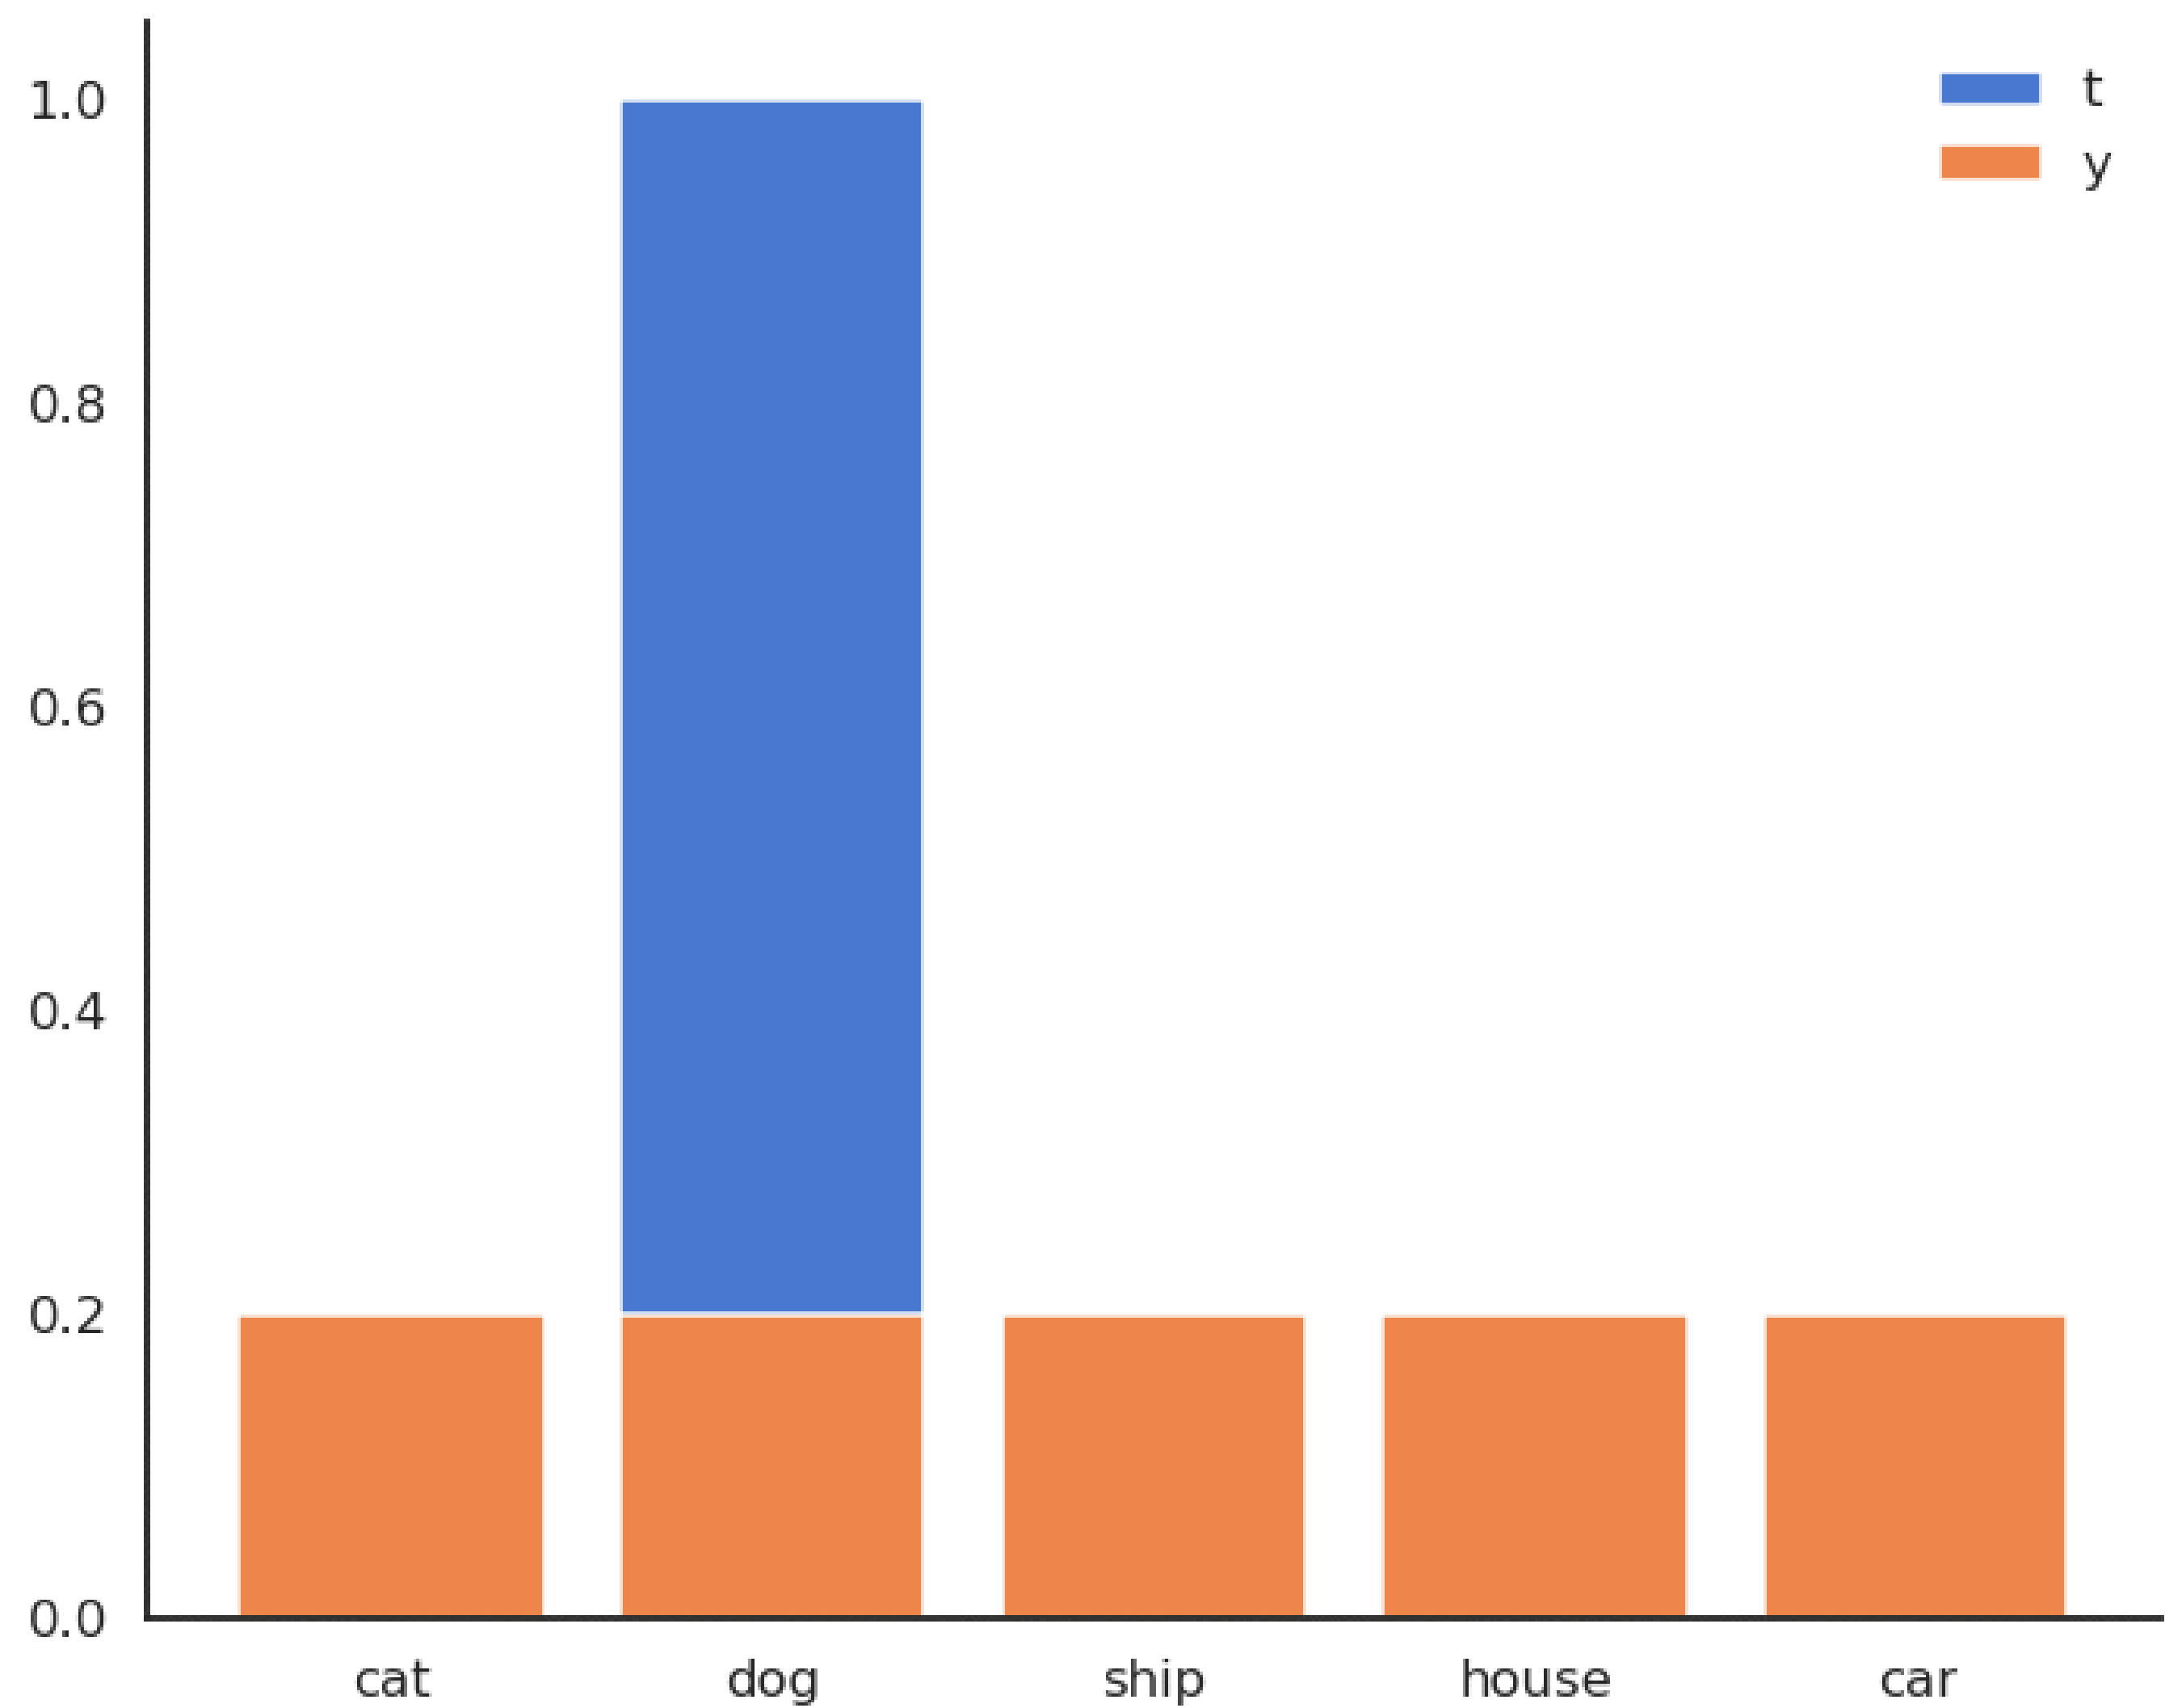
- The minimum of  $-\log y$  is obtained when  $y = 1$ :
  - We want to classifier to output a probability 1 for the true label.
- Because of the softmax activation function, the probability for the other classes should become closer from 0.

$$y_j = P(\text{class} = j) = \frac{\exp(z_j)}{\sum_k \exp(z_k)}$$

- Minimizing the cross-entropy / negative log-likelihood pushes the output distribution  $\mathbf{y}|\mathbf{x}$  to be as close as possible to the target distribution  $\mathbf{t}|\mathbf{x}$ .



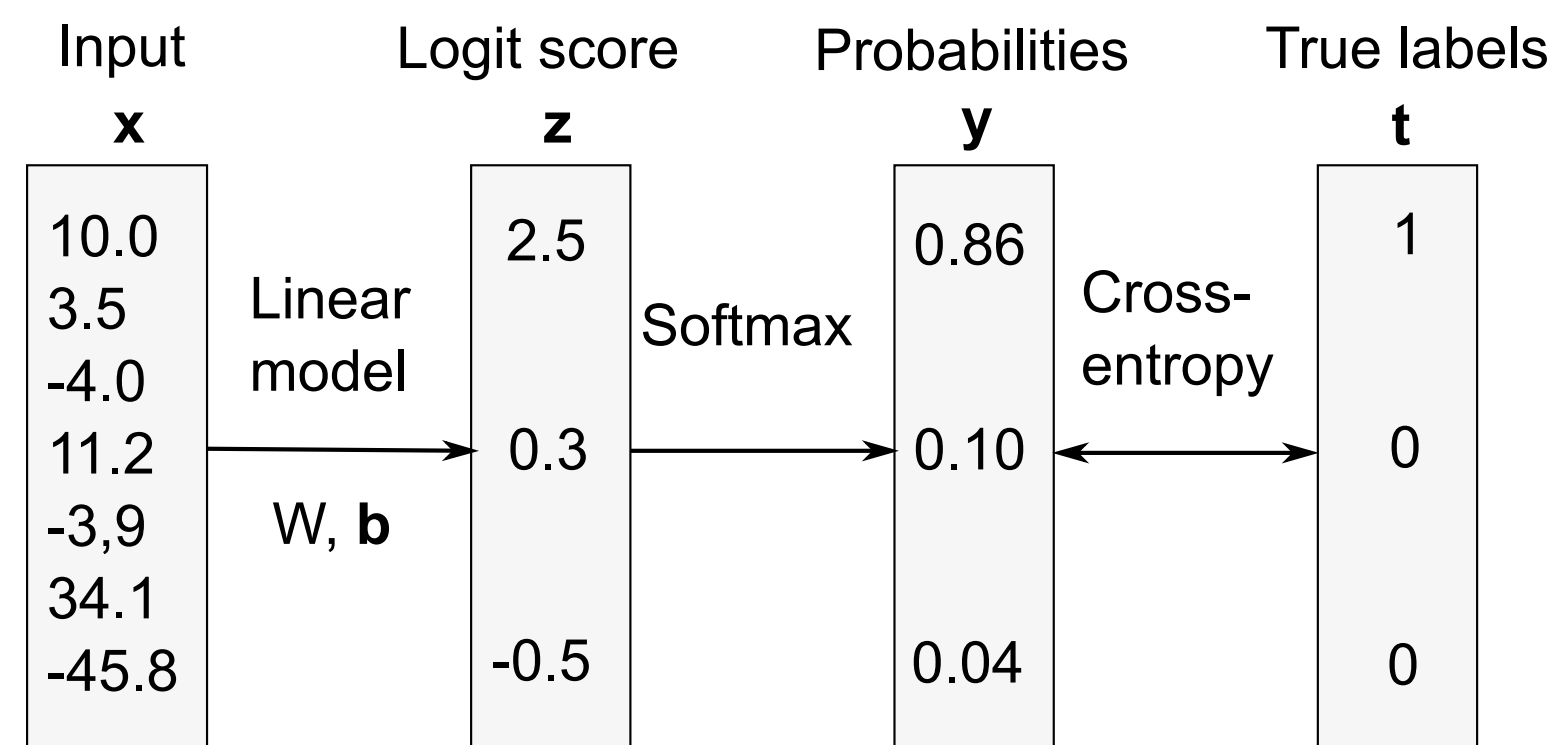
# Cross-entropy loss function



# Cross-entropy loss function

- As  $\mathbf{t}$  is a binary vector  $[0, 1, 0, 0, 0]$ , the cross-entropy / negative log-likelihood can also be noted as the dot product between  $\mathbf{t}$  and  $\log \mathbf{y}$ :

$$l(W, \mathbf{b}) = -\langle \mathbf{t} \cdot \log \mathbf{y} \rangle = -\sum_{j=1}^C t_j \log y_j = -\log y_{j^*}$$



- The **cross-entropy loss function** is then the expectation over the training set of the individual cross-entropies:

$$\mathcal{L}(W, \mathbf{b}) = \mathbb{E}_{\mathbf{x}, \mathbf{t} \sim \mathcal{D}}[-\langle \mathbf{t} \cdot \log \mathbf{y} \rangle] \approx \frac{1}{N} \sum_{i=1}^N -\langle \mathbf{t}_i \cdot \log \mathbf{y}_i \rangle$$

# Cross-entropy loss function

- The nice thing with the **cross-entropy** loss function, when used on a softmax activation function, is that the partial derivative w.r.t the logit score  $\mathbf{z}$  is simple:

$$\begin{aligned}\frac{\partial l(W, \mathbf{b})}{\partial z_i} &= - \sum_j \frac{\partial}{\partial z_i} t_j \log(y_j) = - \sum_j t_j \frac{\partial \log(y_j)}{\partial z_i} = - \sum_j t_j \frac{1}{y_j} \frac{\partial y_j}{\partial z_i} \\ &= - \frac{t_i}{y_i} \frac{\partial y_i}{\partial z_i} - \sum_{j \neq i}^C \frac{t_j}{y_j} \frac{\partial y_j}{\partial z_i} = - \frac{t_i}{y_i} y_i (1 - y_i) - \sum_{j \neq i}^C \frac{t_j}{y_i} (-y_j y_i) \\ &= -t_i + t_i y_i + \sum_{j \neq i}^C t_j y_i = -t_i + \sum_{j=1}^C t_j y_i = -t_i + y_i \sum_{j=1}^C t_j \\ &= -(t_i - y_i)\end{aligned}$$

i.e. the same as with the mse in linear regression!

- Vector notation:

$$\frac{\partial l(W, \mathbf{b})}{\partial \mathbf{z}} = -(\mathbf{t} - \mathbf{y})$$

# Cross-entropy loss function

- As:

$$\mathbf{z} = W \times \mathbf{x} + \mathbf{b}$$

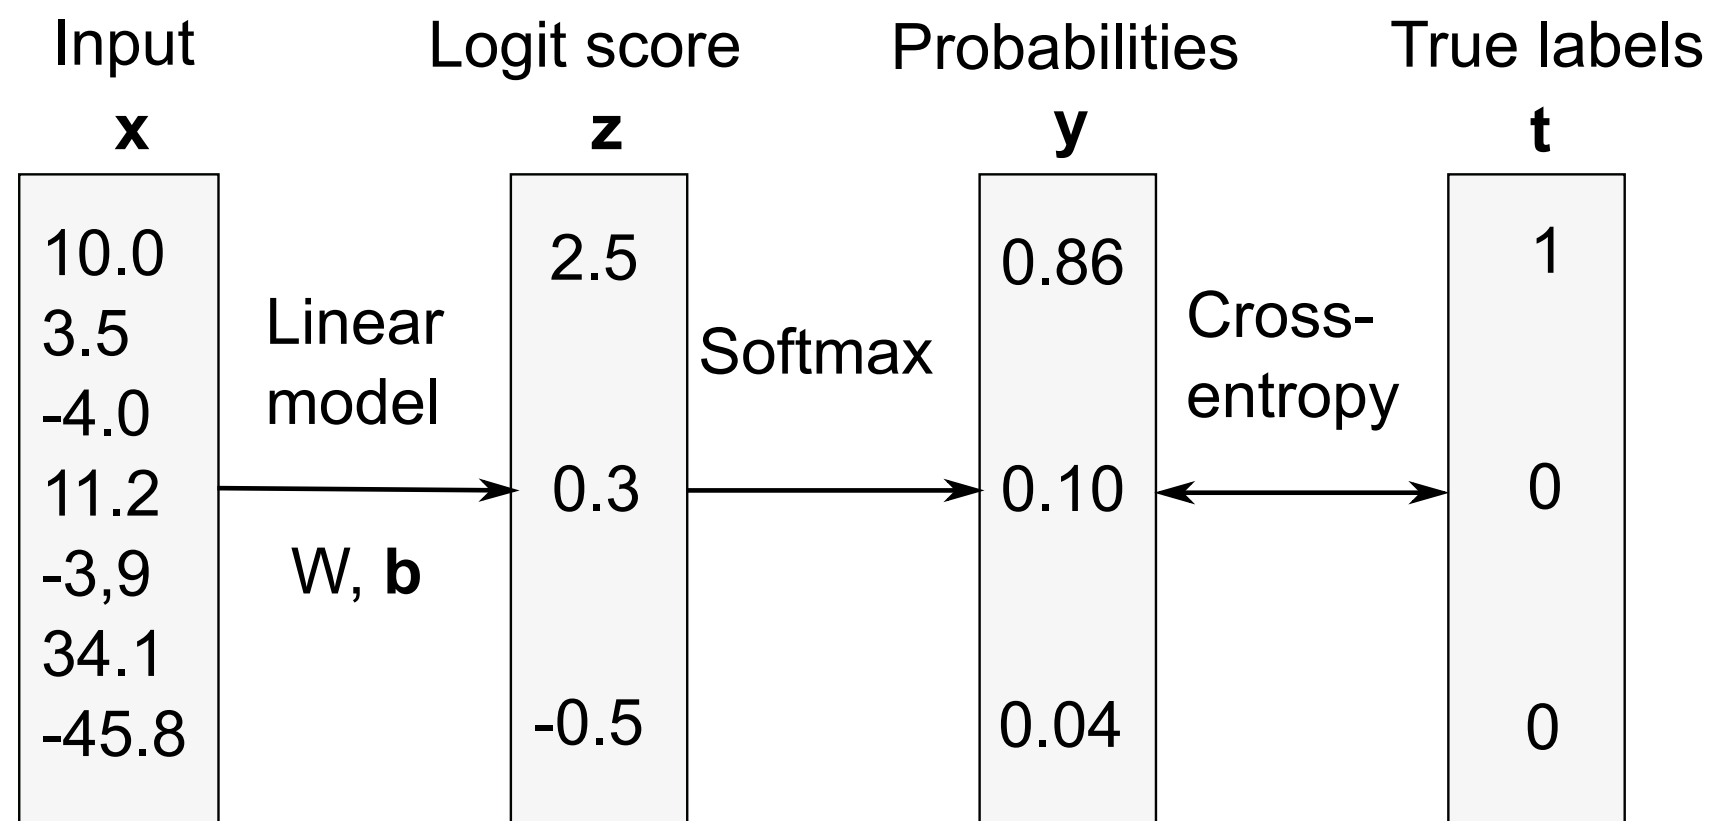
we can obtain the partial derivatives:

$$\begin{cases} \frac{\partial l(W, \mathbf{b})}{\partial W} = \frac{\partial l(W, \mathbf{b})}{\partial \mathbf{z}} \times \frac{\partial \mathbf{z}}{\partial W} = -(\mathbf{t} - \mathbf{y}) \times \mathbf{x}^T \\ \frac{\partial l(W, \mathbf{b})}{\partial \mathbf{b}} = \frac{\partial l(W, \mathbf{b})}{\partial \mathbf{z}} \times \frac{\partial \mathbf{z}}{\partial \mathbf{b}} = -(\mathbf{t} - \mathbf{y}) \end{cases}$$

- So gradient descent leads to the **delta learning rule**:

$$\begin{cases} \Delta W = \eta (\mathbf{t} - \mathbf{y}) \times \mathbf{x}^T \\ \Delta \mathbf{b} = \eta (\mathbf{t} - \mathbf{y}) \end{cases}$$

# Softmax linear classifier



- We first compute the **logit scores** **z** using a linear layer:

$$\mathbf{z} = W \times \mathbf{x} + \mathbf{b}$$

- We turn them into probabilities **y** using the **softmax activation function**:

$$y_j = \frac{\exp(z_j)}{\sum_k \exp(z_k)}$$

- We minimize the **cross-entropy / negative log-likelihood** on the training set:

$$\mathcal{L}(W, \mathbf{b}) = \mathbb{E}_{\mathbf{x}, \mathbf{t} \sim \mathcal{D}} [-\langle \mathbf{t} \cdot \log \mathbf{y} \rangle]$$

which simplifies into the **delta learning rule**:

$$\begin{cases} \Delta W = \eta (\mathbf{t} - \mathbf{y}) \times \mathbf{x}^T \\ \Delta \mathbf{b} = \eta (\mathbf{t} - \mathbf{y}) \end{cases}$$

# Comparison of linear classification and regression

- Classification and regression differ in the nature of their outputs: in classification they are discrete, in regression they are continuous values.
- However, when trying to minimize the mismatch between a model  $\mathbf{y}$  and the real data  $\mathbf{t}$ , we have found the same **delta learning rule**:

$$\begin{cases} \Delta W = \eta (\mathbf{t} - \mathbf{y}) \times \mathbf{x}^T \\ \Delta \mathbf{b} = \eta (\mathbf{t} - \mathbf{y}) \end{cases}$$

- Regression and classification are in the end the same problem for us. The only things that needs to be adapted is the activation function of the output and the **loss function**:
  - For regression, we use regular activation functions and the **mean square error** (mse):


$$\mathcal{L}(W, \mathbf{b}) = \mathbb{E}_{\mathbf{x}, \mathbf{t} \in \mathcal{D}} [||\mathbf{t} - \mathbf{y}||^2]$$

- For classification, we use the softmax activation function and the **cross-entropy** (negative log-likelihood) loss function:

$$\mathcal{L}(W, \mathbf{b}) = \mathbb{E}_{\mathbf{x}, \mathbf{t} \sim \mathcal{D}} [-\langle \mathbf{t} \cdot \log \mathbf{y} \rangle]$$

## 5 - Multi-label classification

# Multi-label classification



GK Hart/Vikki Hart/Getty Images

- What if there is more than one label on the image?
- The target vector  $\mathbf{t}$  does not represent a probability distribution anymore:

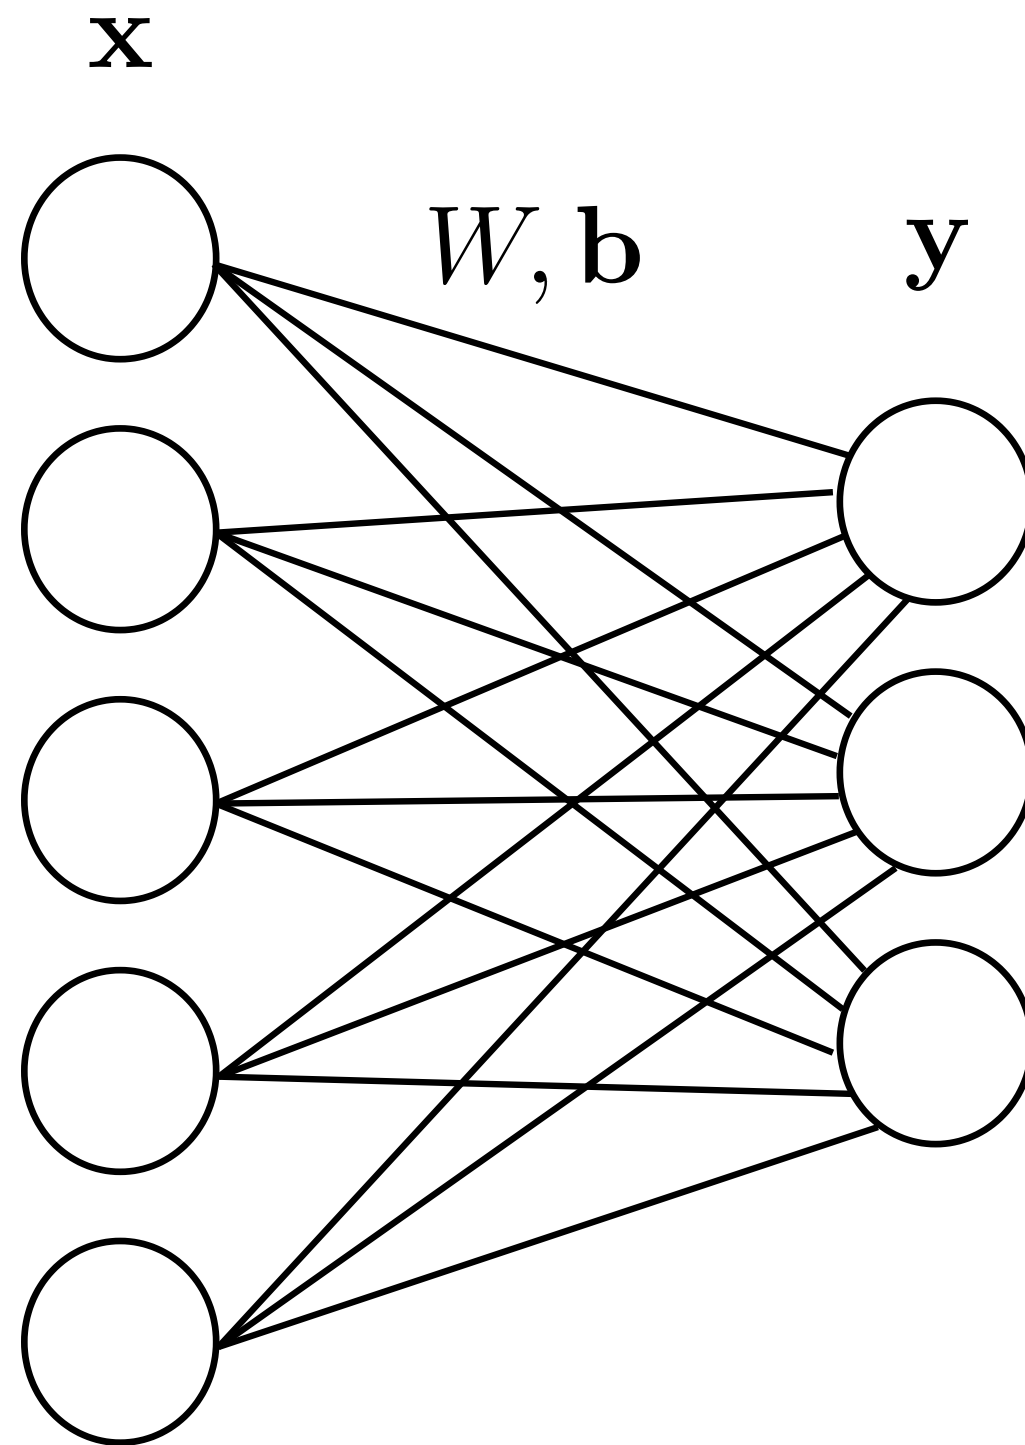
$$\mathbf{t} = [\text{cat}, \text{dog}, \text{ship}, \text{house}, \text{car}] = [1, 1, 0, 0, 0]$$

- Normalizing the vector does not help: it is not a dog **or** a cat, it is a dog **and** a cat.

$$\mathbf{t} = [\text{cat}, \text{dog}, \text{ship}, \text{house}, \text{car}] = [0.5, 0.5, 0, 0, 0]$$



# Multi-label classification



- For multi-label classification, we can simply use the **logistic** activation function for the output neurons:

$$\mathbf{y} = \sigma(W \times \mathbf{x} + \mathbf{b})$$

- The outputs are between 0 and 1, but they do not sum to one. Each output neuron performs **logistic regression for soft classification** on their class:

$$y_j = P(\text{class} = j | \mathbf{x})$$

- Each output neuron  $y_j$  has a binary target  $t_j$  (one-vs-the-rest) and has to minimize the negative log-likelihood:

$$l_j(W, \mathbf{b}) = -t_j \log y_j + (1 - t_j) \log(1 - y_j)$$

- The **binary cross-entropy** loss is the sum of the negative log-likelihood for each class:

$$\mathcal{L}(W, \mathbf{b}) = \mathbb{E}_{\mathcal{D}} \left[ - \sum_{j=1}^C t_j \log y_j + (1 - t_j) \log(1 - y_j) \right]$$