

TECHNISCHE UNIVERSITÄT
CHEMNITZ

Neurocomputing

Restricted Boltzmann Machines


Julien Vitay

Professur für Künstliche Intelligenz - Fakultät für Informatik



<https://tu-chemnitz.de/informatik/KI/edu/neurocomputing>

Restricted Boltzmann Machines

- Auto-encoders are not the only feature extractors that can be stacked.
- **Restricted Boltzmann Machines** (RBM) are generative stochastic artificial neural networks that can learn a probability distribution of their inputs.
- Their neurons form a bipartite graph with two groups of reciprocally connected units:
 - the **visible units \mathbf{v}** (the inputs)
 - the **hidden units \mathbf{h}** (the features or latent space).
- Connections are bidirectional between \mathbf{v} and \mathbf{h} , but the neurons inside the two groups are independent from each other (*restricted*).
- The goal of learning is to find the weights allowing the network to **explain** best the input data.




Restricted Boltzmann Machines



Source : <https://www.edureka.co/blog/restricted-boltzmann-machine-tutorial/>

- RBMs are a form of autoencoder where the input \rightarrow feature weight matrix is the same as the feature \rightarrow output matrix.
- There are two steps:
 - The **forward pass** $P(\mathbf{h}|\mathbf{x})$ propagates the visible units activation to the hidden units.
 - The **backward pass** $P(\mathbf{x}|\mathbf{h})$ reconstructs the visible units from the the hidden units.
- If the weight matrix is correctly chosen, the reconstructed input should “match” the original input: the data is explained.

Restricted Boltzmann Machines



- The visible and units are generally **binary units** (0 or 1), with a probability defined by the weights and biases and the logistic function:

$$P(h_j = 1 | \mathbf{v}) = \sigma\left(\sum_i W_{ij} v_i + c_j\right)$$

$$P(v_i = 1 | \mathbf{h}) = \sigma\left(\sum_j W_{ji} h_j + b_i\right)$$

- The weight matrix W and the biases \mathbf{b}, \mathbf{c} are the parameters θ of a **probability distribution** over the activation of the visible and hidden units.

- The goal is to find the parameters which explain best the data (visible units), i.e. the ones maximizing the **log-likelihood** of the model for the data $(\mathbf{v}_1, \dots, \mathbf{v}_N)$.
- We use **maximum likelihood estimation** (MLE) to maximize the log-likelihood of the model:

$$\max_{\theta} \mathcal{L}(\theta) = \mathbb{E}_{\mathbf{v} \sim D} [\log P_{\theta}(\mathbf{v})]$$

Restricted Boltzmann Machines

- In practice, MLE is not tractable in a RBM, as we cannot estimate the joint probability $P(\mathbf{v}, \mathbf{h})$ of \mathbf{v} and \mathbf{h} (too many combinations are possible).

$$P(\mathbf{v}) = \sum_{\mathbf{h}} P(\mathbf{v}, \mathbf{h})$$

- The main trick in **energy-based models** is to rewrite the probabilities using an energy function $E(\mathbf{v}, \mathbf{h})$:

$$P(\mathbf{v}) = \sum_{\mathbf{h}} P(\mathbf{v}, \mathbf{h}) = \frac{\sum_{\mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{v}, \mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})}} = \frac{1}{Z} \sum_{\mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})}$$

where:

$$Z = \sum_{\mathbf{v}, \mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})} = \sum_{\mathbf{v}} P(\mathbf{v}) \sum_{\mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})}$$

is the **partition function** (a normalizing term).

- The probabilities come from a **Gibbs distribution** (or Boltzmann distribution) parameterized by the energy of the system.
- This is equivalent to a simple **softmax** over the energy...

Restricted Boltzmann Machines

- Having reformulated the probabilities in terms of energy:

$$P(\mathbf{v}) = \frac{1}{Z} \sum_{\mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})}$$

we can introduce the **free energy** of the model for a sample \mathbf{v} (how surprising is the input \mathbf{v} for the model):

$$\mathcal{F}(\mathbf{v}) = -\log \sum_{\mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})}$$

- The log-likelihood of the model for a sample \mathbf{v} of the training data $(\mathbf{v}_1, \dots, \mathbf{v}_N)$ becomes:

$$\log P(\mathbf{v}) = \log \frac{1}{Z} \sum_{\mathbf{h}} \exp^{-E(\mathbf{v}, \mathbf{h})} = -\mathcal{F}(\mathbf{v}) + \log Z = -\mathcal{F}(\mathbf{v}) + \sum_{\mathbf{v}} P(\mathbf{v}) \mathcal{F}(\mathbf{v})$$

- Note that the second term sums over all possible inputs \mathbf{v} .
- Maximizing the log-likelihood of the model on the training data can be done using gradient ascent by following this gradient:

$$\nabla_{\theta} \mathcal{L}(\theta) = \mathbb{E}_{\mathbf{v}}[\nabla_{\theta} \log P(\mathbf{v}_i)] = \mathbb{E}_{\mathbf{v}}[-\nabla_{\theta} \mathcal{F}(\mathbf{v}) + \sum_{\mathbf{v}} P(\mathbf{v}) \nabla_{\theta} \mathcal{F}(\mathbf{v})]$$

Restricted Boltzmann Machines

- The free energy for a RBM with binary neurons is fortunately known analytically:

$$\mathcal{F}(\mathbf{v}) = - \sum_i b_i v_i - \sum_j \log(1 + \exp^{\sum_i W_{ij} v_i + c_j})$$

so finding the gradient w.r.t $\theta = (W, \mathbf{b}, \mathbf{c})$ of the first term on the r.h.s (the free energy of the sample) is easy:


$$\nabla_{\theta} \log P(\mathbf{v}) = -\nabla_{\theta} \mathcal{F}(\mathbf{v}) + \sum_{\mathbf{v}} P(\mathbf{v}) \nabla_{\theta} \mathcal{F}(\mathbf{v})$$

- In particular, the gradient w.r.t the matrix W is the outer product between \mathbf{v} and $P(\mathbf{h}|\mathbf{v})$:

$$\nabla_W \mathcal{F}(\mathbf{v}) = -\mathbf{v} \times P(\mathbf{h}|\mathbf{v})$$

- The problem is the second term: we would need to integrate over all possible values of the inputs \mathbf{v} , what is not tractable.
- We will therefore make an approximation using **Gibbs sampling** (a variant of **Monte-Carlo Markov Chain** sampling - MCMC) to estimate that second term.

Gibbs sampling



Source : <https://towardsdatascience.com/deep-learning-meets-physics-restricted-boltzmann-machines-part-i-6df5c4918c15>

- Gibbs sampling consists of repeatedly applying the encoder $P(\mathbf{h}|\mathbf{v})$ and the decoder $P(\mathbf{v}|\mathbf{h})$ on the input.
 - We start by setting $\mathbf{v}_0 = \mathbf{v}$ using a training sample.
 - We obtain \mathbf{h}_0 by computing $P(\mathbf{h}|\mathbf{v}_0)$ and sampling it.
 - We obtain \mathbf{v}_1 by computing $P(\mathbf{v}|\mathbf{h}_0)$ and sampling it.
 - ...
 - We obtain \mathbf{v}_k by computing $P(\mathbf{v}|\mathbf{h}_{k-1})$ and sampling it.
- After enough iterations k , we should have a good estimate of $P(\mathbf{v}, \mathbf{h})$.
- The k iterations have generated enough **reconstructions** of \mathbf{v} to cover the distribution of \mathbf{v} .

Contrastive divergence


- We set $\mathbf{v}_0 = \mathbf{v}$ on a training sample and let Gibbs sampling iterate for k iterations until we obtain $\mathbf{v}_k = \mathbf{v}^*$.
- **Contrastive divergence** (CD- k) shows that the gradient of the log-likelihood can be approximated by:

$$\begin{aligned}\nabla_W \log P(\mathbf{v}) &= -\nabla_W \mathcal{F}(\mathbf{v}) + \sum_{\mathbf{v}} P(\mathbf{v}) \nabla_W \mathcal{F}(\mathbf{v}) \\ &\approx \mathbf{v} \times P(\mathbf{h}|\mathbf{v}) - \mathbf{v}^* \times P(\mathbf{h}|\mathbf{v}^*)\end{aligned}$$

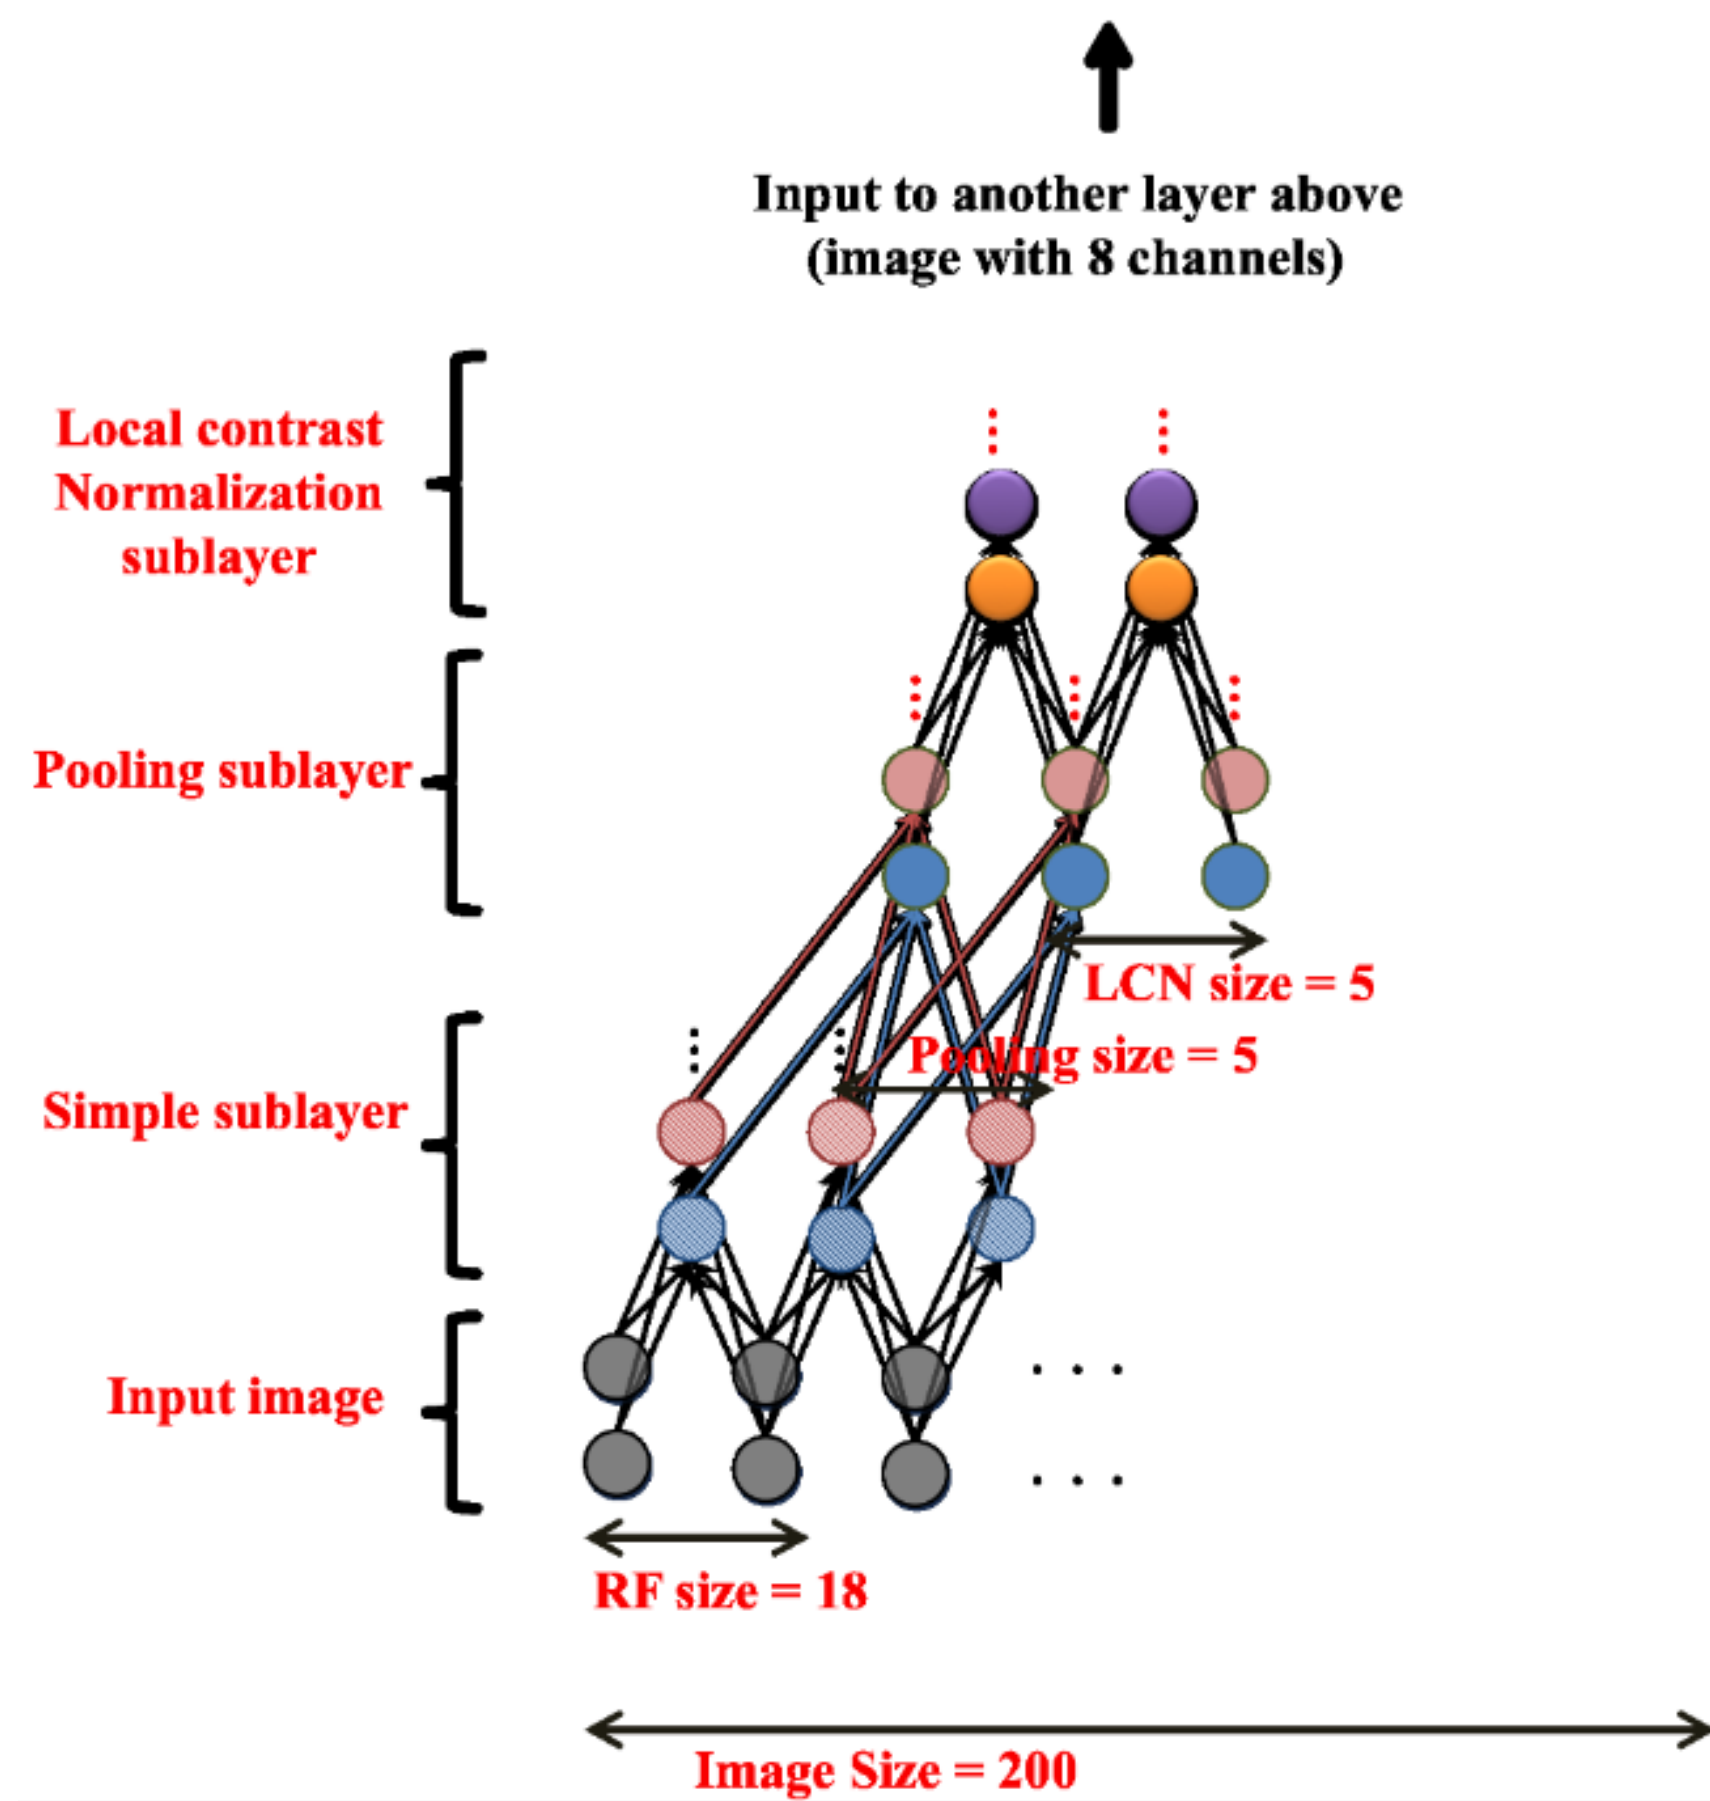
- The gradient of the log-likelihood is the difference between the initial explanation of \mathbf{v} by the model, and its explanation after k iterations (relaxation).
- If the model is good, the reconstruction \mathbf{v}^* is the same as the input \mathbf{v} , so the gradient is zero.
- An input \mathbf{v} is likely under the RBM model if it is able to reconstruct it, i.e. when it is not surprising (the free energy is low).
- In practice, $k = 1$ gives surprisingly good results, but RBMs are very painful to train (hyperparameters)...

Deep Belief Networks = stacked RBMs

- A **Deep Belief Network** (DBM) is a simple stack of RBMs, trained using greedy layer-wise learning.
- The “bottom” parts of the DBM become unidirectional when learning the top part.

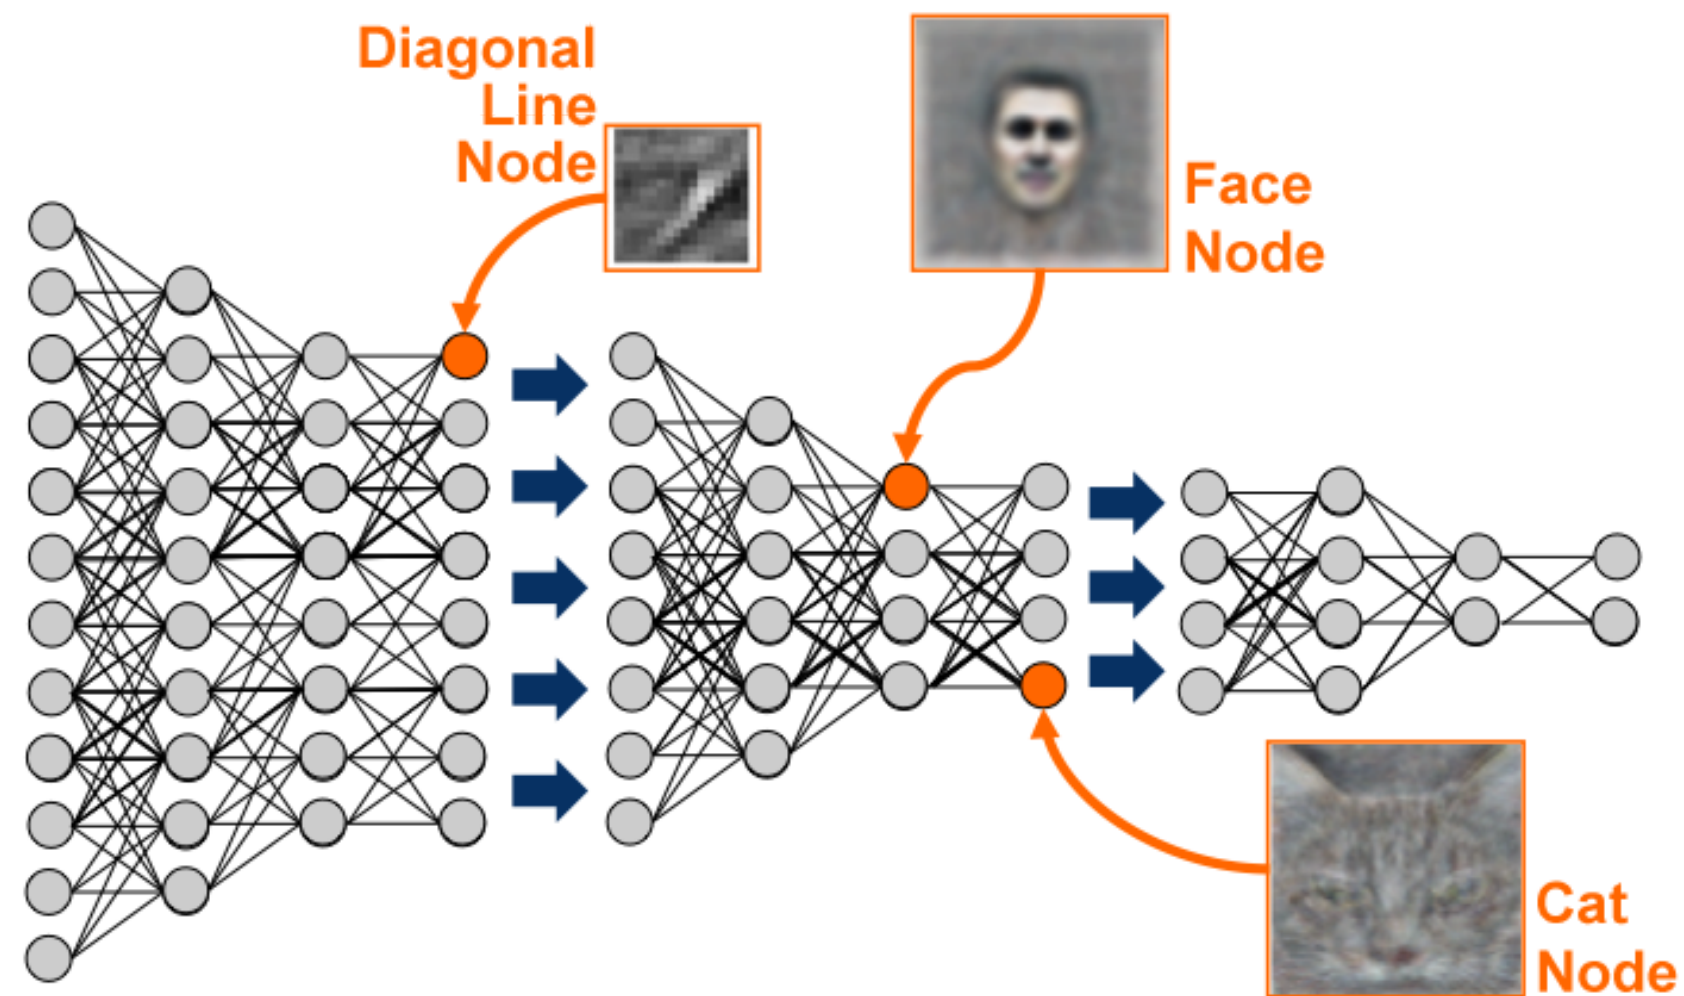


Application: Finding cats on the internet



- Andrew Ng and colleagues (Google, Stanford) used a similar technique to train a deep belief network on color images (200x200) taken from 10 million random unlabeled Youtube videos.
- Each layer was trained greedily. They used a particular form of autoencoder called **restricted Boltzmann machines** (RBM) and a couple of other tricks (receptive fields, contrast normalization).
- Training was distributed over 1000 machines (16.000 cores) and lasted for three days.
- There was absolutely no task: the network just had to watch youtube videos.
- After learning, they visualized what the neurons had learned.

Application: Finding cats on the internet



- After training, some neurons had learned to respond uniquely to faces, or to cats, without ever having been instructed to.
- The network can then be fine-tuned for classification tasks, improving the pre-AlexNet state-of-the-art on ImageNet by 70%.