



UNIVERSITY OF TECHNOLOGY  
IN THE EUROPEAN CAPITAL OF CULTURE  
CHEMNITZ

# Neurocomputing

Attentional neural networks


Julien Vitay

Professur für Künstliche Intelligenz - Fakultät für Informatik

<https://tu-chemnitz.de/informatik/KI/edu/neurocomputing>

# Attentional recurrent networks


- The problem with seq2seq is that it **compresses** the complete input sentence into a single state vector.



- For long sequences, the beginning of the sentence may not be present in the final state vector:
  - Truncated BPTT, vanishing gradients.
  - When predicting the last word, the beginning of the paragraph might not be necessary.
- Consequence: there is not enough information in the state vector to start translating.

# Attentional recurrent networks

- A solution would be to concatenate the **state vectors** of all steps of the encoder and pass them to the decoder.



- **Problem 1:** it would make a lot of elements in the state vector of the decoder (which should be constant).
- **Problem 2:** the state vector of the decoder would depend on the length of the input sequence.

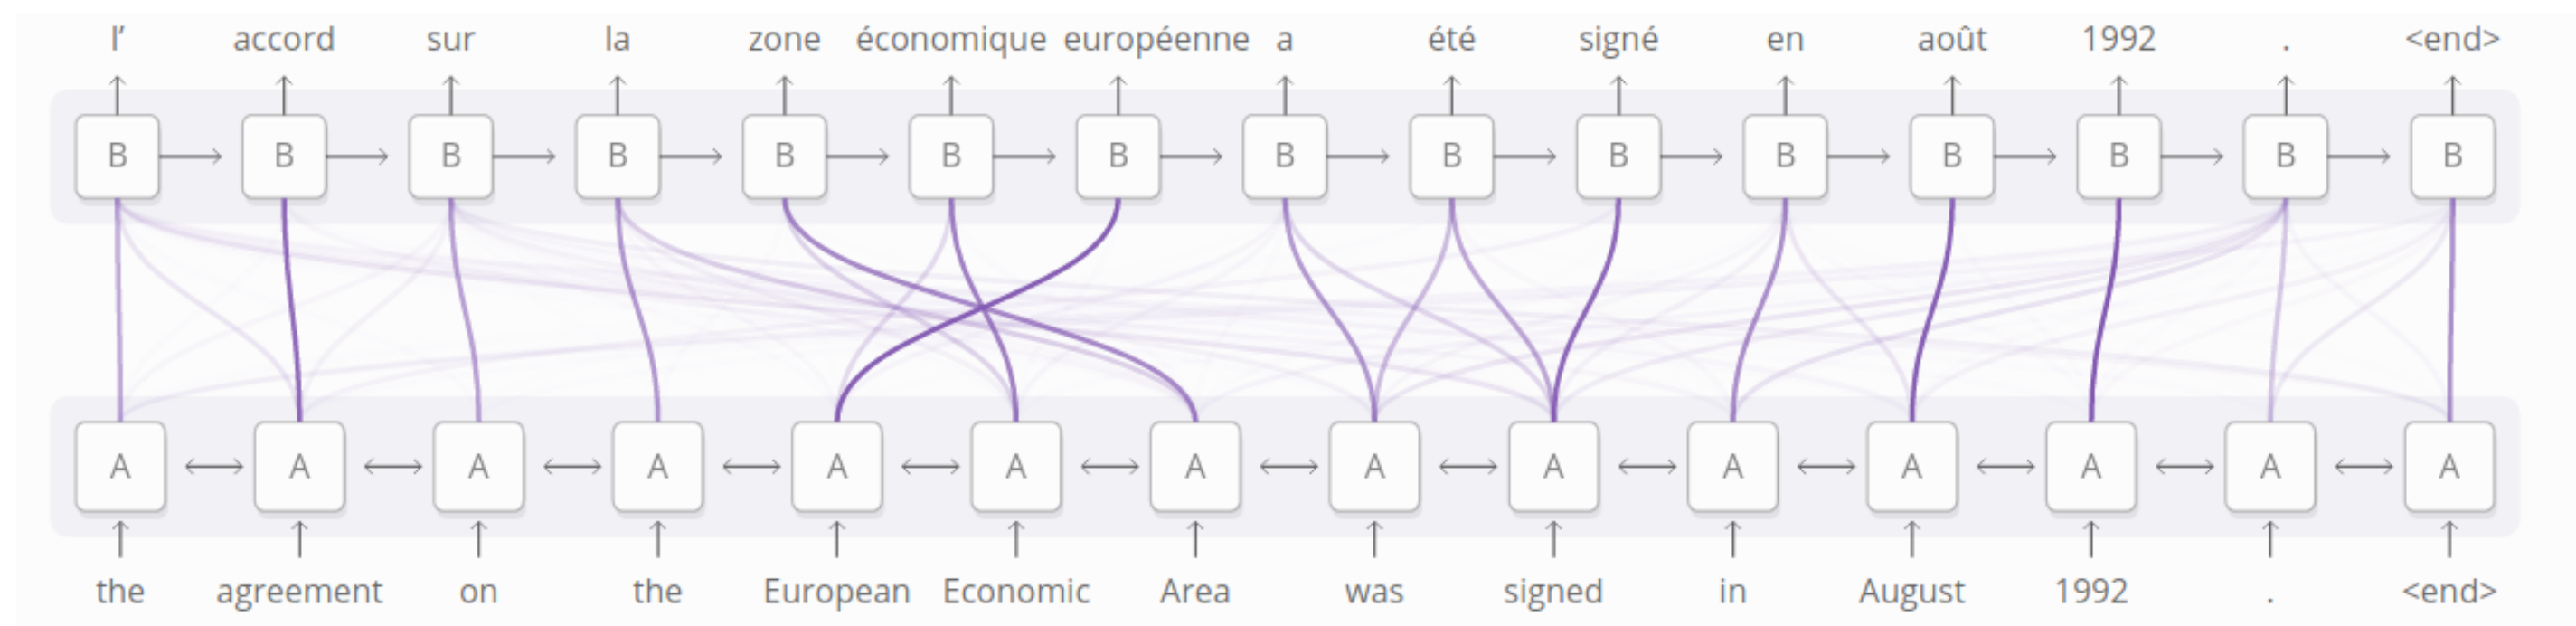
# Attentional recurrent networks

- Attentional mechanisms let the decoder decide (by learning) which state vectors it needs to generate each word at each step.

- The **attentional context vector** of the decoder  $A_t^{\text{decoder}}$  at time  $t$  is a weighted average of all state vectors  $C_i^{\text{encoder}}$  of the encoder.

$$A_t^{\text{decoder}} = \sum_{i=0}^T a_i C_i^{\text{encoder}}$$

- The coefficients  $a_i$  are called the **attention scores** : how much attention is the decoder paying to each of the encoder's state vectors?



# Attentional recurrent networks

- The attention scores  $a_i$  are computed as a **softmax** over the scores  $e_i$  (in order to sum to 1):

$$a_i = \frac{\exp e_i}{\sum_j \exp e_j} \Rightarrow A_t^{\text{decoder}} = \sum_{i=0}^T \frac{\exp e_i}{\sum_j \exp e_j} C_i^{\text{encoder}}$$

- The score  $e_i$  is computed using:
  - the previous output of the decoder  $\mathbf{h}_{t-1}^{\text{decoder}}$ .
  - the corresponding state vector  $C_i^{\text{encoder}}$  of the encoder at step  $i$ .
  - attentional weights  $W_a$ .


$$e_i = \tanh(W_a [\mathbf{h}_{t-1}^{\text{decoder}}; C_i^{\text{encoder}}])$$

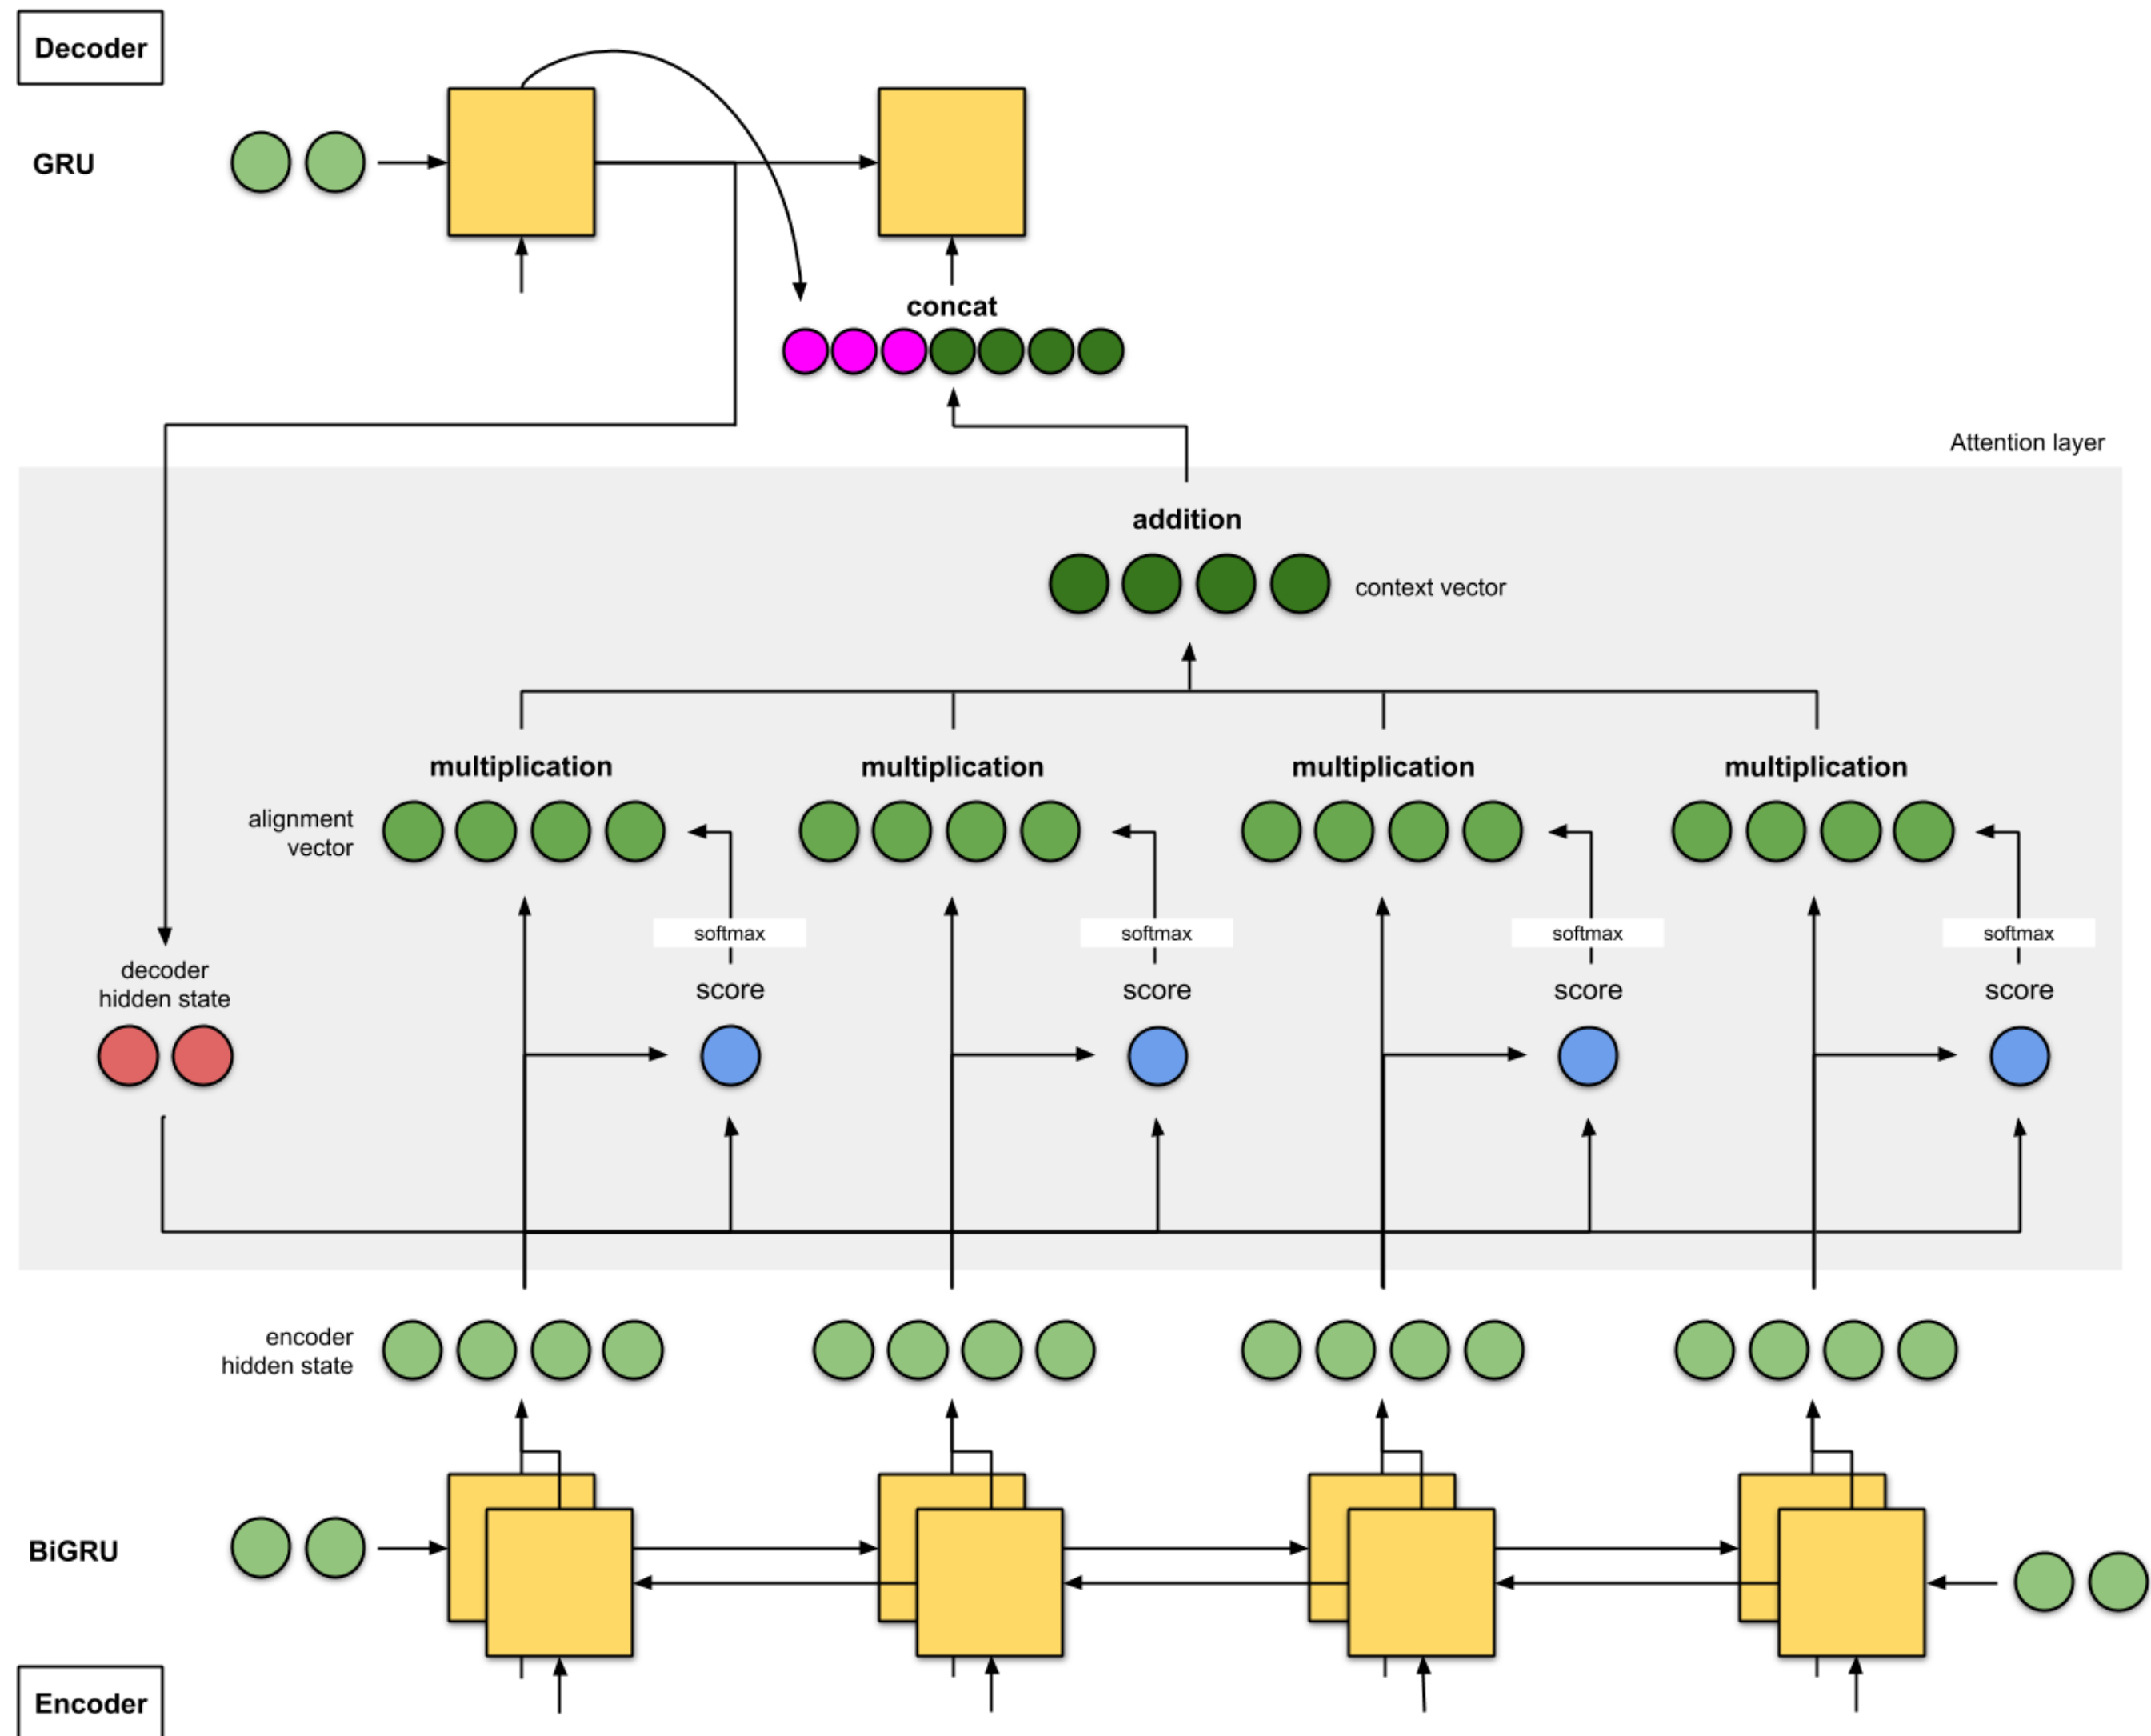
- Everything is differentiable, these attentional weights can be learned with BPTT.

# Attentional recurrent networks

- The attentional context vector  $A_t^{\text{decoder}}$  is concatenated with the previous output  $\mathbf{h}_{t-1}^{\text{decoder}}$  and used as the next input  $\mathbf{x}_t^{\text{decoder}}$  of the decoder:


$$\mathbf{x}_t^{\text{decoder}} = [\mathbf{h}_{t-1}^{\text{decoder}}; A_t^{\text{decoder}}]$$





# Attentional recurrent networks

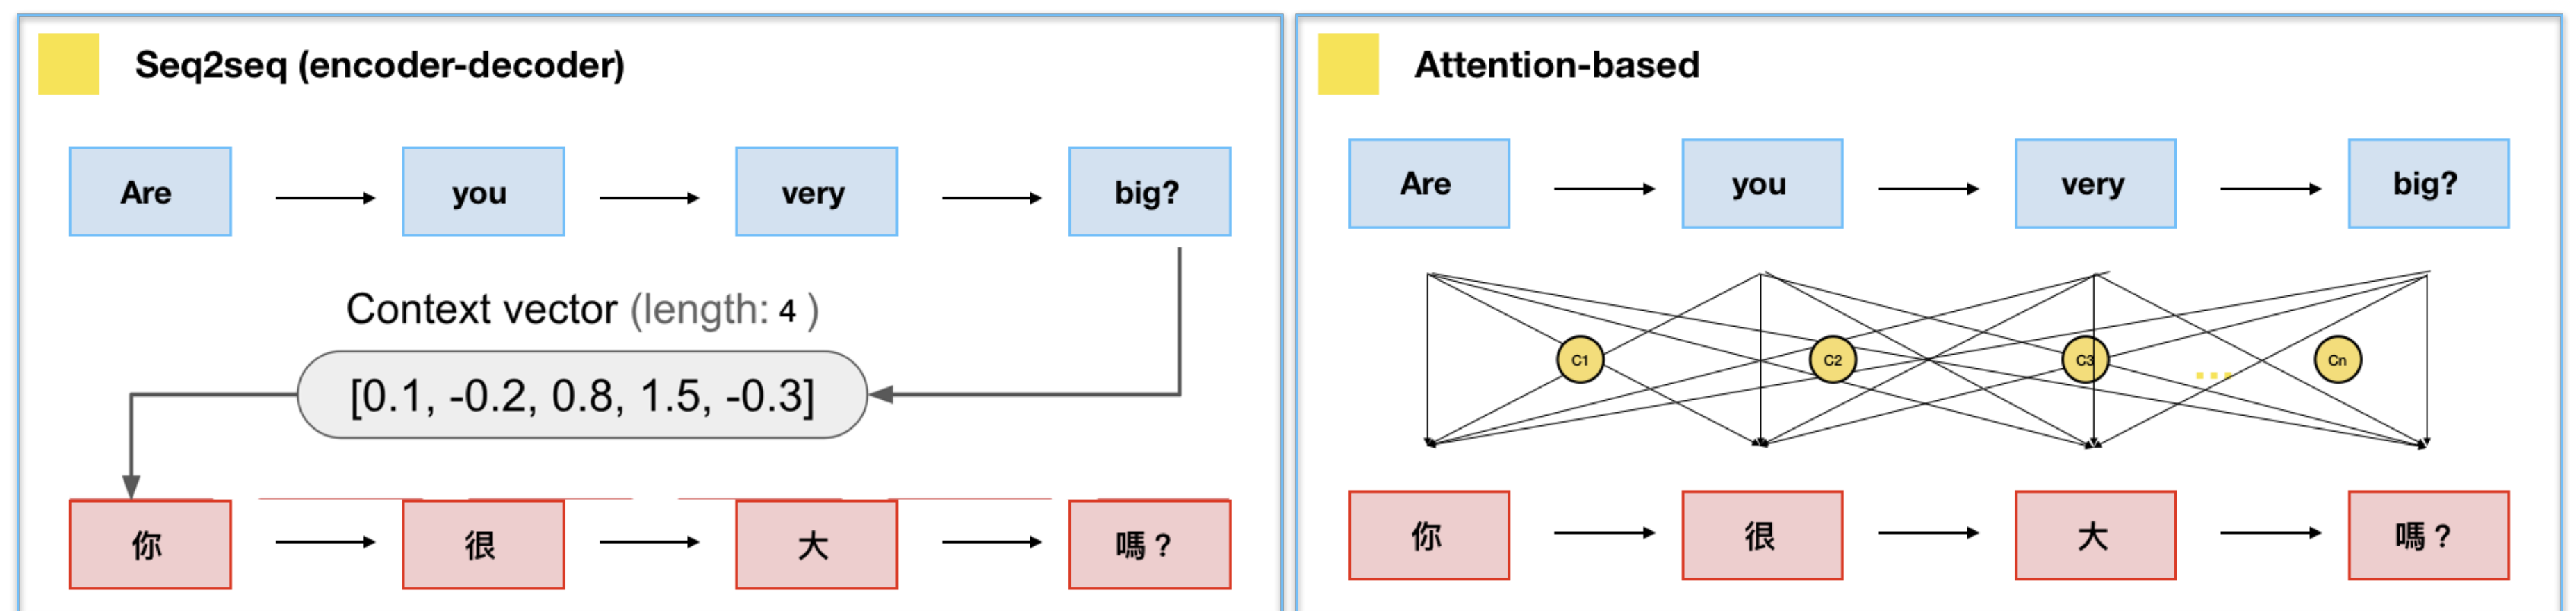
- The attention scores or **alignment scores**  $a_i$  are useful to interpret what happened.
- They show which words in the original sentence are the most important to generate the next word.





# Attentional recurrent networks

- **Attentional mechanisms** are now central to NLP.



- The whole **history** of encoder states is passed to the decoder, which learns to decide which part is the most important using **attention**.
- This solves the bottleneck of seq2seq architectures, at the cost of much more operations.
- They require to use fixed-length sequences (generally 50 words).

# Google Neural Machine Translation (GNMT)

- Google Neural Machine Translation (GNMT) uses an attentional recurrent NN, with bidirectional GRUs, 8 recurrent layers on 8 GPUs for both the encoder and decoder.

