

TECHNISCHE UNIVERSITÄT  
CHEMNITZ

# Neurocomputing

Basics in mathematics

Julien Vitay

Professur für Künstliche Intelligenz - Fakultät für Informatik

<https://tu-chemnitz.de/informatik/KI/edu/neurocomputing>

# Outline

1. Linear algebra
2. Calculus
3. Probability theory
4. Statistics
5. Information theory

# 1 - Linear algebra

# Mathematical objects

- **Scalars**  $x$  are 0-dimensional values. They can either take real values ( $x \in \mathfrak{R}$ , e.g.  $x = 1.4573$ , floats in CS) or natural values ( $x \in \mathbb{N}$ , e.g.  $x = 3$ , integers in CS).
- **Vectors**  $\mathbf{x}$  are 1-dimensional arrays of length  $d$ .
- The bold notation  $\mathbf{x}$  will be used in this course, but you may also be accustomed to the arrow notation  $\vec{x}$  used on the blackboard. When using real numbers, the **vector space** with  $d$  dimensions is noted  $\mathfrak{R}^d$ , so we can note  $\mathbf{x} \in \mathfrak{R}^d$ .
- Vectors are typically represented vertically to outline their  $d$  elements  $x_1, x_2, \dots, x_d$ :

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$

# Mathematical objects

- **Matrices**  $A$  are 2-dimensional arrays of size (or shape)  $m \times n$  ( $m$  rows,  $n$  columns,  $A \in \mathfrak{R}^{m \times n}$ ).
- They are represented by a capital letter to distinguish them from scalars (classically also in bold  $\mathbf{A}$  but not here). The element  $a_{ij}$  of a matrix  $A$  is the element on the  $i$ -th row and  $j$ -th column.

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

- **Tensors**  $\mathcal{A}$  are arrays with more than two dimensions. We will not really do math on these objects, but they are useful internally (hence the name of the [tensorflow](#) library).


# Vectors

- A vector can be thought of as the **coordinates of a point** in an Euclidean space (such the 2D space), relative to the origin.
- A vector space relies on two fundamental operations, which are that:
- Vectors can be added:

$$\mathbf{x} + \mathbf{y} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} + \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_d \end{bmatrix} = \begin{bmatrix} x_1 + y_1 \\ x_2 + y_2 \\ \vdots \\ x_d + y_d \end{bmatrix}$$

- Vectors can be multiplied by a scalar:

$$a \mathbf{x} = a \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} = \begin{bmatrix} a x_1 \\ a x_2 \\ \vdots \\ a x_d \end{bmatrix}$$



Source: [https://mathinsight.org/image/vector\\_2d\\_add](https://mathinsight.org/image/vector_2d_add)

# Properties of vector spaces

- These two operations generate a lot of nice properties (see [https://en.wikipedia.org/wiki/Vector\\_space](https://en.wikipedia.org/wiki/Vector_space) for a full list), including:

- associativity:

$$\mathbf{x} + (\mathbf{y} + \mathbf{z}) = (\mathbf{x} + \mathbf{y}) + \mathbf{z}$$

- commutativity:

$$\mathbf{x} + \mathbf{y} = \mathbf{y} + \mathbf{x}$$

- the existence of a zero vector

$$\mathbf{x} + \mathbf{0} = \mathbf{x}$$


- inversion:

$$\mathbf{x} + (-\mathbf{x}) = \mathbf{0}$$

- distributivity:


$$a(\mathbf{x} + \mathbf{y}) = a\mathbf{x} + a\mathbf{y}$$

# Norm of a vector




- Vectors have a **norm** (or length)  $||\mathbf{x}||$ . The most intuitive one (if you know the Pythagoras theorem) is the **Euclidean norm** or  $L^2$ -norm, which sums the square of each element:

$$||\mathbf{x}||_2 = \sqrt{x_1^2 + x_2^2 + \dots + x_d^2}$$



- Other norms exist, distinguished by the subscript. The  $L^1$ -**norm** (also called the Manhattan norm) sums the absolute value of each element:

$$||\mathbf{x}||_1 = |x_1| + |x_2| + \dots + |x_d|$$



- The **p-norm** generalizes the Euclidean norm to other powers  $p$ :

$$||\mathbf{x}||_p = (|x_1|^p + |x_2|^p + \dots + |x_d|^p)^{\frac{1}{p}}$$

- The **infinity norm** (or maximum norm)  $L^\infty$  returns the maximum element of the vector:

$$||\mathbf{x}||_\infty = \max(|x_1|, |x_2|, \dots, |x_d|)$$

Source:

[https://en.wikipedia.org/wiki/Norm\\_\(mathematics\)](https://en.wikipedia.org/wiki/Norm_(mathematics))



# Dot product

- One important operation for vectors is the **dot product** (also called scalar product or inner product) between two vectors:

$$\langle \mathbf{x} \cdot \mathbf{y} \rangle = \left\langle \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \cdot \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_d \end{bmatrix} \right\rangle = x_1 y_1 + x_2 y_2 + \dots + x_d y_d$$

- The dot product basically sums one by one the product of the elements of each vector. The angular brackets are sometimes omitted ( $\mathbf{x} \cdot \mathbf{y}$ ) but we will use them in this course for clarity.
- One can notice immediately that the dot product is **symmetric**:

$$\langle \mathbf{x} \cdot \mathbf{y} \rangle = \langle \mathbf{y} \cdot \mathbf{x} \rangle$$

and **linear**:

$$\langle (a \mathbf{x} + b \mathbf{y}) \cdot \mathbf{z} \rangle = a \langle \mathbf{x} \cdot \mathbf{z} \rangle + b \langle \mathbf{y} \cdot \mathbf{z} \rangle$$


# Dot product

- The dot product is an indirect measurement of the **angle**  $\theta$  between two vectors:

$$\langle \mathbf{x} \cdot \mathbf{y} \rangle = \|\mathbf{x}\|_2 \|\mathbf{y}\|_2 \cos(\theta)$$

- If you normalize the two vectors by dividing them by their norm (which is a scalar), we indeed have the cosine of the angle between them
- The higher the normalized dot product, the more the two vectors point towards the same direction (**cosine distance** between two vectors).

$$\left\langle \frac{\mathbf{x}}{\|\mathbf{x}\|_2} \cdot \frac{\mathbf{y}}{\|\mathbf{y}\|_2} \right\rangle = \cos(\theta)$$



Source: [https://mathinsight.org/image/dot\\_product\\_projection\\_unit\\_vector](https://mathinsight.org/image/dot_product_projection_unit_vector)

# Matrices

- Matrices are derived from vectors, so most of the previous properties will be true. Let's consider this 4x3 matrix:

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \\ a_{41} & a_{42} & a_{43} \end{bmatrix}$$

- Each column of the matrix is a vector with 4 elements:

$$\mathbf{a}_1 = \begin{bmatrix} a_{11} \\ a_{21} \\ a_{31} \\ a_{41} \end{bmatrix} \quad \mathbf{a}_2 = \begin{bmatrix} a_{12} \\ a_{22} \\ a_{32} \\ a_{42} \end{bmatrix} \quad \mathbf{a}_3 = \begin{bmatrix} a_{13} \\ a_{23} \\ a_{33} \\ a_{43} \end{bmatrix}$$

- A  $m \times n$  matrix is therefore a collection of  $n$  vectors of size  $m$  put side by side column-wise:

$$A = \begin{bmatrix} \mathbf{a}_1 & \mathbf{a}_2 & \mathbf{a}_3 \end{bmatrix}$$

## Properties of matrix spaces

- All properties of the vector spaces (associativity, commutativity, distributivity) also apply to matrices, as additions and multiplications with a scalar are defined.

$$\alpha A + \beta B = \begin{bmatrix} \alpha a_{11} + \beta b_{11} & \alpha a_{12} + \beta b_{12} & \alpha a_{13} + \beta b_{13} \\ \alpha a_{21} + \beta b_{21} & \alpha a_{22} + \beta b_{22} & \alpha a_{23} + \beta b_{23} \\ \alpha a_{31} + \beta b_{31} & \alpha a_{32} + \beta b_{32} & \alpha a_{33} + \beta b_{33} \\ \alpha a_{41} + \beta b_{41} & \alpha a_{42} + \beta b_{42} & \alpha a_{43} + \beta b_{43} \end{bmatrix}$$

**Note:** Beware, you can only add matrices of the same dimensions  $m \times n$ . You cannot add a  $2 \times 3$  matrix to a  $5 \times 4$  one.

# Transposition

- The **transpose**  $A^T$  of a  $m \times n$  matrix  $A$  is a  $n \times m$  matrix, where the row and column indices are swapped:

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}, \quad A^T = \begin{bmatrix} a_{11} & a_{21} & \cdots & a_{m1} \\ a_{12} & a_{22} & \cdots & a_{m2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \cdots & a_{mn} \end{bmatrix}$$

- This is also true for vectors, which become horizontal after transposition:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}, \quad \mathbf{x}^T = [x_1 \quad x_2 \quad \dots \quad x_d]$$

# Matrix multiplication

- If  $A$  is a  $m \times n$  matrix and  $B$  a  $n \times p$  matrix:

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}, \quad B = \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1p} \\ b_{21} & b_{22} & \cdots & b_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{np} \end{bmatrix}$$

we can multiply them to obtain a  $m \times p$  matrix:

$$C = A \times B = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1p} \\ c_{21} & c_{22} & \cdots & c_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mp} \end{bmatrix}$$

where each element  $c_{ij}$  is the dot product of the  $i$ th row of  $A$  and  $j$ th column of  $B$ :

$$c_{ij} = \langle A_{i,:} \cdot B_{:,j} \rangle = a_{i1}b_{1j} + a_{i2}b_{2j} + \cdots + a_{in}b_{nj} = \sum_{k=1}^n a_{ik}b_{kj}$$

**Note:**  $n$ , the number of columns of  $A$  and rows of  $B$ , must be the same!

# Matrix multiplication

- The element  $c_{ij}$  of  $C = A \times B$  is the dot product between the  $i$ th row of  $A$  and the  $j$ th column of  $B$ .

$$c_{11} = a_{11}b_{11} + a_{12}b_{21} + a_{13}b_{31} + a_{14}b_{41}$$

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & b_{13} \\ b_{21} & b_{22} & b_{23} \\ b_{31} & b_{32} & b_{33} \\ b_{41} & b_{42} & b_{43} \end{bmatrix} = \begin{bmatrix} c_{11} & c_{12} & c_{13} \\ c_{21} & c_{22} & c_{23} \end{bmatrix}$$

$2 \times 4 \qquad \qquad 4 \times 3 \qquad \qquad 2 \times 3$

$$c_{22} = a_{21}b_{12} + a_{22}b_{22} + a_{23}b_{32} + a_{24}b_{42}$$

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & b_{13} \\ b_{21} & b_{22} & b_{23} \\ b_{31} & b_{32} & b_{33} \\ b_{41} & b_{42} & b_{43} \end{bmatrix} = \begin{bmatrix} c_{11} & c_{12} & c_{13} \\ c_{21} & c_{22} & c_{23} \end{bmatrix}$$

Source:

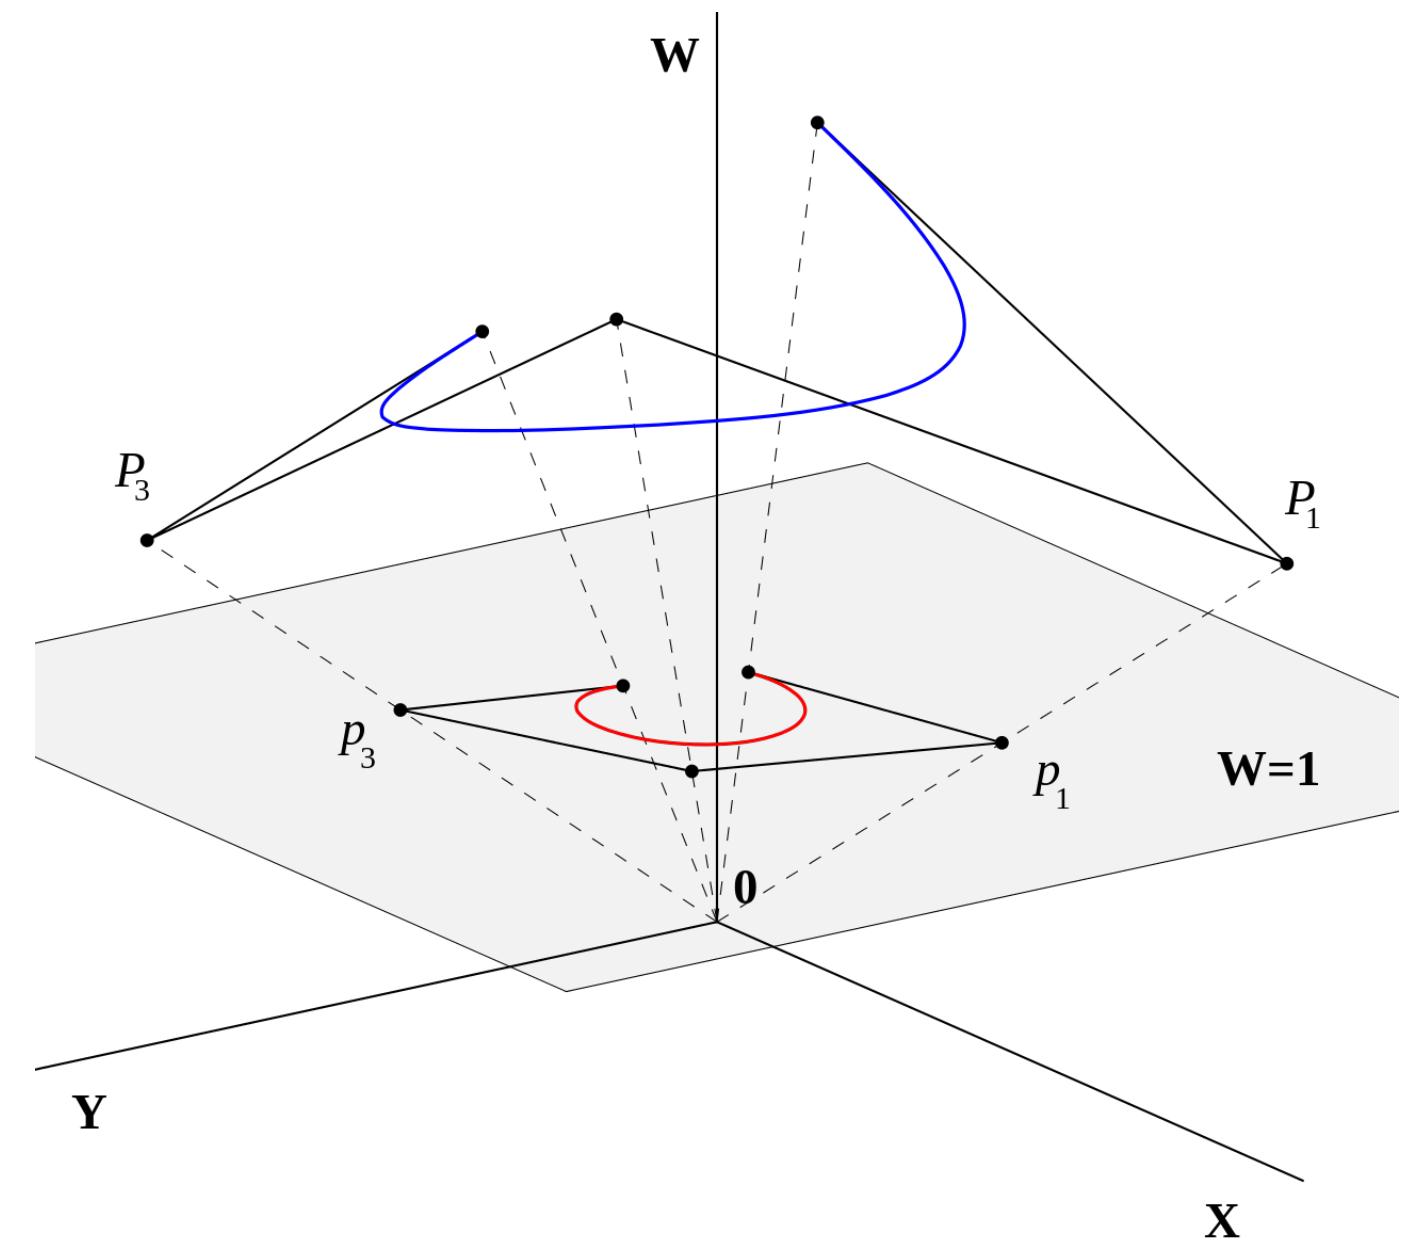
[https://chem.libretexts.org/Bookshelves/Physical\\_and\\_Theoretical\\_Chemistry\\_Textbook\\_Maps/Book%3A\\_Mathematical\\_Methods\\_in\\_Chemistry\\_\(Levitus\)/15%3A\\_Matrices](https://chem.libretexts.org/Bookshelves/Physical_and_Theoretical_Chemistry_Textbook_Maps/Book%3A_Mathematical_Methods_in_Chemistry_(Levitus)/15%3A_Matrices)  
CC BY-NC-SA; Marcia Levitus

# Matrix-vector multiplication

- Thinking of vectors as  $n \times 1$  matrices, we can multiply a matrix  $m \times n$  with a vector:

$$\mathbf{y} = \mathbf{A} \times \mathbf{x} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

- The result  $\mathbf{y}$  is a vector of size  $m$ .
- In that sense, a matrix  $A$  can transform a vector of size  $n$  into a vector of size  $m$ :
  - $A$  represents a **projection** from  $\mathfrak{R}^n$  to  $\mathfrak{R}^m$ .



Source:

[https://en.wikipedia.org/wiki/Homogeneous\\_coordinate](https://en.wikipedia.org/wiki/Homogeneous_coordinate)



## Dot product

- Note that the **dot product** between two vectors of size  $n$  is the matrix multiplication between the transpose of the first vector and the second one:

$$\mathbf{x}^T \times \mathbf{y} = \begin{bmatrix} x_1 & x_2 & \dots & x_n \end{bmatrix} \times \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = x_1 y_1 + x_2 y_2 + \dots + x_n y_n = \langle \mathbf{x} \cdot \mathbf{y} \rangle$$

# Matrix inversion

- Square matrices of size  $n \times n$  can be inverted.
- The **inverse**  $A^{-1}$  of a matrix  $A$  is defined by:

$$A \times A^{-1} = A^{-1} \times A = I$$

where  $I$  is the identity matrix (a matrix with ones on the diagonal and 0 otherwise).

- Matrix inversion allows to solve linear systems of equations. Given the problem:

$$\begin{cases} a_{11} x_1 + a_{12} x_2 + \dots + a_{1n} x_n = b_1 \\ a_{21} x_1 + a_{22} x_2 + \dots + a_{2n} x_n = b_2 \\ \dots \\ a_{n1} x_1 + a_{n2} x_2 + \dots + a_{nn} x_n = b_n \end{cases}$$

which is equivalent to:

$$A \times \mathbf{x} = \mathbf{b}$$

- We can multiply both sides to the left with  $A^{-1}$  (if it exists) and obtain:

$$\mathbf{x} = A^{-1} \times \mathbf{b}$$


## 2 - Calculus

# Univariate functions

- A **univariate function**  $f$  associates to any real number  $x \in \mathfrak{R}$  (or a subset of  $\mathfrak{R}$  called the support of the function) another (unique) real number  $f(x)$ :

$$\begin{aligned} f: \quad \mathfrak{R} &\rightarrow \mathfrak{R} \\ x &\mapsto f(x), \end{aligned}$$

$$f(x) = x^2 - 2x + 1$$




# Multivariate functions

- A **multivariate function**  $f$  associates to any vector  $\mathbf{x} \in \mathfrak{R}^n$  (or a subset) a real number  $f(\mathbf{x})$ :

$$\begin{aligned} f: \quad \mathfrak{R}^n &\rightarrow \mathfrak{R} \\ \mathbf{x} &\mapsto f(\mathbf{x}), \end{aligned}$$

- The variables of the function are the elements of the vector.
- For low-dimensional vector spaces, it is possible to represent each element explicitly, for example:

$$\begin{aligned} f: \quad \mathfrak{R}^3 &\rightarrow \mathfrak{R} \\ x, y, z &\mapsto f(x, y, z), \end{aligned}$$




Source: [https://en.wikipedia.org/wiki/Function\\_of\\_several\\_real\\_variables](https://en.wikipedia.org/wiki/Function_of_several_real_variables)

# Vector fields

- **Vector fields** associate to any vector  $\mathbf{x} \in \mathfrak{R}^n$  (or a subset) another vector (possibly of different size):

$$\begin{aligned}\vec{f}: \quad \mathfrak{R}^n &\rightarrow \mathfrak{R}^m \\ \mathbf{x} &\mapsto \vec{f}(\mathbf{x}),\end{aligned}$$



Source: [https://en.wikipedia.org/wiki/Vector\\_field](https://en.wikipedia.org/wiki/Vector_field)

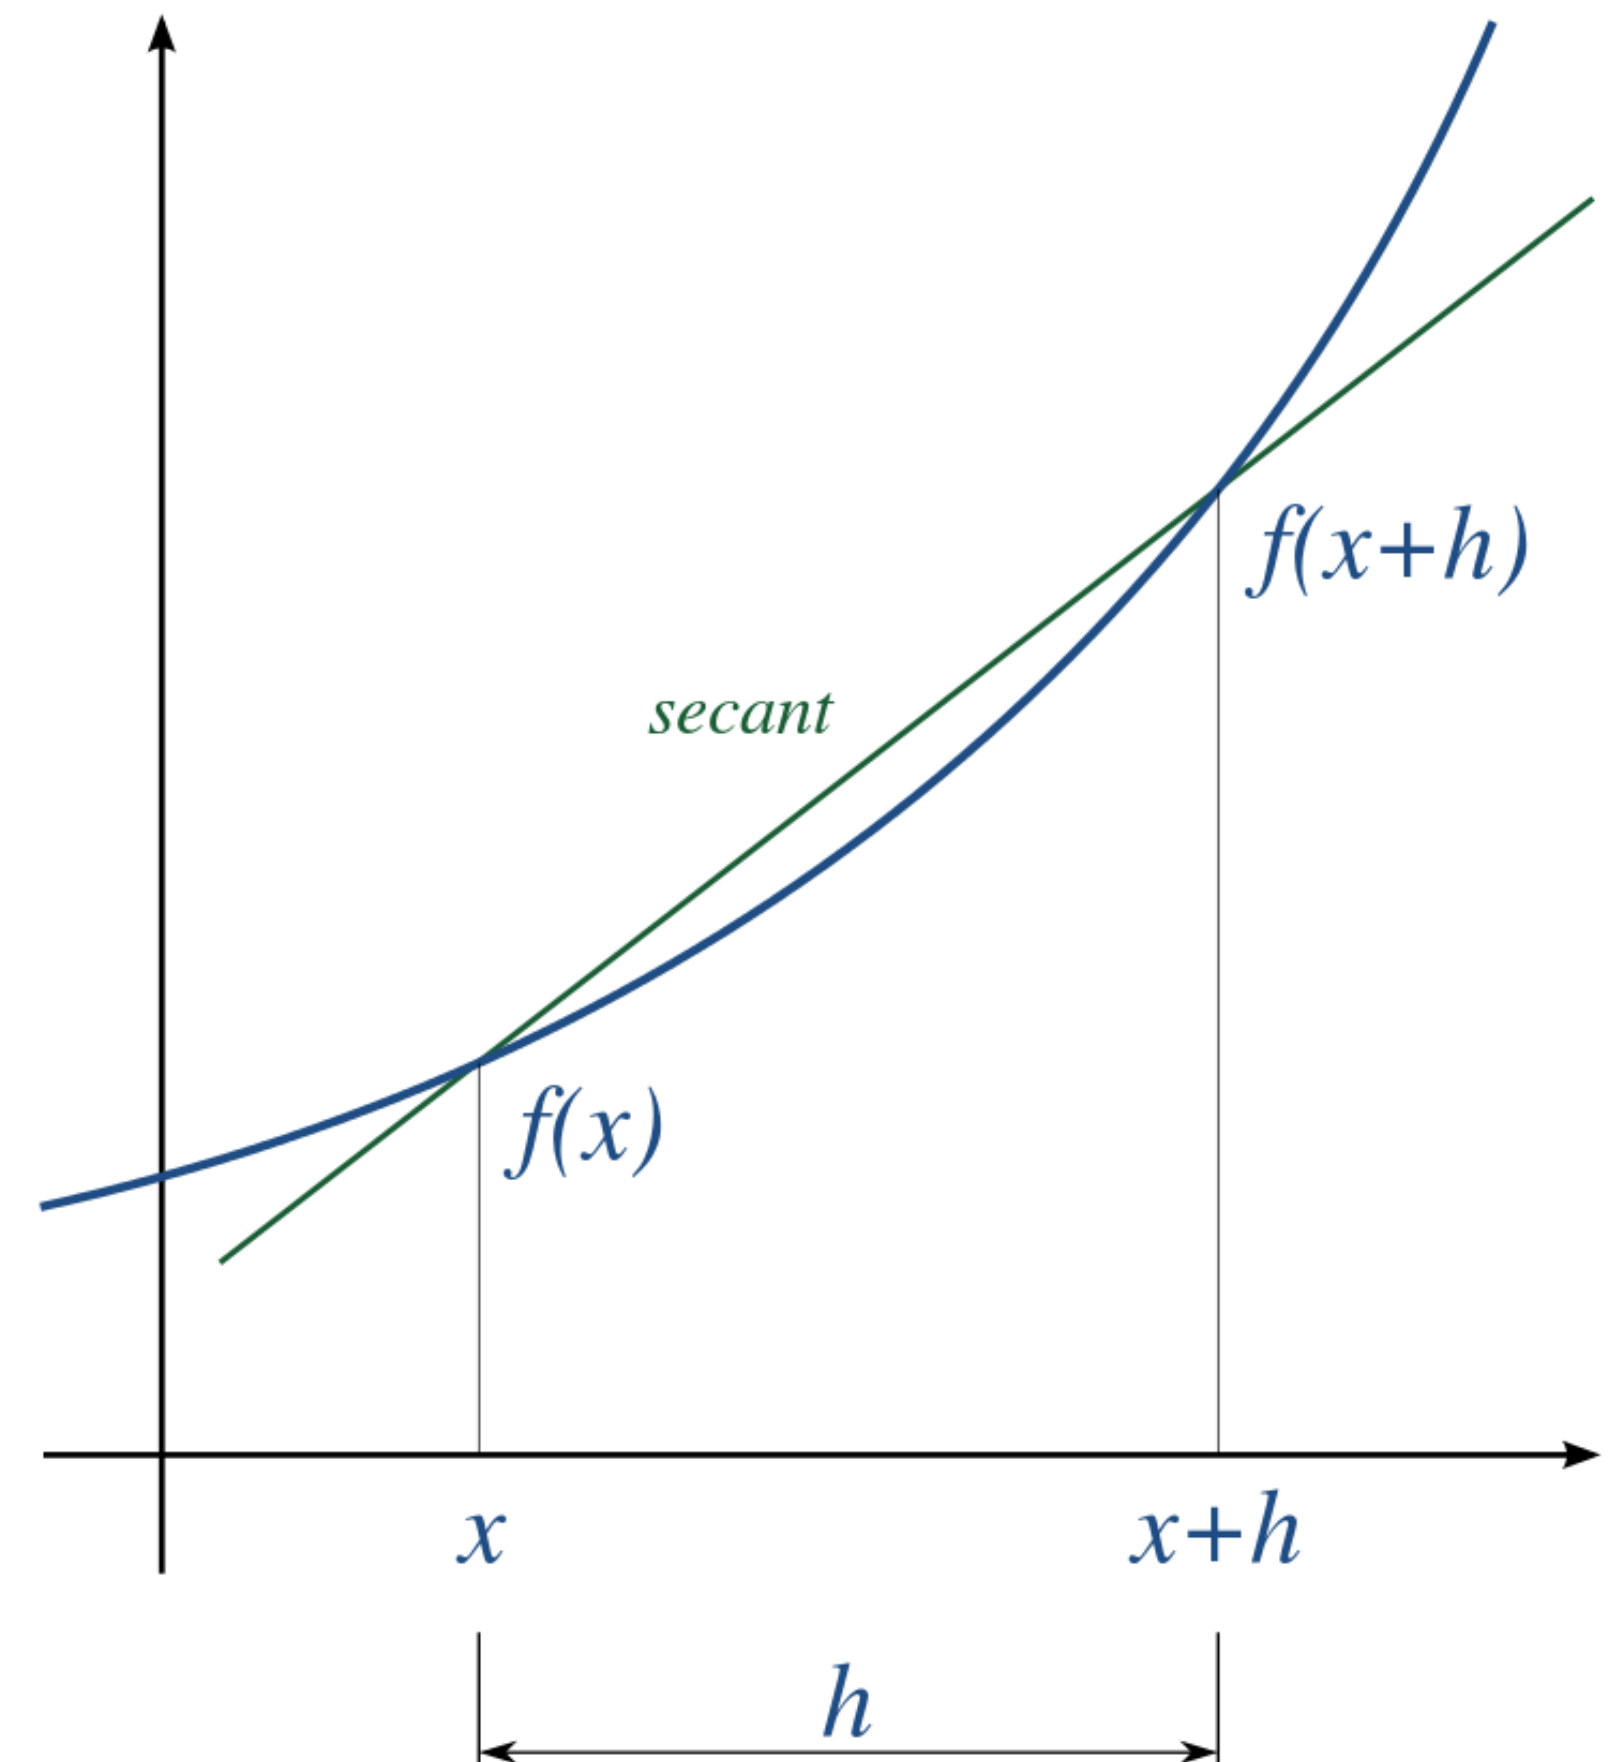
**Note:** The matrix-vector multiplication  $\mathbf{y} = A \times \mathbf{x}$  is a linear vector field, mapping any vector  $\mathbf{x}$  into another vector  $\mathbf{y}$ .

# Differentiation

- Differential calculus deals with the **derivative** of a function, a process called differentiation.
- The derivative  $f'(x)$  or  $\frac{df(x)}{dx}$  of a univariate function  $f(x)$  is defined as the local *slope* of the tangent to the function for a given value of  $x$ :


$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- The line passing through the points  $(x, f(x))$  and  $(x+h, f(x+h))$  becomes tangent to the function when  $h$  becomes very small.



# Differentiation

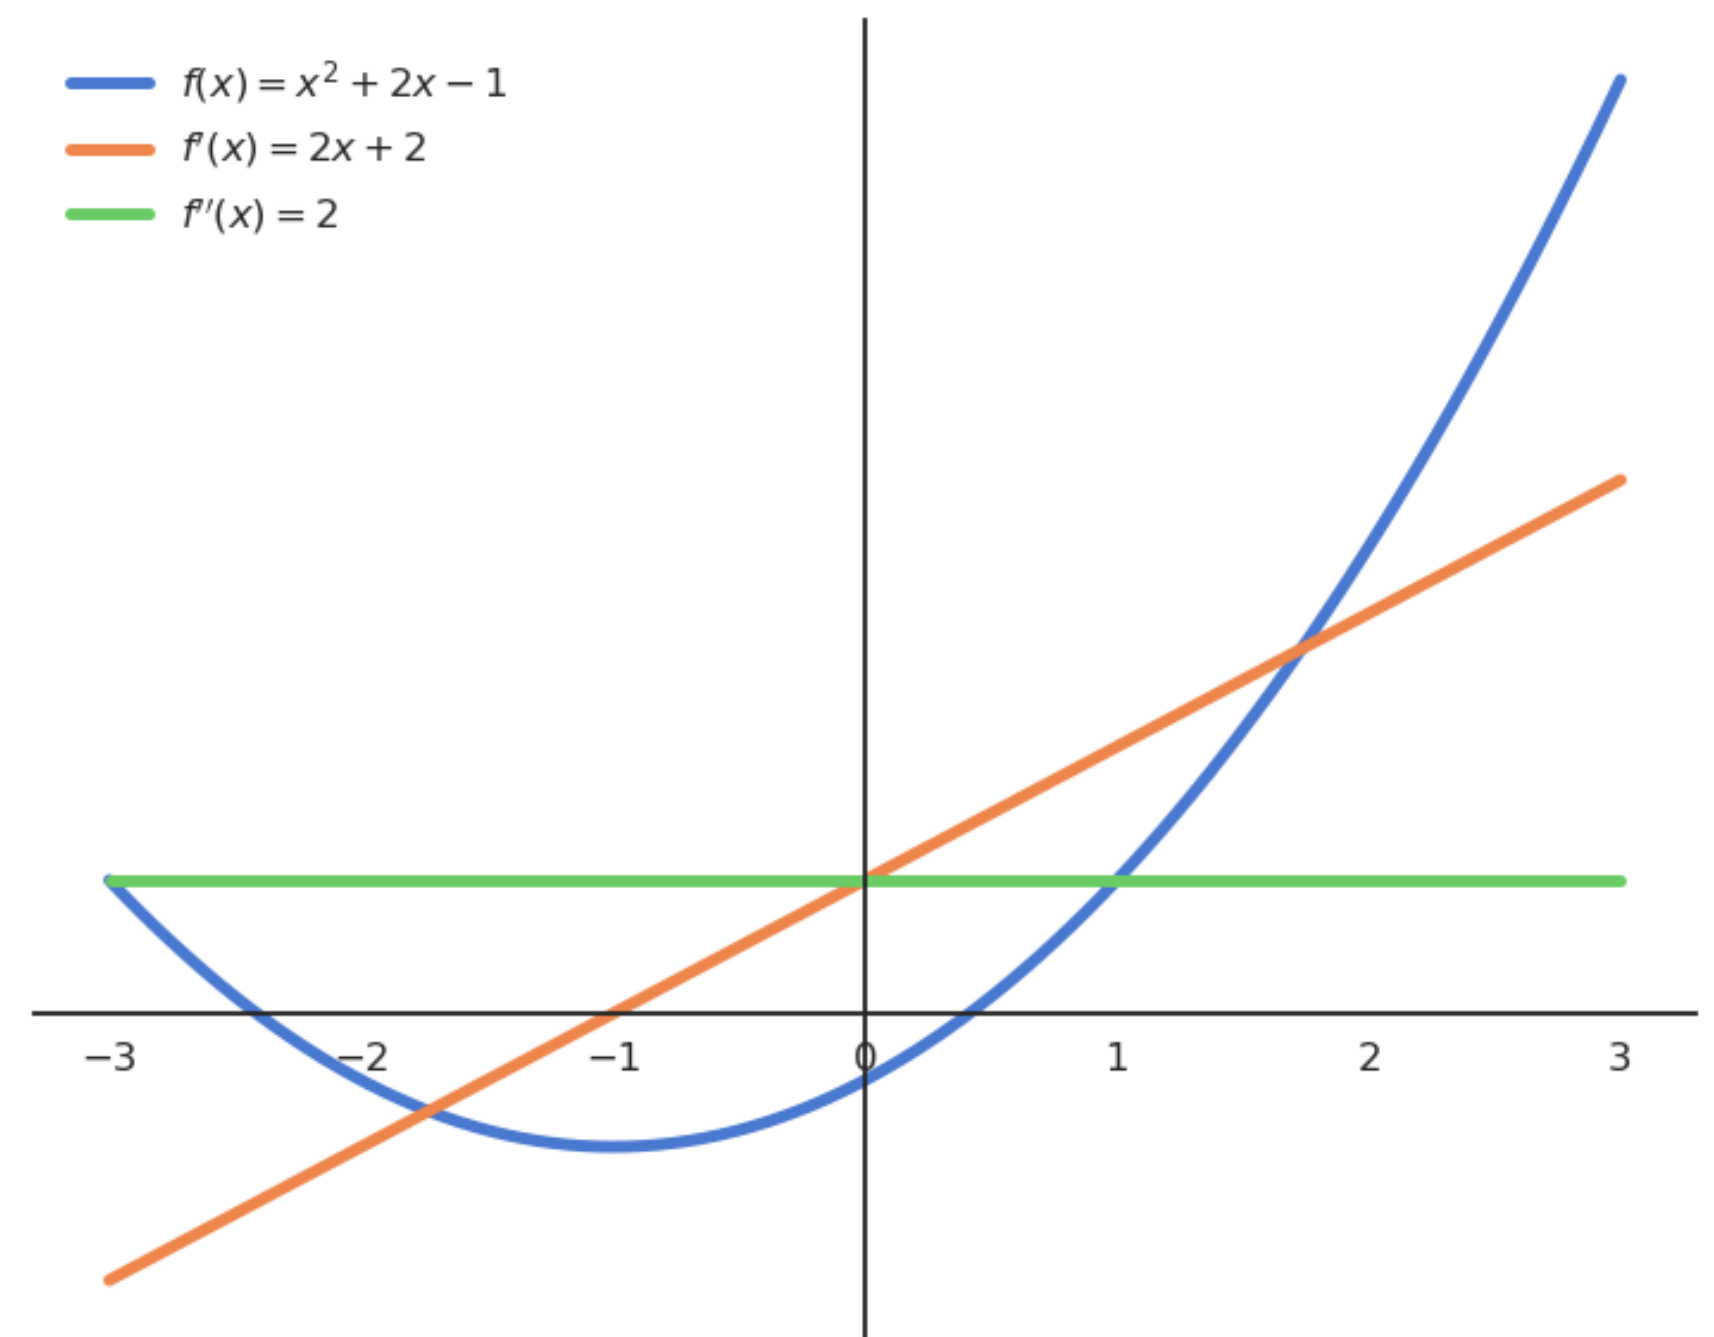
- The sign of the derivative tells you how the function behaves locally:
  - If the derivative is positive, increasing a little bit  $x$  increases the function  $f(x)$ , so the function is **locally increasing**.
  - If the derivative is negative, increasing a little bit  $x$  decreases the function  $f(x)$ , so the function is **locally decreasing**.
- It basically allows you to measure the local influence of  $x$  on  $f(x)$ : if I change a little bit the value  $x$ , what happens to  $f(x)$ ? This will be very useful in machine learning.





# Extrema

- A special case is when the derivative is equal to 0 in  $x$ .  $x$  is then called an **extremum** (or optimum) of the function, i.e. it can be a maximum or minimum.
- You can tell whether an extremum is a maximum or a minimum by looking at its second-order derivative:
  - If  $f''(x) > 0$ , the extremum is a **minimum**.
  - If  $f''(x) < 0$ , the extremum is a **maximum**.
  - If  $f''(x) = 0$ , the extremum is a **saddle point**.



# Gradients

- The derivative of a **multivariate function**  $f(\mathbf{x})$  is a vector of partial derivatives called the **gradient of the function**  $\nabla_{\mathbf{x}} f(\mathbf{x})$ :

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f(\mathbf{x})}{\partial x_1} \\ \frac{\partial f(\mathbf{x})}{\partial x_2} \\ \dots \\ \frac{\partial f(\mathbf{x})}{\partial x_n} \end{bmatrix}$$

- The subscript to the  $\nabla$  operator denotes *with respect to* (w.r.t) which variable the differentiation is done.

# Partial derivatives

- A **partial derivative** w.r.t. to particular variable (or element of the vector) is simply achieved by differentiating the function while considering all other variables to be **constant**. For example the function:

$$f(x, y) = x^2 + 3xy + 4xy^2 - 1$$

can be partially differentiated w.r.t.  $x$  and  $y$  as:

$$\begin{cases} \frac{\partial f(x, y)}{\partial x} = 2x + 3y + 4y^2 \\ \frac{\partial f(x, y)}{\partial y} = 3x + 8xy \end{cases}$$

# Jacobian

- The gradient can be generalized to **vector fields**, where the **Jacobian** or **Jacobi matrix** is a matrix containing all partial derivatives.

$$J = \left[ \frac{\partial \mathbf{f}}{\partial x_1} \quad \dots \quad \frac{\partial \mathbf{f}}{\partial x_n} \right] = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \dots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \dots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

# Analytical properties

- Differentiation is linear, which means that if we define the function:

$$h(x) = a f(x) + b g(x)$$

its derivative is:

$$h'(x) = a f'(x) + b g'(x)$$

- A product of functions can also be differentiated analytically (product rule):

$$(f(x) \times g(x))' = f'(x) \times g(x) + f(x) \times g'(x)$$

**Example:**

$$f(x) = x^2 e^x$$

$$f'(x) = 2x e^x + x^2 \cdot e^x$$

# Chain rule

- A very important concept for neural networks is the **chain rule**, which tells how to differentiate **function compositions** (functions of a function) of the form:

$$(f \circ g)(x) = f(g(x))$$

- The derivative of  $f \circ g$  is:

$$(f \circ g)'(x) = (f' \circ g)(x) \times g'(x)$$

- The chain rule may be more understandable using Leibniz's notation:

$$\frac{d(f \circ g)(x)}{dx} = \frac{df(g(x))}{dg(x)} \times \frac{dg(x)}{dx}$$

- By posing  $y = g(x)$  as an intermediary variable, it becomes:

$$\frac{df(y)}{dx} = \frac{df(y)}{dy} \times \frac{dy}{dx}$$

# Chain rule

- The function :

$$h(x) = \frac{1}{2x + 1}$$

is the function composition of  $g(x) = 2x + 1$  and  $f(x) = \frac{1}{x}$ , whose derivatives are known:

$$g'(x) = 2$$

$$f'(x) = -\frac{1}{x^2}$$

- Its derivative according to the **chain rule** is:

$$h'(x) = f'(g(x)) \times g'(x) = -\frac{1}{(2x + 1)^2} \times 2$$

## Chain rule

- The chain rule also applies to partial derivatives:

$$\frac{\partial f \circ g(x, y)}{\partial x} = \frac{\partial f \circ g(x, y)}{\partial g(x, y)} \times \frac{\partial g(x, y)}{\partial x}$$

and gradients:

$$\nabla_{\mathbf{x}} f \circ g(\mathbf{x}) = \nabla_{g(\mathbf{x})} f \circ g(\mathbf{x}) \times \nabla_{\mathbf{x}} g(\mathbf{x})$$



# Integrals

- The opposite operation of differentiation is **integration**. Given a function  $f(x)$ , we search a function  $F(x)$  whose *derivative* is  $f(x)$ :

$$F'(x) = f(x)$$

- The **integral** of  $f$  is noted:

$$F(x) = \int f(x) dx$$


$dx$  being an infinitesimal interval (similar to  $h$  in the definition of the derivative).

- There are tons of formal definitions of integrals (Riemann, Lebesgue, Darboux...) and we will not get into details here as we will not use integrals a lot.

# Integrals

- The most important to understand for now is maybe that the integral of a function is the **area under the curve**.
- The area under the curve of a function  $f$  on the interval  $[a, b]$  is:


$$S = \int_a^b f(x) dx$$



Source: <https://www.math24.net/riemann-sums-definite-integral/>

# Integrals

- One way to approximate this surface is to split the interval  $[a, b]$  into  $n$  intervals of width  $dx$  with the points  $x_1, x_2, \dots, x_n$ .
- This defines  $n$  rectangles of width  $dx$  and height  $f(x_i)$ , so their surface is  $f(x_i) dx$ .
- The area under the curve can then be approximated by the sum of the surfaces of all these rectangles.




Source: <https://www.math24.net/riemann-sums-definite-integral/>

# Integrals

- When  $n \rightarrow \infty$ , or equivalently  $dx \rightarrow 0$ , the sum of these rectangular areas (called the Riemann sum) becomes exactly the area under the curve. This is the definition of the definite integral:

$$\int_a^b f(x) dx = \lim_{dx \rightarrow 0} \sum_{i=1}^n f(x_i) dx$$


- Very roughly speaking, the integral can be considered as the equivalent of a sum for continuous functions.



Source: <https://www.math24.net/riemann-sums-definite-integral/>

## 3 - Probability theory

# Discrete probability distributions



Credit:

<https://commons.wikimedia.org/wiki/File:2-Dice-Icon.svg>

- Let's note  $X$  a **discrete random variable** with  $n$  realizations (or outcomes)  $x_1, \dots, x_n$ .
- The **probability** that  $X$  takes the value  $x_i$  is defined by the *relative frequency of occurrence*, i.e. the proportion of samples having the value  $x_i$ , when the total number  $N$  of samples tends to infinity:

$$P(X = x_i) = \frac{\text{Number of favorable cases}}{\text{Total number of samples}}$$

- The set of probabilities  $\{P(X = x_i)\}_{i=1}^n$  define the **probability distribution** for the random variable (or probability mass function, pmf).
- By definition, we have  $0 \leq P(X = x_i) \leq 1$  and the probabilities **have** to respect:

$$\sum_{i=1}^n P(X = x_i) = 1$$

# Mathematical expectation and variance

- An important metric for a random variable is its **mathematical expectation** or expected value, i.e. its “mean” realization weighted by the probabilities:

$$\mathbb{E}[X] = \sum_{i=1}^n P(X = x_i) x_i$$

- The expectation does not even need to be a valid realization:

$$\mathbb{E}[\text{Coin}] = \frac{1}{2} 0 + \frac{1}{2} 1 = 0.5$$

$$\mathbb{E}[\text{Dice}] = \frac{1}{6} (1 + 2 + 3 + 4 + 5 + 6) = 3.5$$

- We can also compute the mathematical expectation of **functions of** a random variable:

$$\mathbb{E}[f(X)] = \sum_{i=1}^n P(X = x_i) f(x_i)$$

# Mathematical expectation and variance

- The **variance** of a random variable is the squared deviation around the mean:

$$\text{Var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \sum_{i=1}^n P(X = x_i) (x_i - \mathbb{E}[X])^2$$

- Variance of a coin:


$$\text{Var}(\text{Coin}) = \frac{1}{2} (0 - 0.5)^2 + \frac{1}{2} (1 - 0.5)^2 = 0.25$$

- Variance of a dice:

$$\text{Var}(\text{Dice}) = \frac{1}{6} ((1 - 3.5)^2 + (2 - 3.5)^2 + (3 - 3.5)^2 + (4 - 3.5)^2 + (5 - 3.5)^2 + (6 - 3.5)^2) = \frac{105}{36}$$



# Continuous probability distributions



Source: [https://en.wikipedia.org/wiki/Normal\\_distribution](https://en.wikipedia.org/wiki/Normal_distribution)

- **Continuous random variables** can take an infinity of continuous values, e.g.  $\mathcal{R}$  or some subset.
- The closed set of values they can take is called the **support**  $\mathcal{D}_X$  of the probability distribution.
- The probability distribution is described by a **probability density function** (pdf)  $f(x)$ .
- The pdf of a distribution must be positive ( $f(x) \geq 0 \forall x \in \mathcal{D}_X$ ) and its integral must be equal to 1:

$$\int_{x \in \mathcal{D}_X} f(x) dx = 1$$

- The pdf does not give the probability of taking a particular value  $x$  (it is 0), but allows to get the probability that a value lies in a specific interval:

$$P(a \leq X \leq b) = \int_a^b f(x) dx$$

- One can however think of the pdf as the **likelihood** that a value  $x$  comes from that distribution.

# Expectation and variance of continuous distributions

- The mathematical expectation is now defined by an integral instead of a sum:

$$\mathbb{E}[X] = \int_{x \in \mathcal{D}_X} f(x) x \, dx$$

the variance:

$$\text{Var}(X) = \int_{x \in \mathcal{D}_X} f(x) (x - \mathbb{E}[X])^2 \, dx$$

or a function of the random variable:

$$\mathbb{E}[g(X)] = \int_{x \in \mathcal{D}_X} f(x) g(x) \, dx$$

- Note that the expectation operator is **linear**:

$$\mathbb{E}[a X + b Y] = a \mathbb{E}[X] + b \mathbb{E}[Y]$$

## Some parameterized probability distributions

- Probability distributions can in principle have any form:  $f(x)$  is unknown.
- However, specific parameterized distributions can be very useful: their pmf/pdf is fully determined by a couple of parameters.
- The **Bernoulli** distribution is a binary (discrete, 0 or 1) distribution with a parameter  $p$  specifying the probability to obtain the outcome 1:

$$P(X = 1) = p \text{ and } P(X = 0) = 1 - p$$

$$P(X = x) = p^x (1 - p)^{1-x}$$


$$\mathbb{E}[X] = p$$

- The **Multinoulli** or **categorical** distribution is a discrete distribution with  $k$  realizations. Each realization  $x_i$  is associated with a parameter  $p_i > 0$  representing its probability. We have  $\sum_i p_i = 1$ .

$$P(X = x_i) = p_i$$

- Knowing  $p$  or the  $p_i$  tells us everything about the discrete distributions.

# The uniform distribution



- The **uniform distribution** has an equal and constant probability of returning values between  $a$  and  $b$ , never outside this range.
- It is parameterized by two parameters:
  - the start of the range  $a$ .
  - the end of the range  $b$ .
- Its support is  $[a, b]$ .


Credit: [https://en.wikipedia.org/wiki/Uniform\\_distribution\\_\(continuous\)](https://en.wikipedia.org/wiki/Uniform_distribution_(continuous))

- The pdf of the uniform distribution  $\mathcal{U}(a, b)$  is defined on  $[a, b]$  as:

$$f(x; a, b) = \frac{1}{b - a}$$

- Knowing  $a$  and  $b$  completely defines the distribution.

# The normal or Gaussian distribution



- For continuous distributions, the **normal distribution** is the most frequently encountered one.
- It is parameterized by two parameters:
  - the mean  $\mu$ .
  - the variance  $\sigma^2$  (or standard deviation  $\sigma$ ).
- Its support is  $\mathfrak{R}$ .


Credit: [https://en.wikipedia.org/wiki/Normal\\_distribution](https://en.wikipedia.org/wiki/Normal_distribution)

- The pdf of the normal distribution  $\mathcal{N}(\mu, \sigma)$  is defined on  $\mathfrak{R}$  as:

$$f(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- Knowing  $\mu$  and  $\sigma$  completely defines the distribution.

# The exponential distribution



- The **exponential distribution** is the probability distribution of the time between events in a Poisson point process, i.e., a process in which events occur continuously and independently at a constant average rate.
- It is parameterized by one parameter:
  - the rate  $\lambda$ .
- Its support is  $\mathfrak{R}^+$  ( $x > 0$ ).

Credit: [https://en.wikipedia.org/wiki/Exponential\\_distribution](https://en.wikipedia.org/wiki/Exponential_distribution)

- The pdf of the exponential distribution is defined on  $\mathfrak{R}^+$  as:

$$f(x; \lambda) = \lambda e^{-\lambda x}$$

- Knowing  $\lambda$  completely defines the distribution.

# Joint probabilities

- Let's now suppose that we have two random variables  $X$  and  $Y$  with different probability distributions  $P(X)$  and  $P(Y)$ .
- The **joint probability**  $P(X, Y)$  denotes the probability of observing the realizations  $x$  **and**  $y$  at the same time:

$$P(X = x, Y = y)$$

- If the random variables are **independent**, we have:

$$P(X = x, Y = y) = P(X = x) P(Y = y)$$

- If you know the joint probability, you can compute the **marginal probability distribution** of each variable:

$$P(X = x) = \sum_y P(X = x, Y = y)$$

- The same is true for continuous probability distributions:

$$f(x) = \int f(x, y) dy$$

# Conditional probabilities

- Some useful information between two random variables is the **conditional probability**.
- $P(X = x|Y = y)$  is the conditional probability that  $X = x$ , **given** that  $Y = y$  is observed.
- $Y = y$  is not random anymore: it is a **fact** (at least theoretically).
- You wonder what happens to the probability distribution of  $X$  now that you know the value of  $Y$ .
- Conditional probabilities are linked to the joint probability by:


$$P(X = x|Y = y) = \frac{P(X = x, Y = y)}{P(Y = y)}$$

- If  $X$  and  $Y$  are **independent**, we have  $P(X = x|Y = y) = P(X = x)$  (knowing  $Y$  does not change anything to the probability distribution of  $X$ ).
- We can use the same notation for the complete probability distributions:

$$P(X|Y) = \frac{P(X, Y)}{P(Y)}$$



# Joint and conditional probabilities: using a Venn diagram



Credit: <https://www.elewise.co.uk/g-e-m-h-5-u.html>

- You ask 50 people whether they like cats or dogs:
  - 18 like both cats and dogs.
  - 21 like only dogs.
  - 5 like only cats.
  - 6 like none of them.
- We consider loving cats and dogs as random variables (and that our sample size is big enough to use probabilities...)
- We have  $P(\text{dog}) = \frac{18+21}{50}$  and  $P(\text{cat}) = \frac{18+5}{50}$ .
- Among the 23 who love cats, which proportion also loves dogs?

- The joint probability of loving both cats and dogs is  $P(\text{cat}, \text{dog}) = \frac{18}{50}$ .
- The conditional probability of loving dogs given one loves cats is:

$$P(\text{dog}|\text{cat}) = \frac{P(\text{cat}, \text{dog})}{P(\text{cat})} = \frac{\frac{18}{50}}{\frac{23}{50}} = \frac{18}{23}$$

# Bayes' rule

- Noticing that the definition of conditional probabilities is symmetric:

$$P(X, Y) = P(X|Y) P(Y) = P(Y|X) P(X)$$

we can obtain the **Bayes' rule**:

$$P(Y|X) = \frac{P(X|Y) P(Y)}{P(X)}$$

- It is very useful when you already know  $P(X|Y)$  and want to obtain  $P(Y|X)$  (**Bayesian inference**).
  - $P(Y|X)$  is called the **posterior probability**.
  - $P(X|Y)$  is called the **likelihood**.
  - $P(Y)$  is called the **prior probability** (belief).
  - $P(X)$  is called the **model evidence** or **marginal likelihood**.

## Bayes' rule : example

- Let's consider a disease  $D$  (binary random variable) and a medical test  $T$  (also binary). The disease affects 10% of the general population:

$$P(D = 1) = 0.1$$

$$P(D = 0) = 0.9$$

- When a patient has the disease, the test is positive 80% of the time:

$$P(T = 1|D = 1) = 0.8$$

$$P(T = 0|D = 1) = 0.2$$

- When a patient does not have the disease, the test is still positive 10% of the time:

$$P(T = 1|D = 0) = 0.1$$

$$P(T = 0|D = 0) = 0.9$$

- Given that the test is positive, what is the probability that the patient is ill?

## Bayes' rule : example


$$\begin{aligned}P(D = 1|T = 1) &= \frac{P(T = 1|D = 1) P(D = 1)}{P(T = 1)} \\&= \frac{P(T = 1|D = 1) P(D = 1)}{P(T = 1|D = 1) P(D = 1) + P(T = 1|D = 0) P(D = 0)} \\&= \frac{0.8 \times 0.1}{0.8 \times 0.1 + 0.1 \times 0.9} \\&= 0.47\end{aligned}$$

# 4 - Statistics

# Random sampling / Monte Carlo sampling

- In ML, we will deal with random variables whose exact probability distribution is unknown, but we are interested in their expectation or variance anyway.
- **Random sampling** or **Monte Carlo sampling** (MC) consists of taking  $N$  samples  $x_i$  out of the distribution  $X$  (discrete or continuous) and computing the **sample average**:

$$\mathbb{E}[X] = \mathbb{E}_{x \sim X}[x] \approx \frac{1}{N} \sum_{i=1}^N x_i$$



- More samples will be obtained where  $f(x)$  is high ( $x$  is probable), so the average of the sampled data will be close to the expected value of the distribution.

# Random sampling / Monte Carlo sampling

## Law of big numbers

As the number of identically distributed, randomly generated variables increases, their sample mean (average) approaches their theoretical mean.

MC estimates are only correct when:

- the samples are **i.i.d** (independent and identically distributed):
  - independent: the samples must be unrelated with each other.
  - identically distributed: the samples must come from the same distribution  $X$ .
- the number of samples is large enough. Usually  $N > 30$  for simple distributions.




# Random sampling / Monte Carlo sampling

- One can estimate any function of the random variable with random sampling:


$$\mathbb{E}[f(X)] = \mathbb{E}_{x \sim X}[f(x)] \approx \frac{1}{N} \sum_{i=1}^N f(x_i)$$

- Example of Monte Carlo sampling to estimate  $\pi/4$ :


Estimate of pi: 3.116




Estimate of pi: 3.142



Estimate of pi: 3.13872



Estimate of pi: 3.141932





# Central limit theorem

- Suppose we have an unknown distribution  $X$  with expected value  $\mu = \mathbb{E}[X]$  and variance  $\sigma^2$ .
- We can take randomly  $N$  samples from  $X$  to compute the sample average:

$$S_N = \frac{1}{N} \sum_{i=1}^N x_i$$


- The **Central Limit Theorem** (CLT) states that:

The distribution of sample averages is normally distributed with mean  $\mu$  and variance  $\frac{\sigma^2}{N}$ .



$$S_N \sim \mathcal{N}\left(\mu, \frac{\sigma}{\sqrt{N}}\right)$$

# Central limit theorem

- If we perform the sampling multiple times, even with few samples, the average of the sampling averages will be very close to the expected value.
- The more samples we get, the smaller the variance of the estimates.
- Although the distribution  $X$  can be anything, the sampling averages are normally distributed.



population  
distribution



sampling distribution  
of the mean

Credit: [https://en.wikipedia.org/wiki/Central\\_limit\\_theorem](https://en.wikipedia.org/wiki/Central_limit_theorem)


# Estimators

- CLT shows that the sampling average is an **unbiased estimator** of the expected value of a distribution:

$$\mathbb{E}(S_N) = \mathbb{E}(X)$$


- An estimator is a random variable used to measure parameters of a distribution (e.g. its expectation). The problem is that estimators can generally be **biased**.
- Take the example of a thermometer  $M$  measuring the temperature  $T$ .  $T$  is a random variable (normally distributed with  $\mu = 20$  and  $\sigma = 10$ ) and the measurements  $M$  relate to the temperature with the relation:

$$M = 0.95 T + 0.65$$



# Estimators

- The thermometer is not perfect, but do random measurements allow us to estimate the expected value of the temperature?
- We could repeatedly take 100 random samples of the thermometer and see how the distribution of sample averages look like:



- But, as the expectation is linear, we actually have:

$$\mathbb{E}[M] = \mathbb{E}[0.95 T + 0.65] = 0.95 \mathbb{E}[T] + 0.65 = 19.65 \neq \mathbb{E}[T]$$

- The thermometer is a **biased estimator** of the temperature.

# Estimators

- Let's note  $\theta$  a parameter of a probability distribution  $X$  that we want to estimate (it does not have to be its mean).
- An **estimator**  $\hat{\theta}$  is a random variable mapping the sample space of  $X$  to a set of sample estimates.
- The **bias** of an estimator is the mean error made by the estimator:

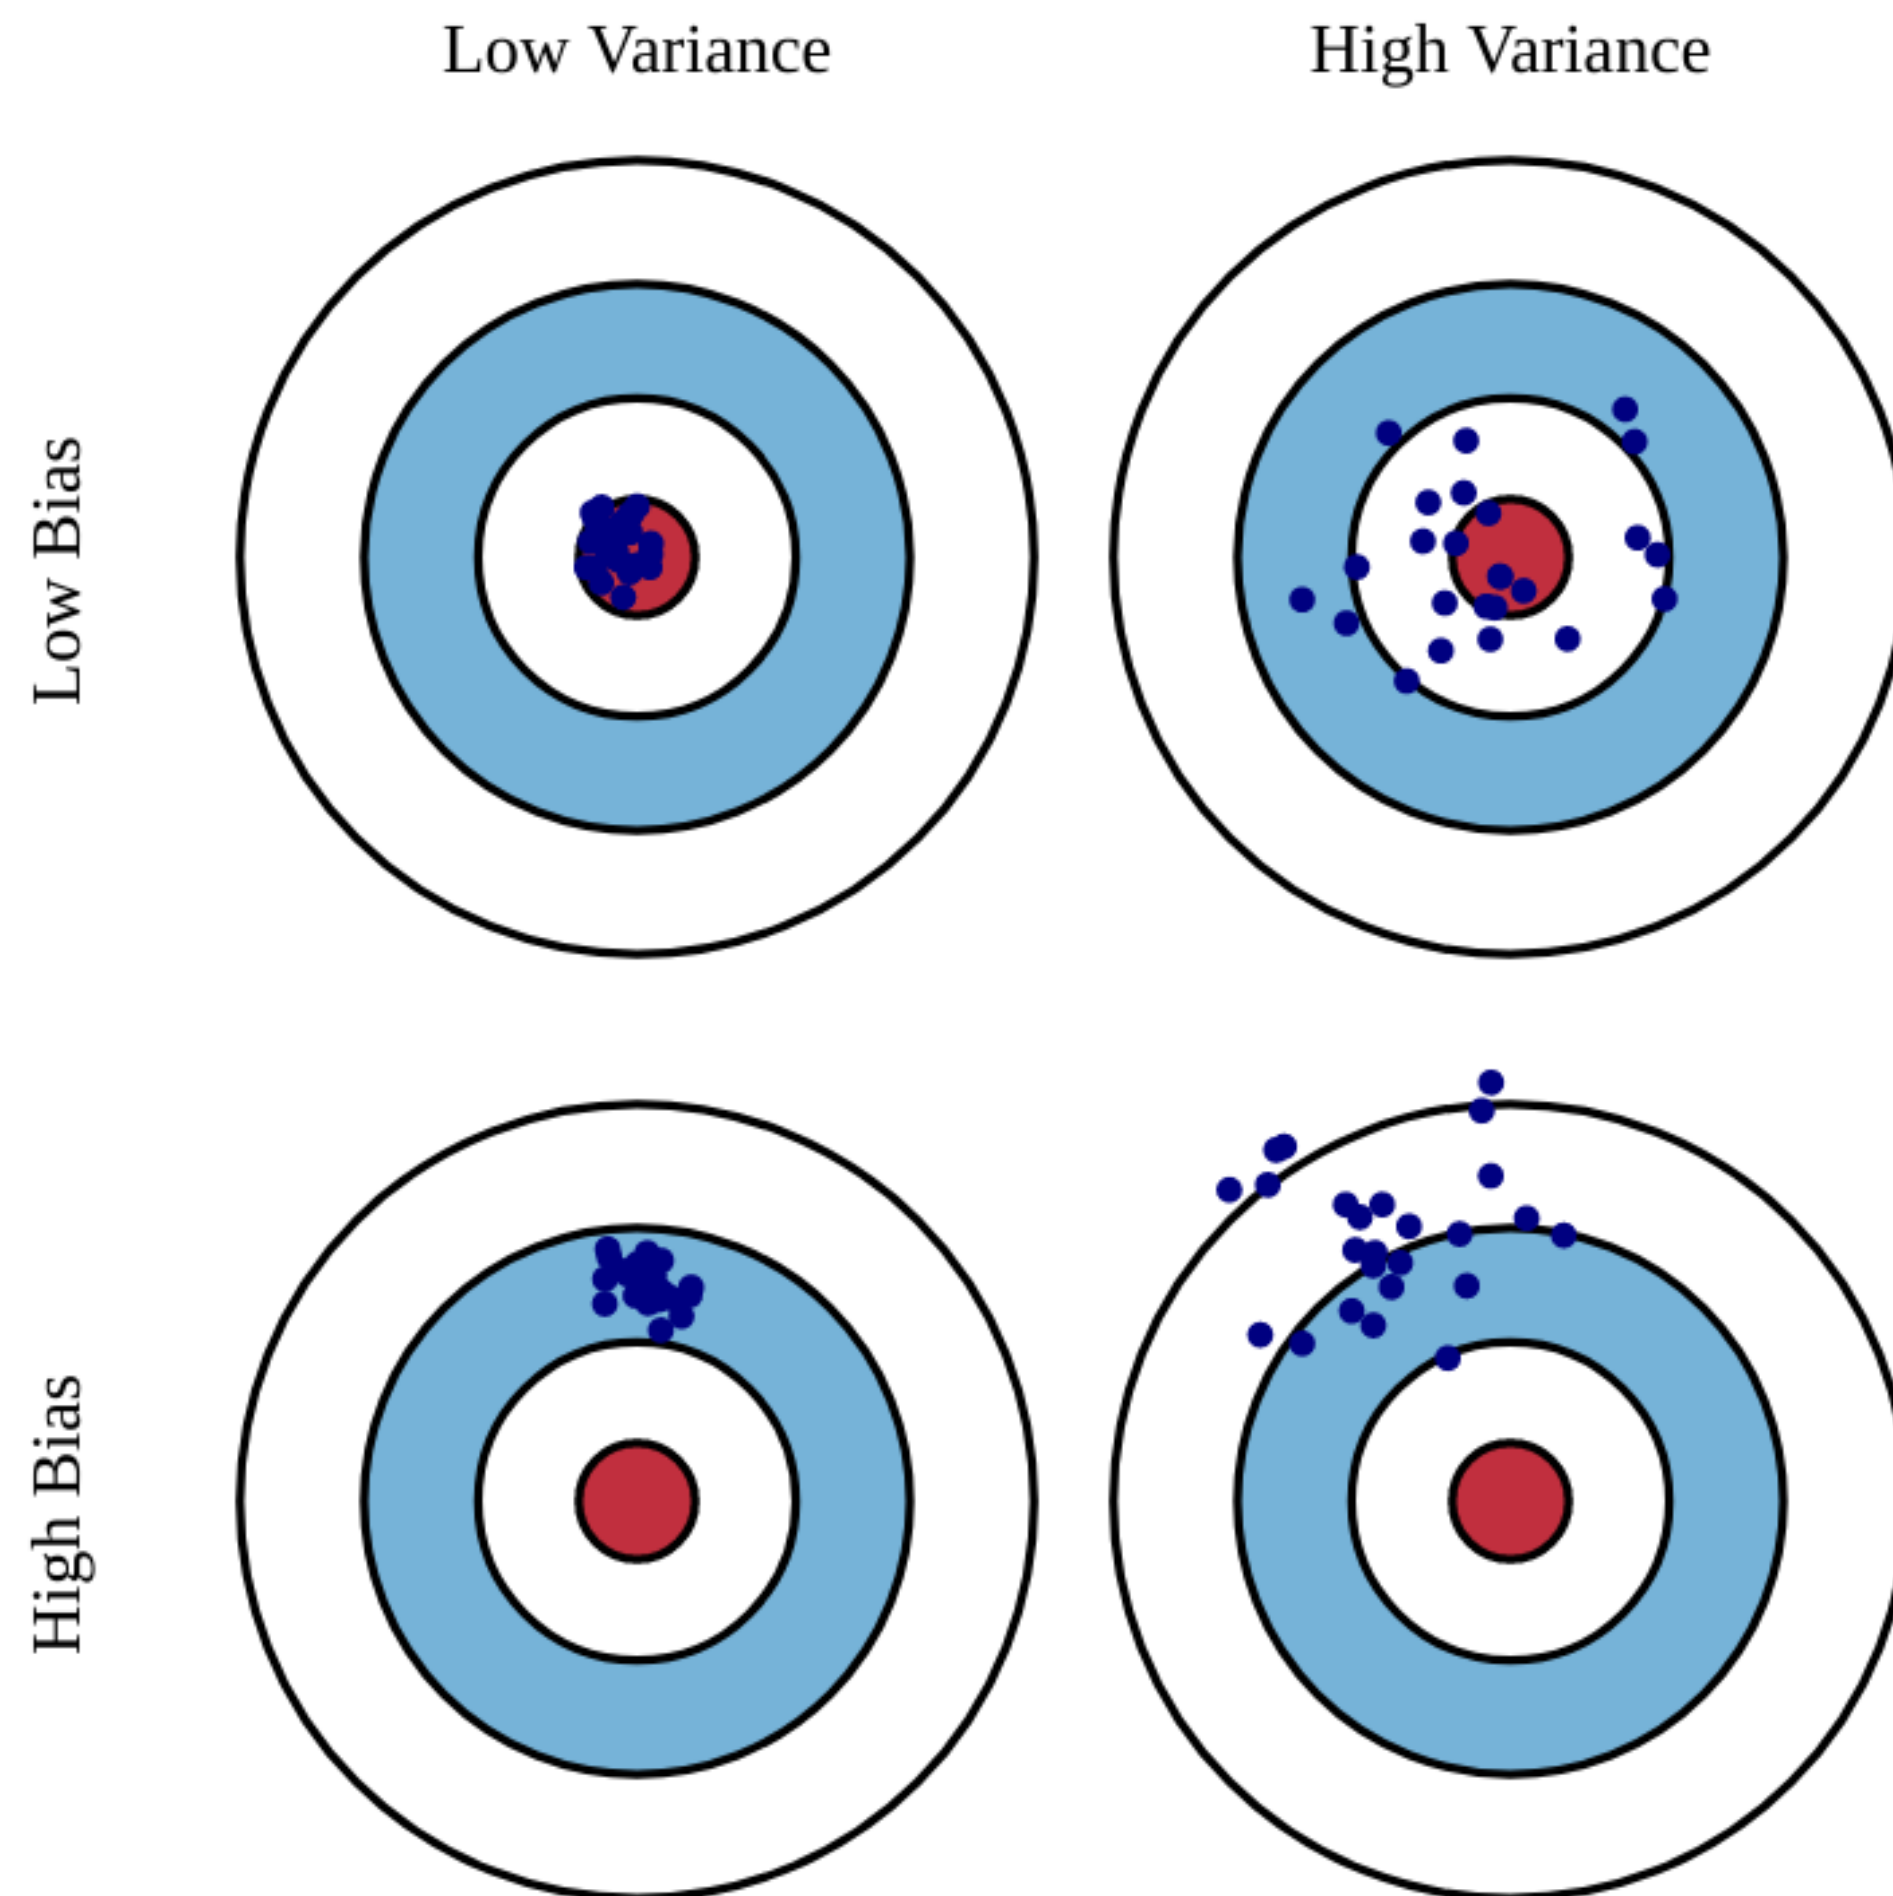
$$\mathcal{B}(\hat{\theta}) = \mathbb{E}[\hat{\theta} - \theta] = \mathbb{E}[\hat{\theta}] - \theta$$

- The **variance** of an estimator is the deviation of the samples around the expected value:

$$\text{Var}(\hat{\theta}) = \mathbb{E}[(\hat{\theta} - \mathbb{E}[\hat{\theta}])^2]$$

- Ideally, we would like estimators with:
  - **low bias**: the estimations are correct on average (= equal to the true parameter).
  - **low variance**: we do not need many estimates to get a correct estimate (CLT:  $\frac{\sigma}{\sqrt{N}}$ )

# Estimators: bias and variance




- Unfortunately, the perfect estimator does not exist.
- Estimators will have a bias and a variance:
  - **Bias:** the estimated values will be wrong, and the policy not optimal.
  - **Variance:** we will need a lot of samples (trial and error) to have correct estimates.
- One usually talks of a **bias/variance** trade-off: if you have a small bias, you will have a high variance, or vice versa.
- In machine learning, bias corresponds to underfitting, variance to overfitting.

## 5 - Information theory



# Information

- **Information theory** (Claude Shannon) asks how much information is contained in a probability distribution.
- Information is related to **surprise** or **uncertainty**: are the outcomes of a random variable surprising?
  - Almost certain outcomes ( $P \sim 1$ ) are not surprising because they happen all the time.
  - Almost impossible outcomes ( $P \sim 0$ ) are very surprising because they are very rare.



- A useful measurement of how surprising is an outcome  $x$  is the **self-information**:

$$I(x) = -\log P(X = x)$$

- Depending on which log is used, self-information has different units:
  - $\log_2$ : bits or shannons.
  - $\log_e = \ln$ : nats.
- But it is just a rescaling, the base never matters.



# Entropy

- The **entropy** (or Shannon entropy) of a probability distribution is the expected value of the self-information of its outcomes:

$$H(X) = \mathbb{E}_{x \sim X}[I(x)] = \mathbb{E}_{x \sim X}[-\log P(X = x)]$$


- It measures the **uncertainty, randomness** or **information content** of the random variable.
- In the discrete case:

$$H(X) = - \sum_x P(x) \log P(x)$$

- In the continuous case:

$$H(X) = - \int_x f(x) \log f(x) dx$$

- The entropy of a Bernoulli variable is maximal when both outcomes are **equiprobable**.
- If a variable is **deterministic**, its entropy is minimal and equal to zero.



# Joint and conditional entropies

- The **joint entropy** of two random variables  $X$  and  $Y$  is defined by:

$$H(X, Y) = \mathbb{E}_{x \sim X, y \sim Y}[-\log P(X = x, Y = y)]$$

- The **conditional entropy** of two random variables  $X$  and  $Y$  is defined by:

$$H(X|Y) = \mathbb{E}_{x \sim X, y \sim Y}[-\log P(X = x|Y = y)] = \mathbb{E}_{x \sim X, y \sim Y}[-\log \frac{P(X = x, Y = y)}{P(Y = y)}]$$

- If the variables are **independent**, we have:

$$H(X, Y) = H(X) + H(Y) \quad \text{or} \quad H(X|Y) = H(X)$$

- Both are related by:

$$H(X|Y) = H(X, Y) - H(Y)$$

- The equivalent of Bayes' rule is:

$$H(Y|X) = H(X|Y) + H(Y) - H(X)$$

# Mutual Information

- The most important information measurement between two variables is the **mutual information** MI (or information gain):

$$I(X, Y) = H(X) - H(X|Y) = H(Y) - H(Y|X)$$

- It measures how much information the variable  $X$  holds on  $Y$ :
  - If the two variables are **independent**, the MI is 0 :  $X$  is as random, whether you know  $Y$  or not.

$$I(X, Y) = 0$$

- If the two variables are **dependent**, knowing  $Y$  gives you information on  $X$ , which becomes less random, i.e. less uncertain / surprising.

$$I(X, Y) > 0$$


- If you can fully predict  $X$  when you know  $Y$ , it becomes deterministic ( $H(X|Y) = 0$ ) so the mutual information is maximal ( $I(X, Y) = H(X)$ ).

# Cross-entropy


- The **cross-entropy** between two distributions  $X$  and  $Y$  is defined as:

$$H(X, Y) = \mathbb{E}_{x \sim X}[-\log P(Y = x)]$$

- Beware that the notation  $H(X, Y)$  is the same as the joint entropy, but it is a different concept!
- The cross-entropy measures the **negative log-likelihood** that a sample  $x$  taken from the distribution  $X$  could also come from the distribution  $Y$ .
- More exactly, it measures how many bits of information one would need to distinguish the two distributions  $X$  and  $Y$ .



# Cross-entropy



$$H(X, Y) = \mathbb{E}_{x \sim X}[-\log P(Y = x)]$$

- If the two distributions are the same *almost anywhere*, one cannot distinguish samples from the two distributions:
  - The cross-entropy is the same as the entropy of  $X$ .
- If the two distributions are completely different, one can tell whether a sample  $Z$  comes from  $X$  or  $Y$ :
  - The cross-entropy is higher than the entropy of  $X$ .

# Kullback-Leibler divergence

- In practice, the **Kullback-Leibler divergence**  $KL(X||Y)$  is a better measurement of the similarity (statistical distance) between two probability distributions:

$$KL(X||Y) = \mathbb{E}_{x \sim X}[-\log \frac{P(Y = x)}{P(X = x)}]$$

- It is linked to the cross-entropy by:

$$KL(X||Y) = H(X, Y) - H(X)$$

- If the two distributions are the same *almost anywhere*:
  - The KL divergence is zero.
- If the two distributions are different:
  - The KL divergence is positive.
- Minimizing the KL between two distributions is the same as making the two distributions “equal”.
- Again, the KL is not a metric, as it is not symmetric.